

Lecture 1:

Why Parallelism? Why Efficiency?

**Parallel Computing
Stanford CS149, Fall 2024**

Hello!



Prof. Kayvon



Prof. Olukotun

One common definition

A parallel computer is a collection of processing elements
that cooperate to solve problems quickly

We care about performance,
and we care about efficiency

We're going to use multiple
processing elements to get it

DEMO 1

(our first parallel program)

Speedup

One major motivation of using parallel processing: achieve a speedup

For a given problem:

speedup(using P processors) =

$$\frac{\text{execution time (using 1 processor)}}{\text{execution time (using P processors)}}$$

Class observations from demo 1

- **Communication limited the maximum speedup achieved**
 - In the demo, the communication was telling each other the partial sums
- **Minimizing the cost of communication improved speedup**
 - Moved students (“processors”) closer together (or let them shout)

DEMO 2

(scaling up to four “processors”)

Class observations from demo 2

- **Imbalance in work assignment limited speedup**
 - Some students (“processors”) ran out work to do (went idle), while others were still working on their assigned task
- **Improving the distribution of work improved speedup**

DEMO 3

(massively parallel execution)

Class observations from demo 3

- The problem I just gave you has a significant amount of communication compared to computation
- Communication costs can dominate a parallel computation, severely limiting speedup

Course theme 1:

Designing and writing parallel programs ... that scale!

- **Parallel thinking**

1. **Decomposing work into pieces that can safely be performed in parallel**
2. **Assigning work to processors**
3. **Managing communication/synchronization between the processors so that it does not limit speedup**

- **Abstractions/mechanisms for performing the above tasks**

- **Writing code in popular parallel programming languages**

Course theme 2:

Parallel computer hardware implementation: how parallel computers work

- **Mechanisms used to implement abstractions efficiently**
 - **Performance characteristics of implementations**
 - **Design trade-offs: performance vs. convenience vs. cost**
- **Why do I need to know about hardware?**
 - **Because the characteristics of the machine really matter
(recall speed of communication issues in earlier demos)**
 - **Because you care about efficiency and performance
(you are writing parallel programs after all!)**

Course theme 3:

Thinking about efficiency

- **FAST != EFFICIENT**
- Just because your program runs faster on a parallel computer, it does not mean it is using the hardware efficiently
 - Is 2x speedup on computer with 10 processors a good result?
- Programmer's perspective: make use of provided machine capabilities
- HW designer's perspective: choosing the right capabilities to put in system (performance/cost, cost = silicon area?, power?, etc.)

Course logistics

Getting started

■ The course web site

- <https://cs149.stanford.edu>

■ Textbook

- **There is no course textbook (the internet is plenty good these days), also see the course web site for suggested references**

Stanford CS149, Fall 2024

PARALLEL COMPUTING

From smart phones, to multi-core CPUs and GPUs, to the world's largest supercomputers and web sites, parallel processing is ubiquitous in modern computing. The goal of this course is to provide a deep understanding of the fundamental principles and engineering trade-offs involved in designing modern parallel computing systems as well as to teach parallel programming techniques necessary to effectively utilize these machines. Because writing good parallel programs requires an understanding of key machine performance characteristics, this course will cover both parallel hardware and software design.

Basic Info

Time: Tues/Thurs 10:30-11:50am

Location: NVIDIA Auditorium

Instructors: **Kayvon Fatahalian** and **Kunle Olukotun**

See the [course info](#) page for more info on policies and logistics.

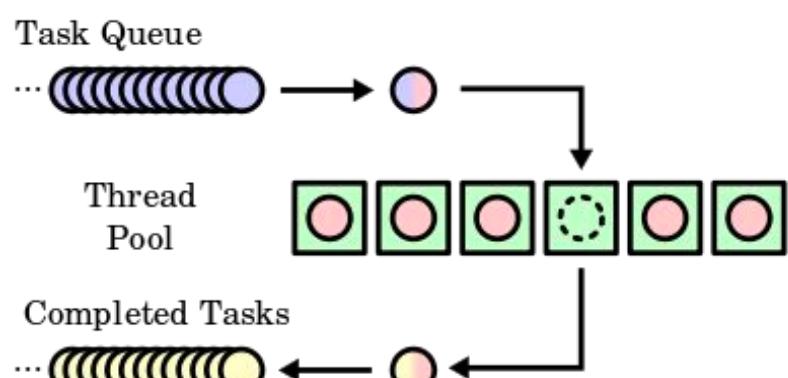
Fall 2024 Schedule

Sep 24	Why Parallelism? Why Efficiency? Challenges of parallelizing code, motivations for parallel chips, processor basics
Sep 26	A Modern Multi-Core Processor Forms of parallelism: multi-core, SIMD, and multi-threading
Oct 01	Multi-Core Architecture Part II + ISPC Programming Abstractions Finish up multi-threaded and latency vs. bandwidth. ISPC programming, abstraction vs. implementation
Oct 03	Parallel Programming Basics Structuring parallel programs. Process of parallelizing a program in data parallel and shared address space models
Oct 08	Performance Optimization I: Work Distribution and Scheduling Achieving good work distribution while minimizing overhead, scheduling Cilk programs with work stealing
Oct 10	Performance Optimization II: Locality, Communication, and Contention Message passing, async vs. blocking sends/receives, pipelining, increasing arithmetic intensity, avoiding contention
Oct 15	GPU Architecture and CUDA Programming CUDA programming abstractions, and how they are implemented on modern GPUs
Oct 17	Data-Parallel Thinking Data-parallel operations like map, reduce, scan, prefix sum, groupByKey
Oct 22	Distributed Data-Parallel Computing Using Spark Producer-consumer locality, RDD abstraction, Spark implementation and scheduling
Oct 24	Efficiently Evaluating DNNs on GPUs Efficiently scheduling DNN layers, mapping convs to matrix-multiplication, transformers, layer fusion

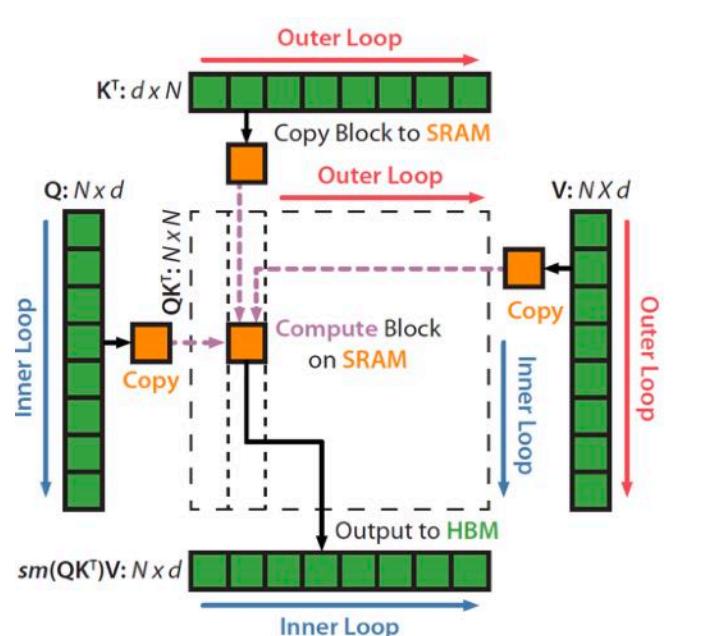
Four programming assignments



**Assignment 1: ISPC programming
on multi-core CPUs**



**Assignment 3: Writing a renderer
in CUDA on NVIDIA GPUs**



**Assignment 4:
optimizing a mini-DNN**



**Optional assignment 5:
(Can be used to boost a prior grade)**

Topic TBD

**Programming assignments can
(optionally) be done with a partner.**

We realize finding a partner can be
stressful. 😱 😓

Fill out our partner request form by
Friday noon and we will find you a
partner! 🎉🤓

Written assignments

- **There will be six written assignments this quarter**
 - **Some of the questions are graded on correctness, others are graded on effort**
- **Written assignments contain modified versions of previous exam questions, so they:**
 - **Give you practice with key course concepts**
 - **Provide practice for the style of questions you will see on an exam**

Late days

- You get **eight late days** for the quarter
 - For use on programming assignments only
- The idea of late days is to give you the flexibility to handle almost all events that arise throughout the quarter
 - Work from other classes, failing behind, most illnesses, athletic/extracurricular events, academic conference travel...
 - We expect to give extra late days only under exceptional circumstances
- Requests for additional late days to accommodate foreseeable exceptional circumstances should be made 72 hours prior to the original assignment deadline.
 - We will deny requests if you could have reasonably planned ahead.

Grades

56% Programming assignments (4)

12% Written assignments (6)

15% Midterm exam

- Nov 14th

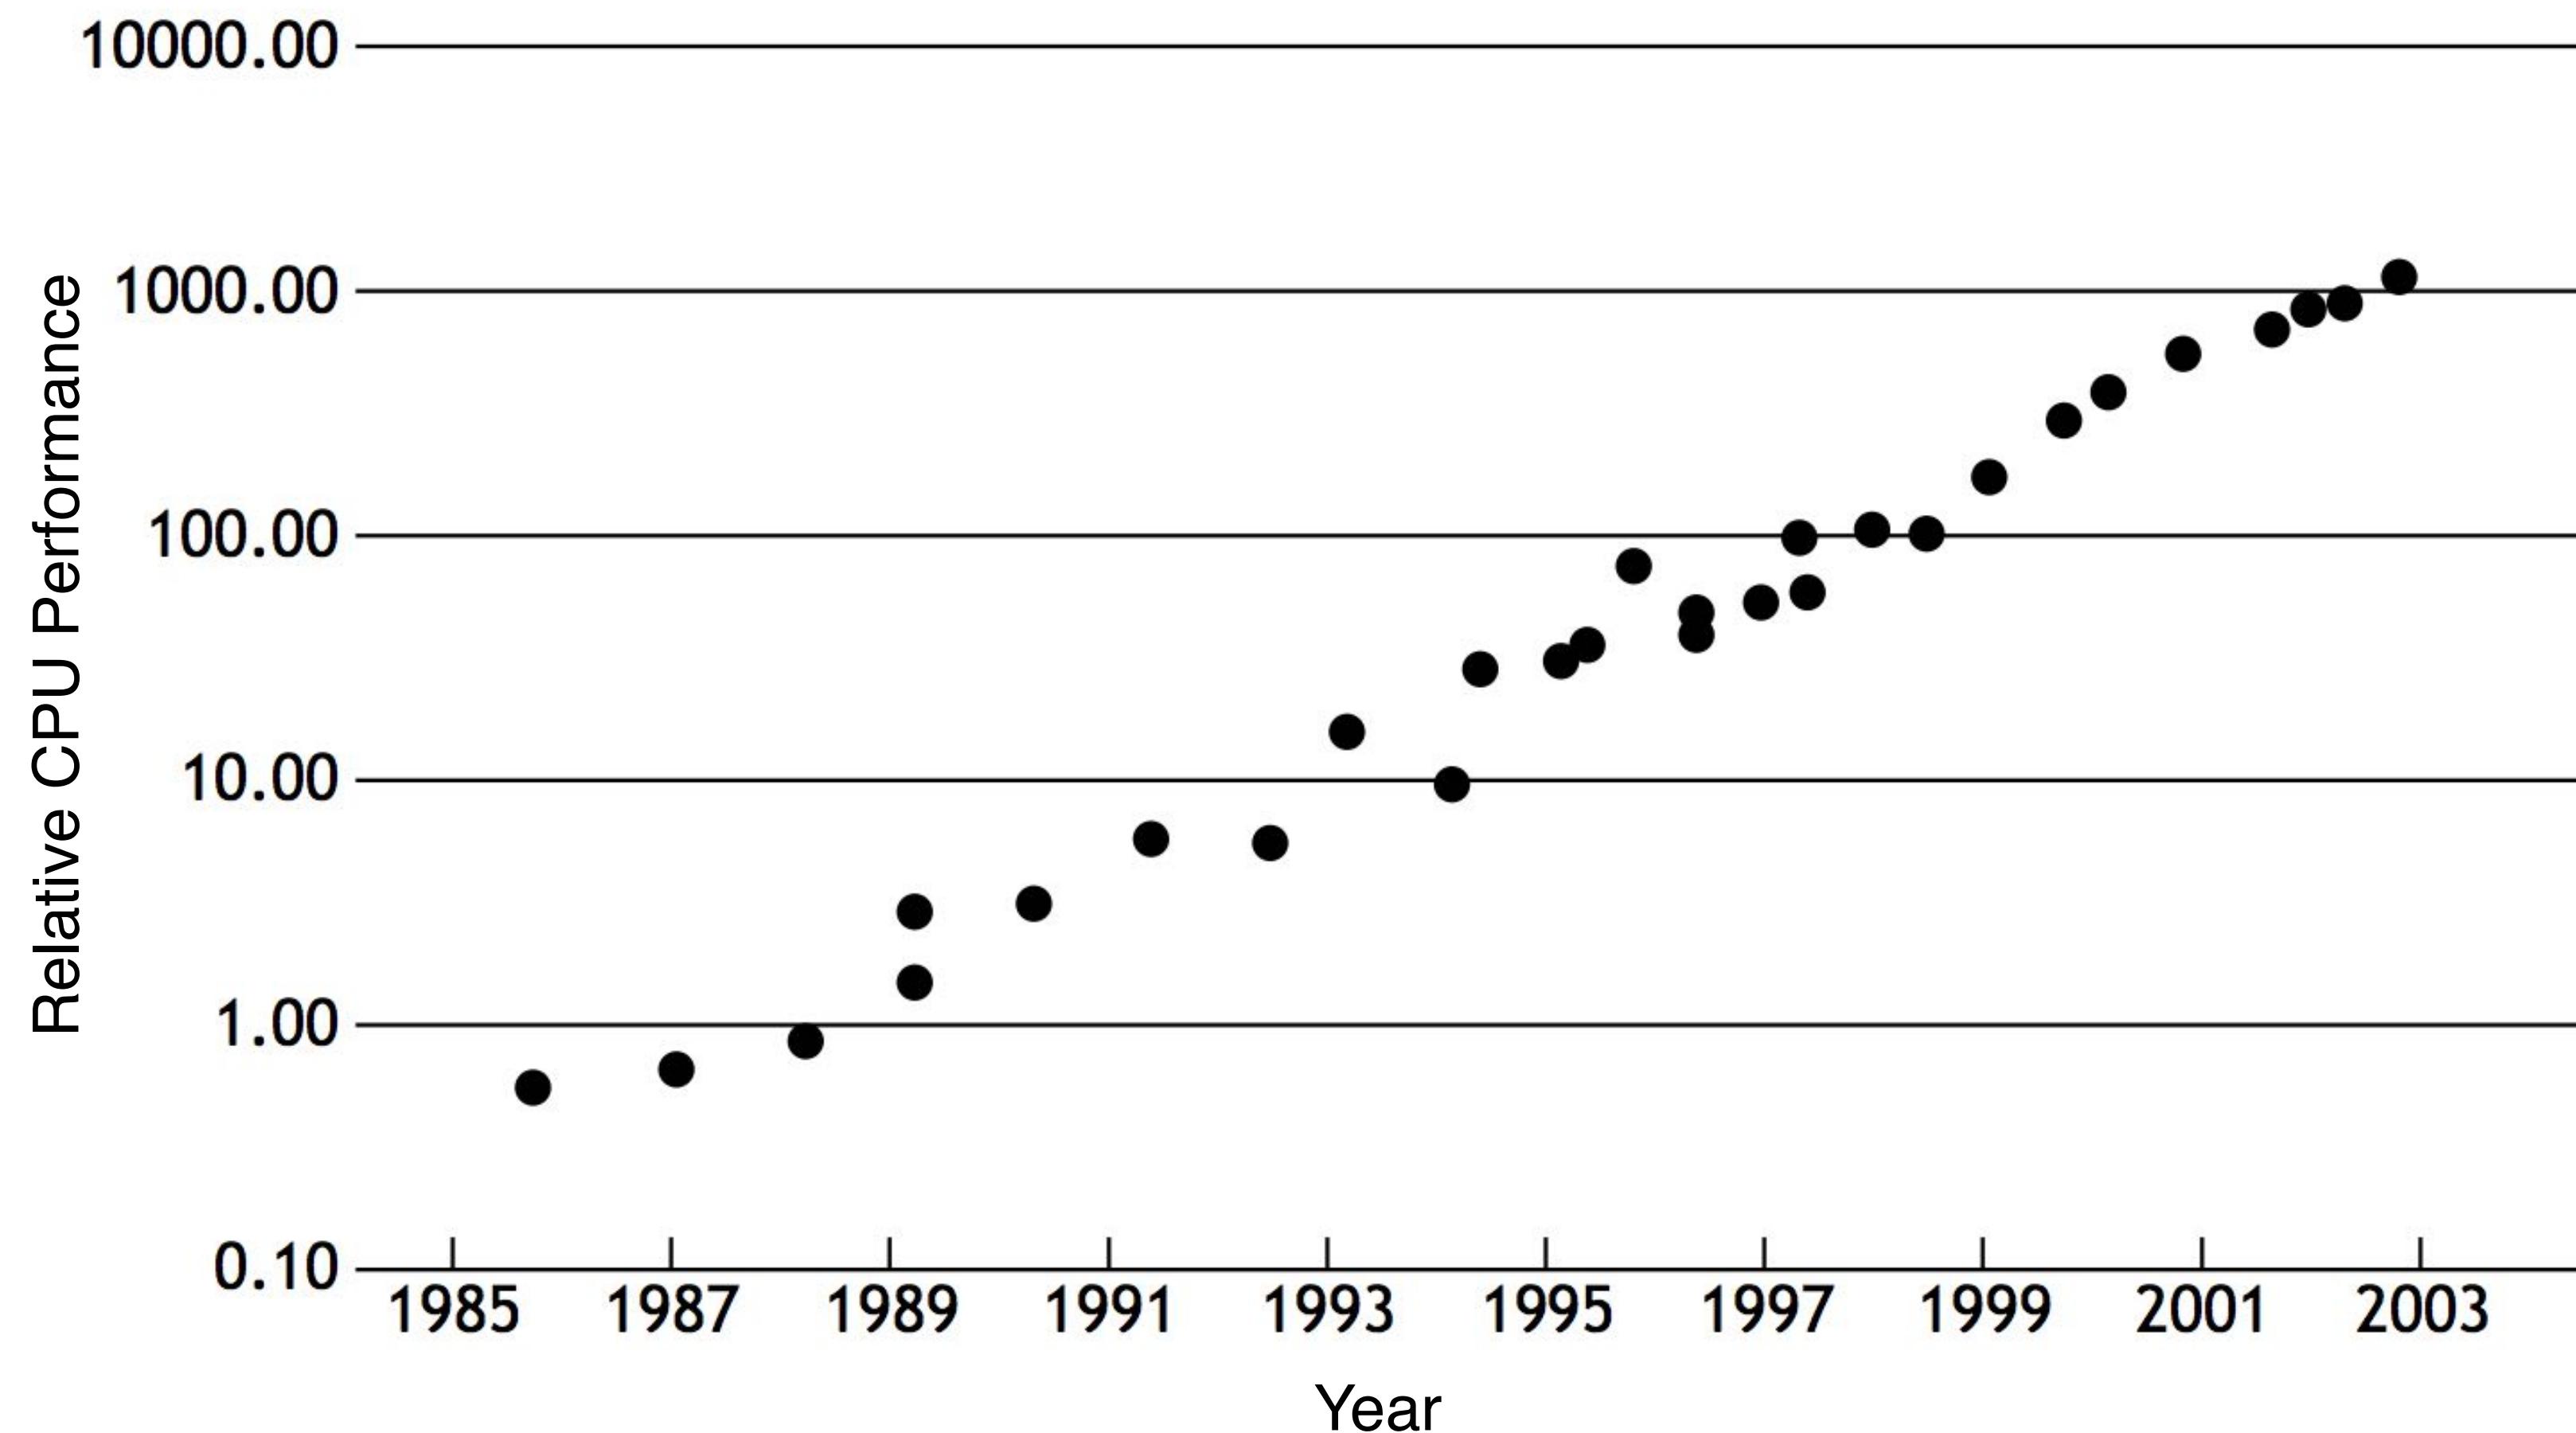
17% Final exam

- During the university-assigned slot: Dec 12th, 3:30pm

Why parallelism?

Some historical context: why avoid parallel processing?

- Single-threaded CPU performance doubling ~ every 18 months
- Implication: working to parallelize your code was often not worth the time
 - Software developer does nothing, code gets faster next year. Woot!



Until ~15 years ago: two significant reasons for processor performance improvement

- 1. Exploiting instruction-level parallelism (superscalar execution)**
- 2. Increasing CPU clock frequency**

What is a computer program?

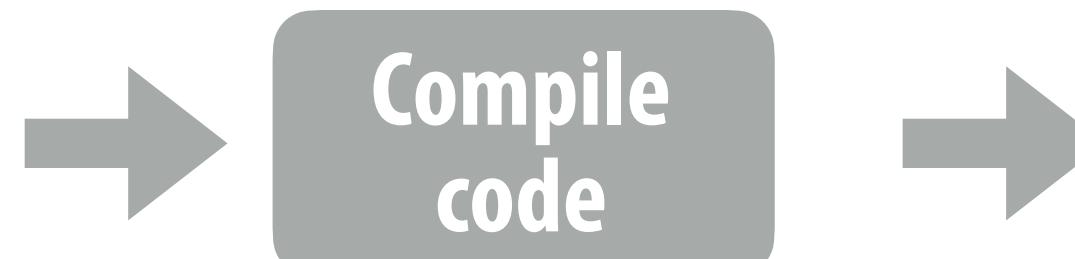
Here is a program written in C

```
int main(int argc, char** argv) {  
    int x = 1;  
  
    for (int i=0; i<10; i++) {  
        x = x + x;  
    }  
  
    printf("%d\n", x);  
  
    return 0;  
}
```

What is a program? (from a processor's perspective)

A program is just a list of processor instructions!

```
int main(int argc, char** argv) {  
  
    int x = 1;  
  
    for (int i=0; i<10; i++) {  
        x = x + x;  
    }  
  
    printf("%d\n", x);  
  
    return 0;  
}
```

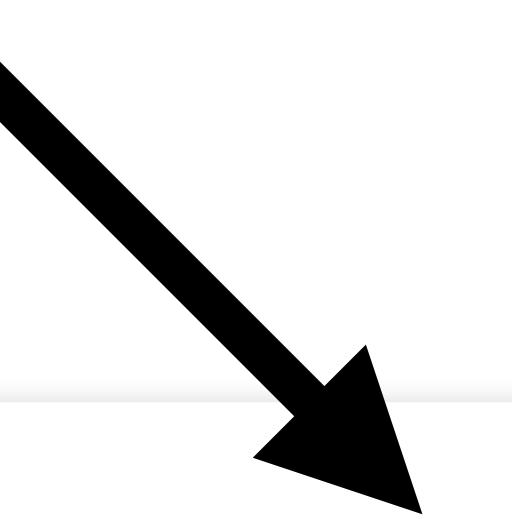


Compile
code

```
main:  
100000f10: pushq %rbp  
100000f11: movq %rsp, %rbp  
100000f14: subq $32, %rsp  
100000f18: movl $0, -4(%rbp)  
100000f1f: movl %edi, -8(%rbp)  
100000f22: movq %rsi, -16(%rbp)  
100000f26: movl $1, -20(%rbp)  
100000f2d: movl $0, -24(%rbp)  
100000f34: cmpl $10, -24(%rbp)  
100000f38: jge 23 <_main+0x45>  
100000f3e: movl -20(%rbp), %eax  
100000f41: addl -20(%rbp), %eax  
100000f44: movl %eax, -20(%rbp)  
100000f47: movl -24(%rbp), %eax  
100000f4a: addl $1, %eax  
100000f4d: movl %eax, -24(%rbp)  
100000f50: jmp -33 <_main+0x24>  
100000f55: leaq 58(%rip), %rdi  
100000f5c: movl -20(%rbp), %esi  
100000f5f: movb $0, %al  
100000f61: callq 14  
100000f66: xorl %esi, %esi  
100000f68: movl %eax, -28(%rbp)  
100000f6b: movl %esi, %eax  
100000f6d: addq $32, %rsp  
100000f71: popq %rbp  
100000f72: rets
```

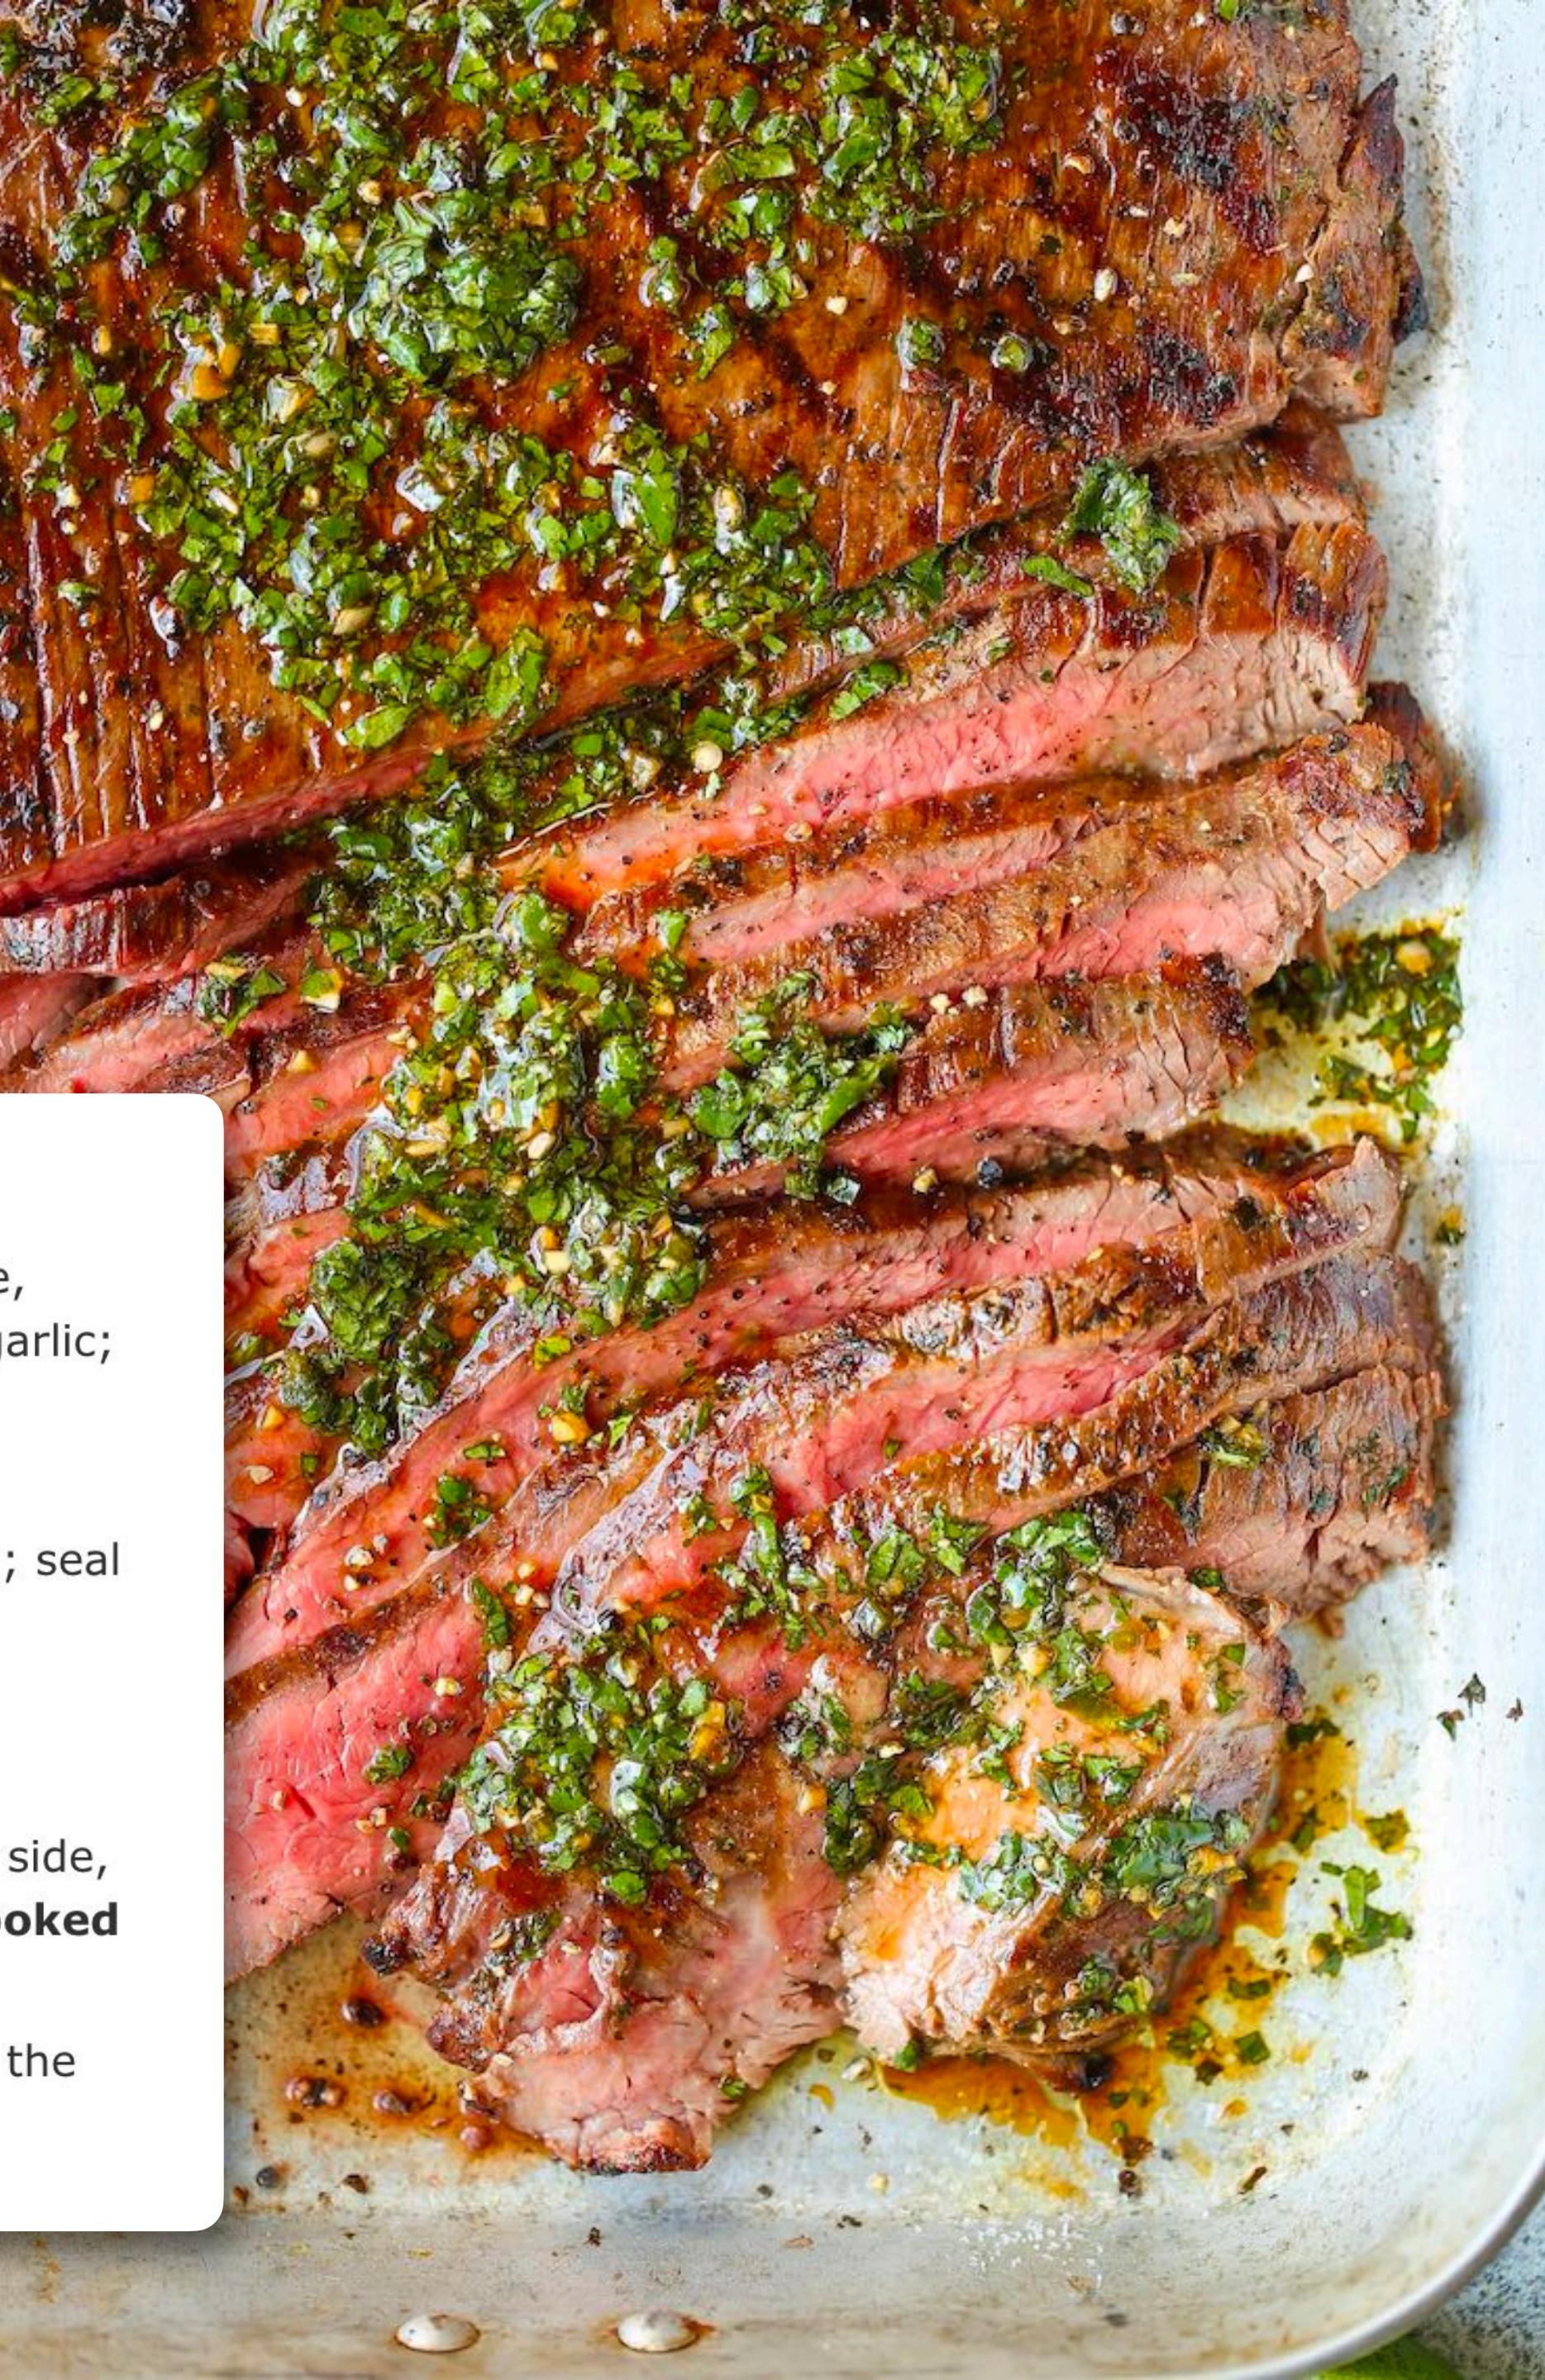
Kind of like the instructions in a recipe for your favorite meals

Mmm, carne asada



Instructions

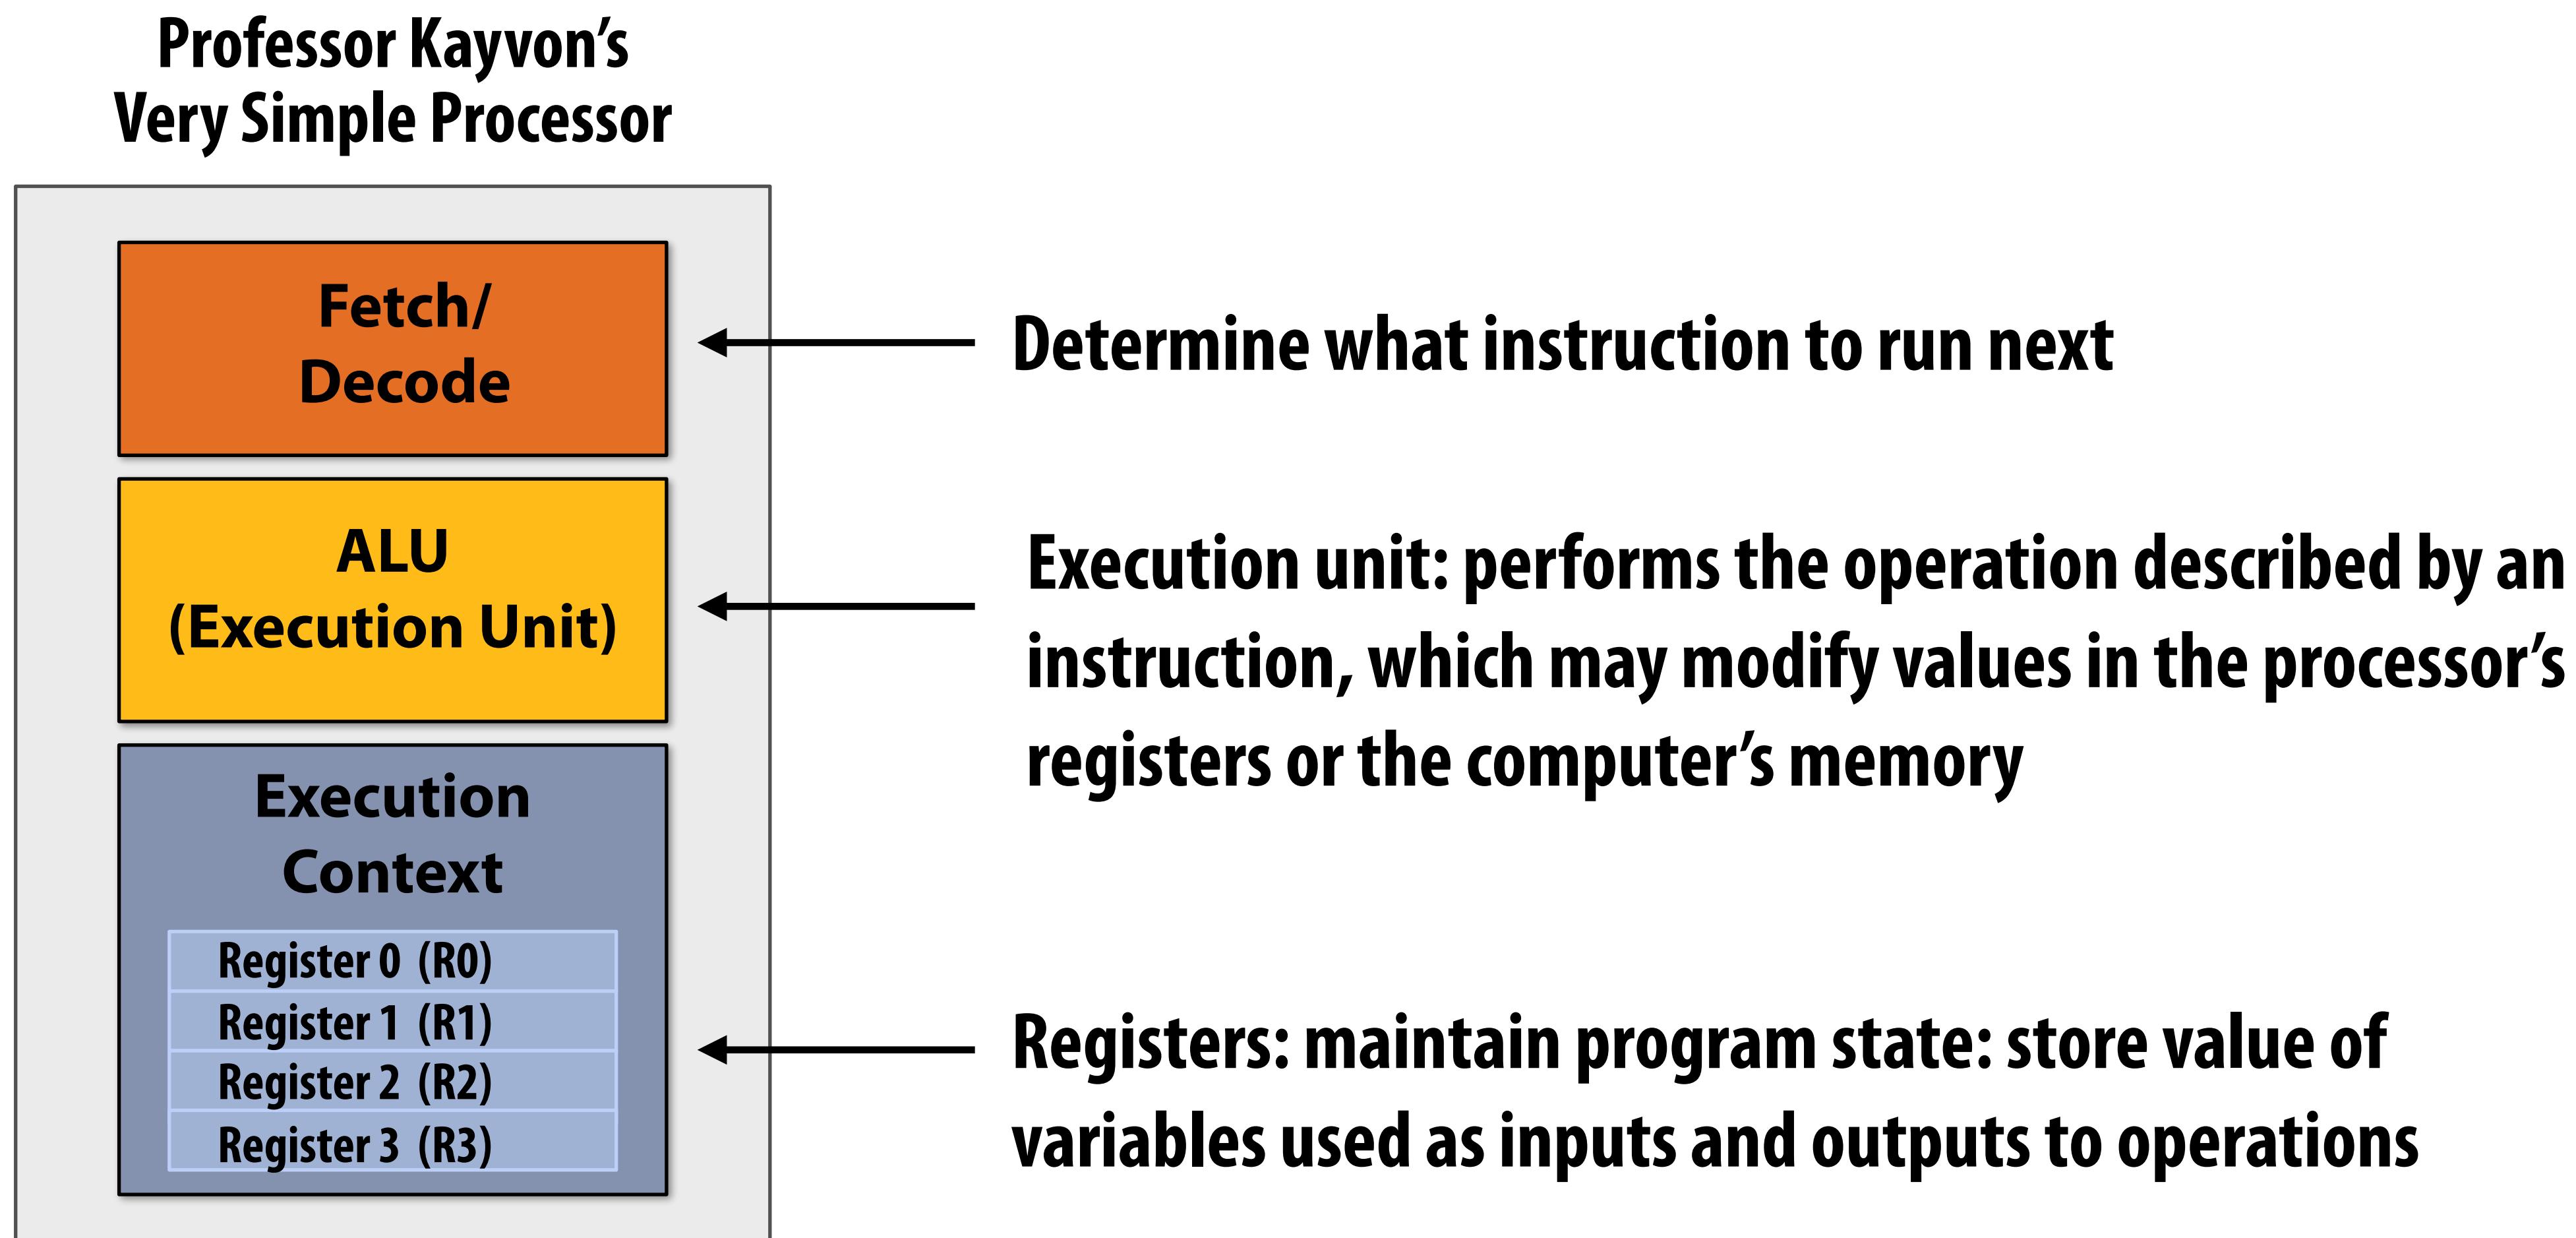
1. In a large mixing bowl combine orange juice, olive oil, cilantro, lime juice, lemon juice, white wine vinegar, cumin, salt and pepper, jalapeno, and garlic; whisk until well combined.
2. Reserve $\frac{1}{3}$ cup of the marinade; cover the rest and refrigerate.
3. Combine remaining marinade and steak in a large resealable freezer bag; seal and refrigerate for at least 2 hours, or overnight.
4. Preheat grill to HIGH heat.
5. Remove steak from marinade and lightly pat dry with paper towels.
6. Add steak to the preheated grill and cook for another 6 to 8 minutes per side, or until desired doneness. **Note that flank steak tastes best when cooked to rare or medium rare because it's a lean cut of steak.**
7. Remove from heat and let rest for 10 minutes. Thinly slice steak against the grain, garnish with reserved cilantro mixture, and serve.



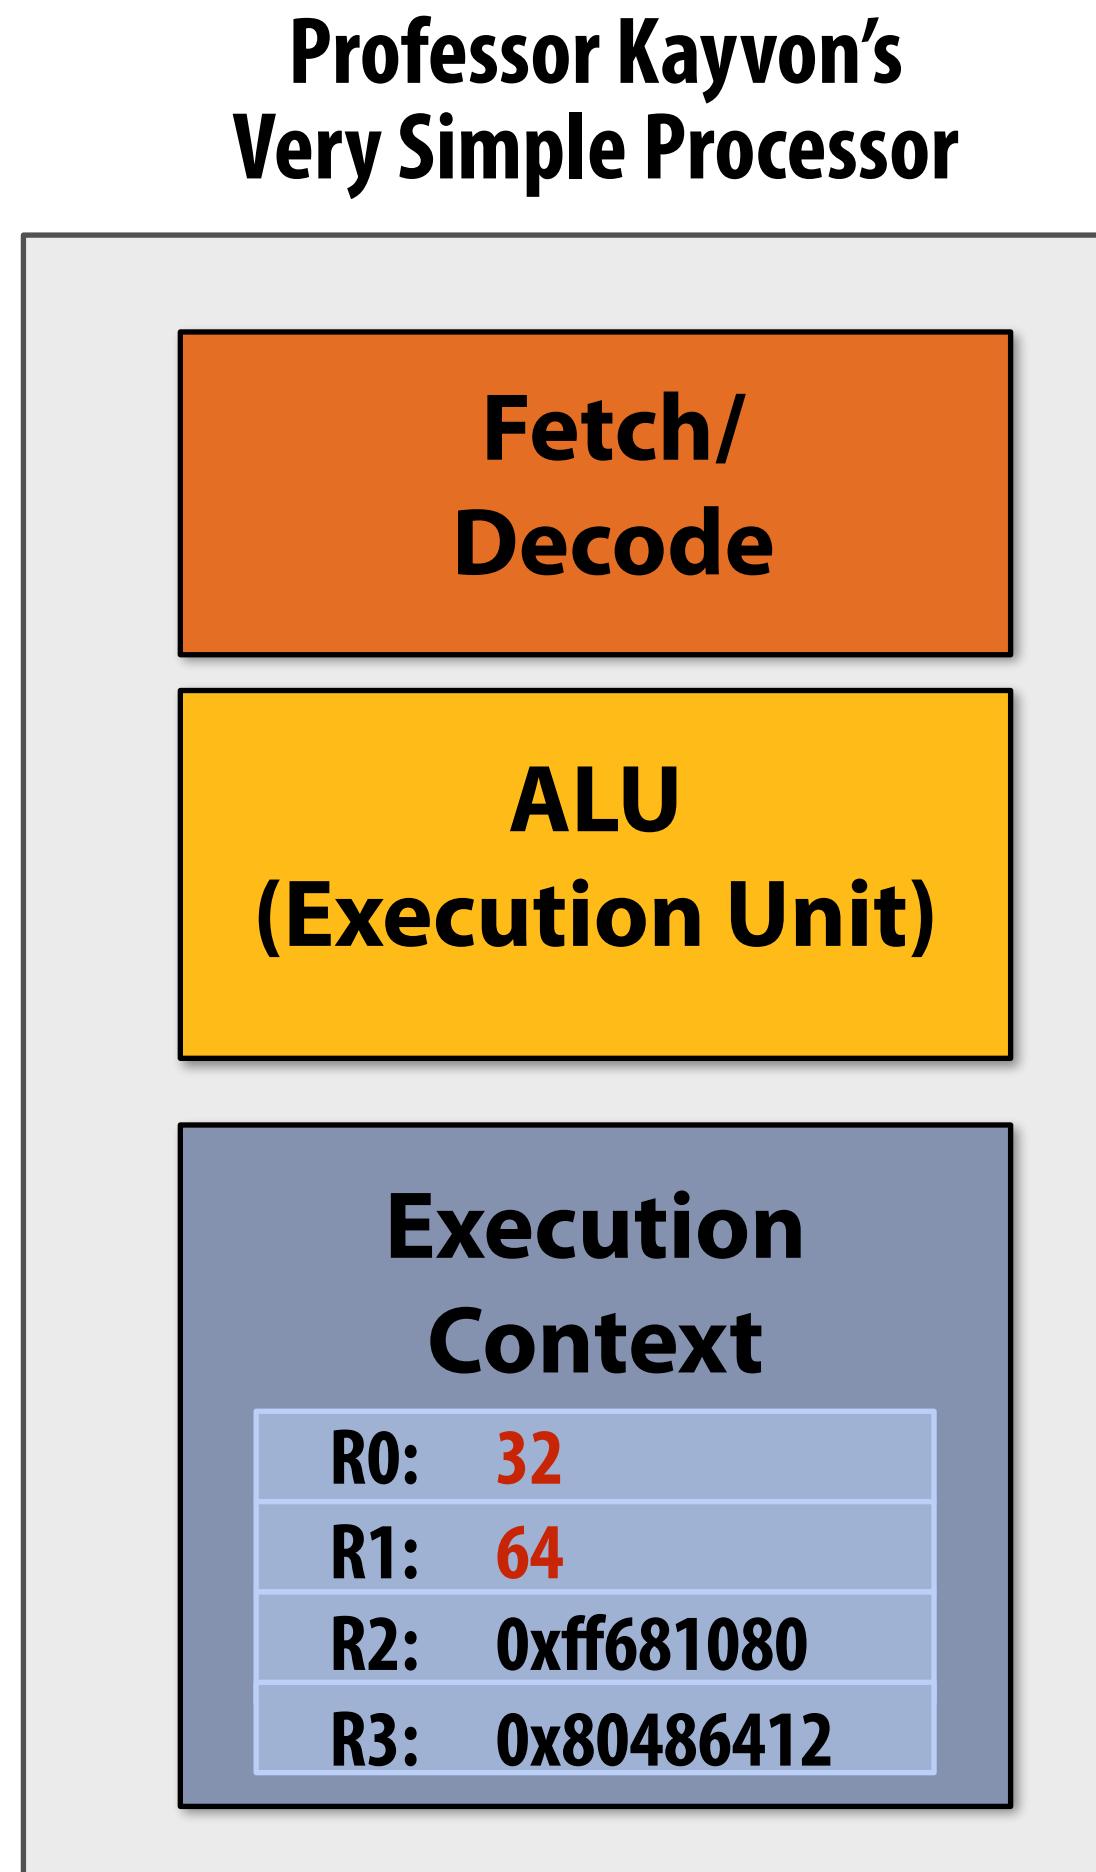
What does a processor do?



A processor executes instructions



One example instruction: add two numbers



Step 1:

Processor gets next program instruction from memory
(figure out what the processor should do next)

add R0 \leftarrow R0, R1

"Please add the contents of register R0 to the contents of register R1 and put the result of the addition into register R0"

Step 2:

Get operation inputs from registers

Contents of R0 input to execution unit: **32**

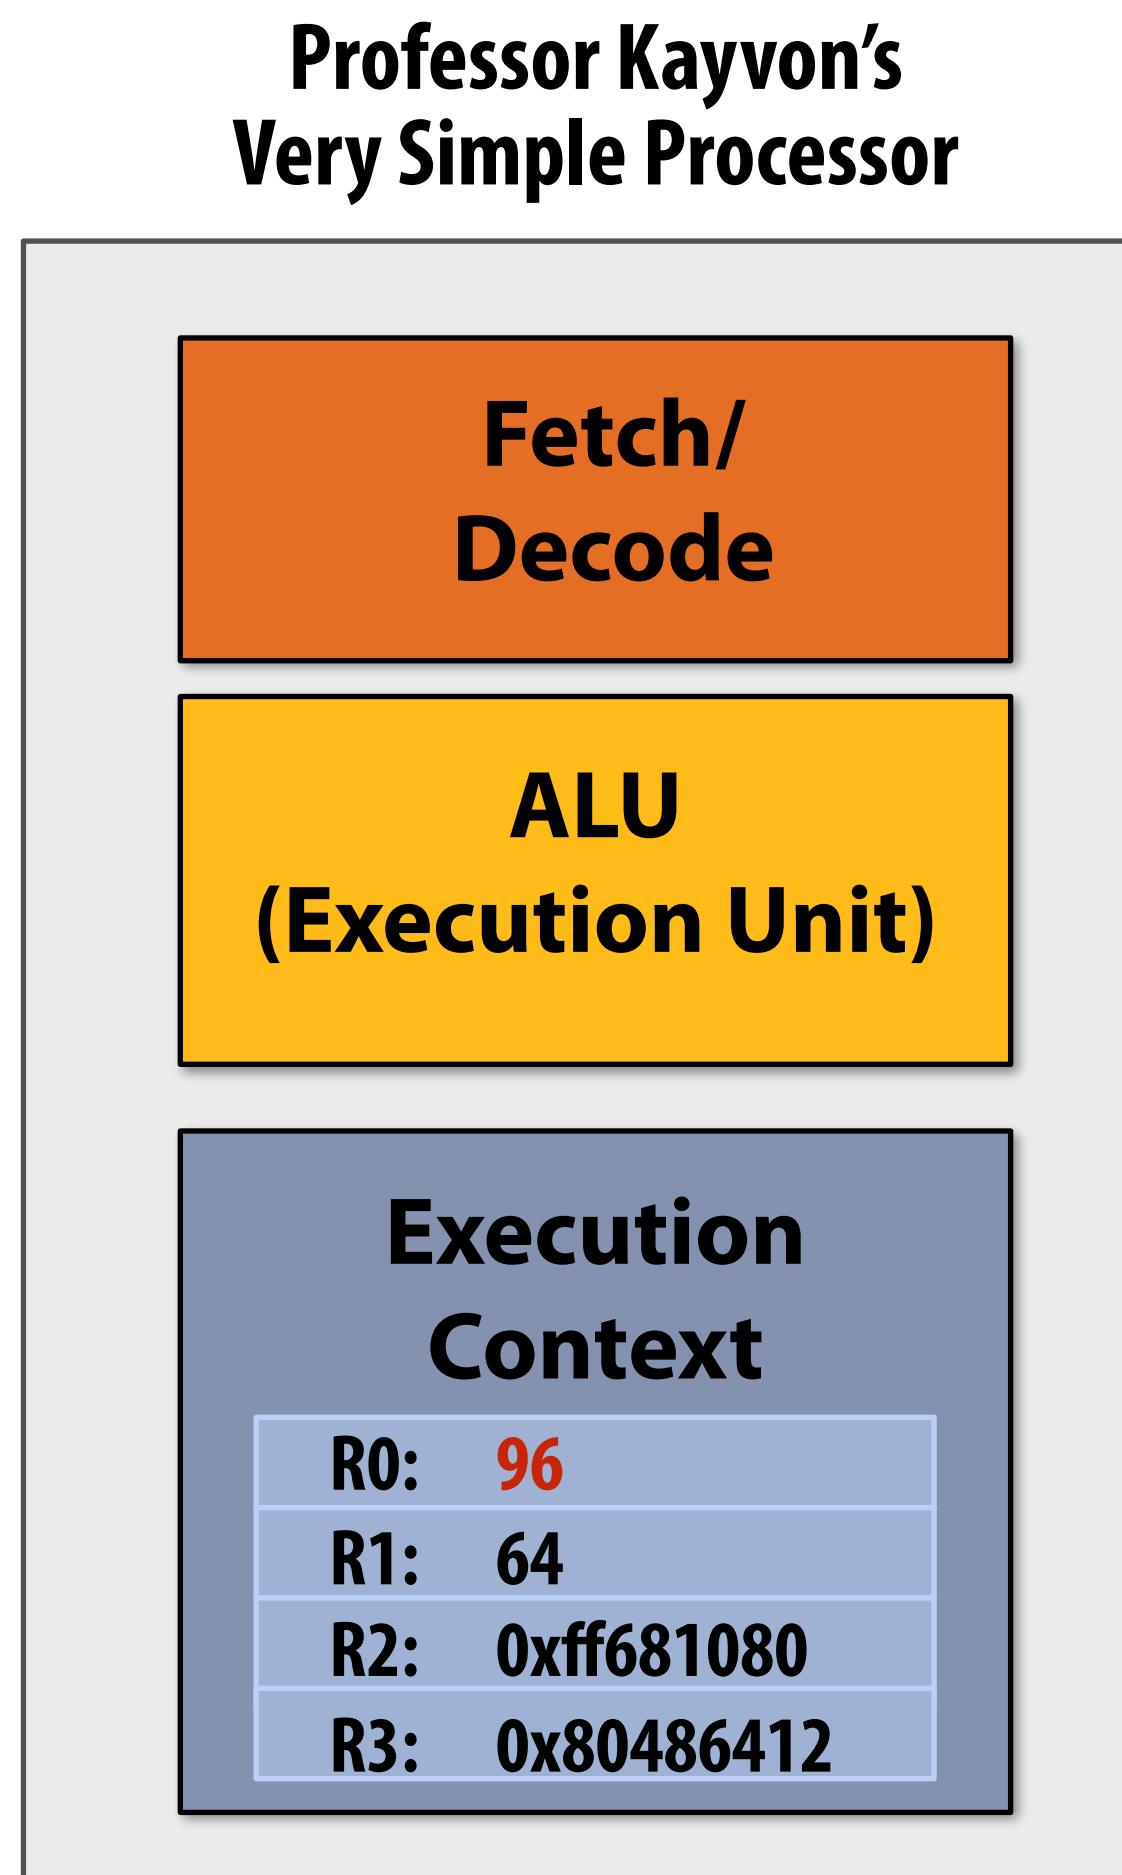
Contents of R1 input to execution unit: **64**

Step 3:

Perform addition operation:

Execution unit performs arithmetic, the result is: **96**

One example instruction: add two numbers



Step 1:
Processor gets next program instruction from memory
(figure out what the processor should do next)

add R0 \leftarrow R0, R1

"Please add the contents of register R0 to the contents of register R1 and put the result of the addition into register R0"

Step 2:
Get operation inputs from registers

Contents of R0 input to execution unit: **32**

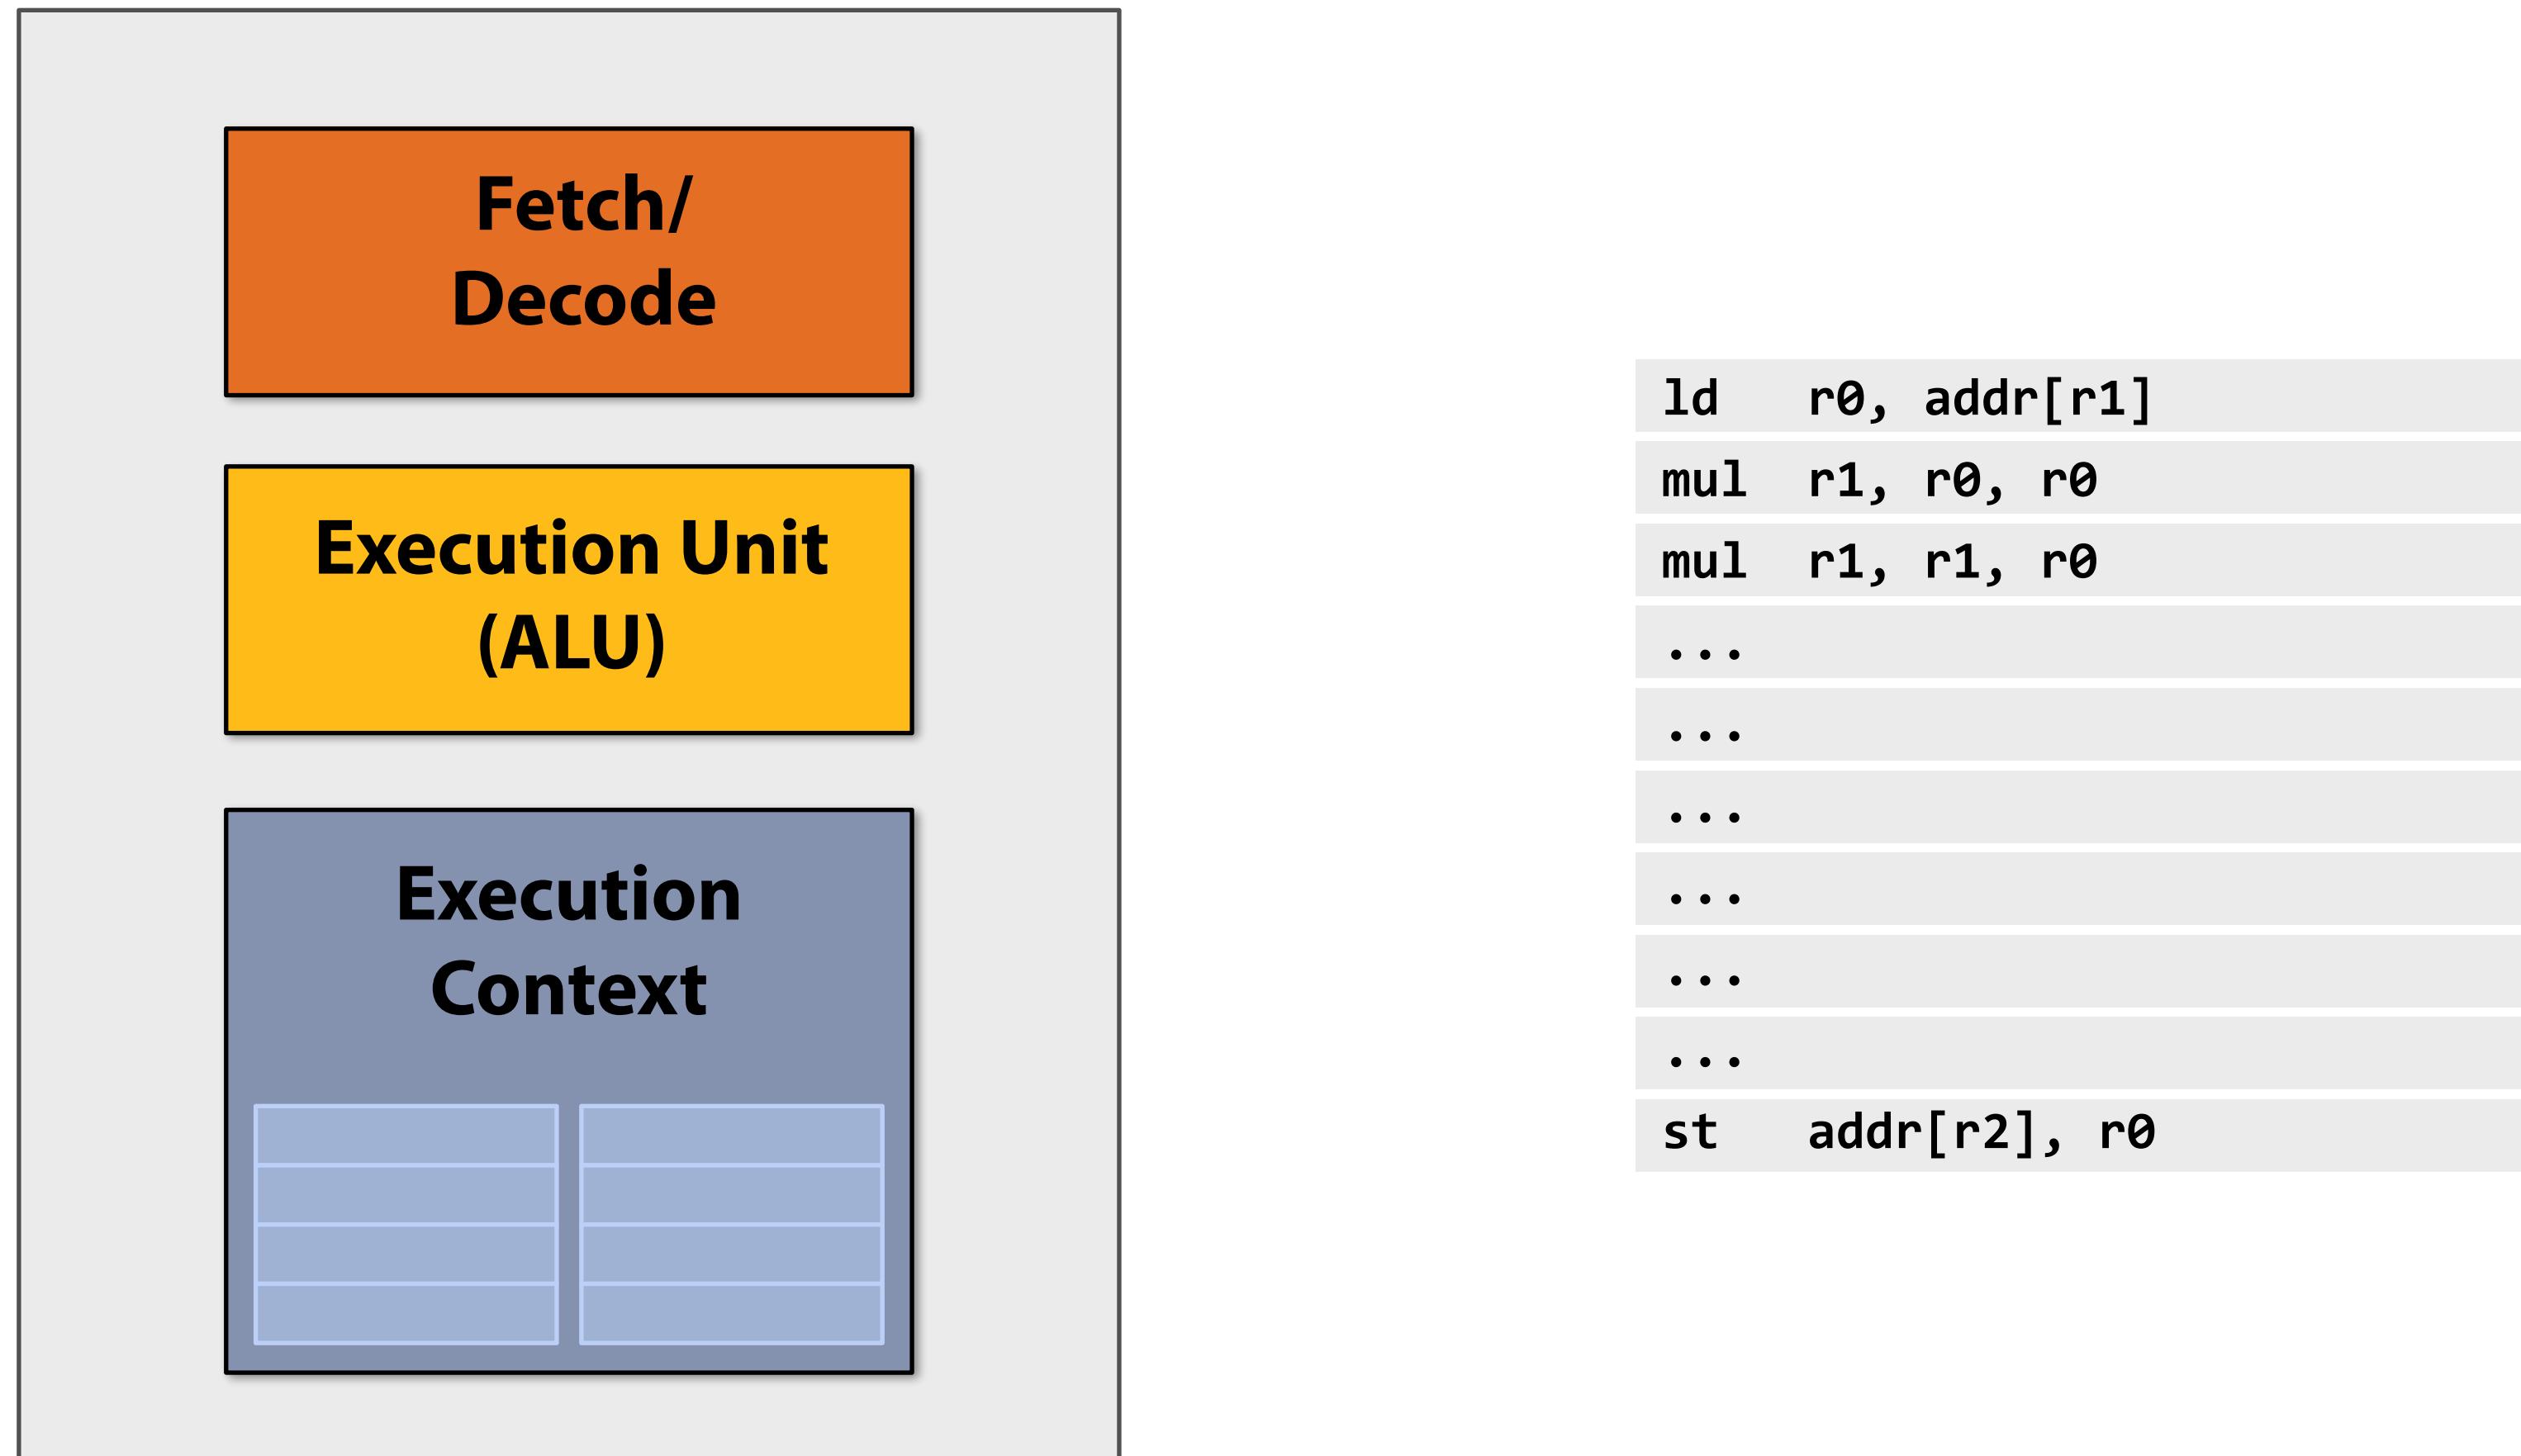
Contents of R1 input to execution unit: **64**

Step 3:
Perform addition operation:
Execution unit performs arithmetic, the result is: **96**

Step 4:
Store result **96** back to register R0

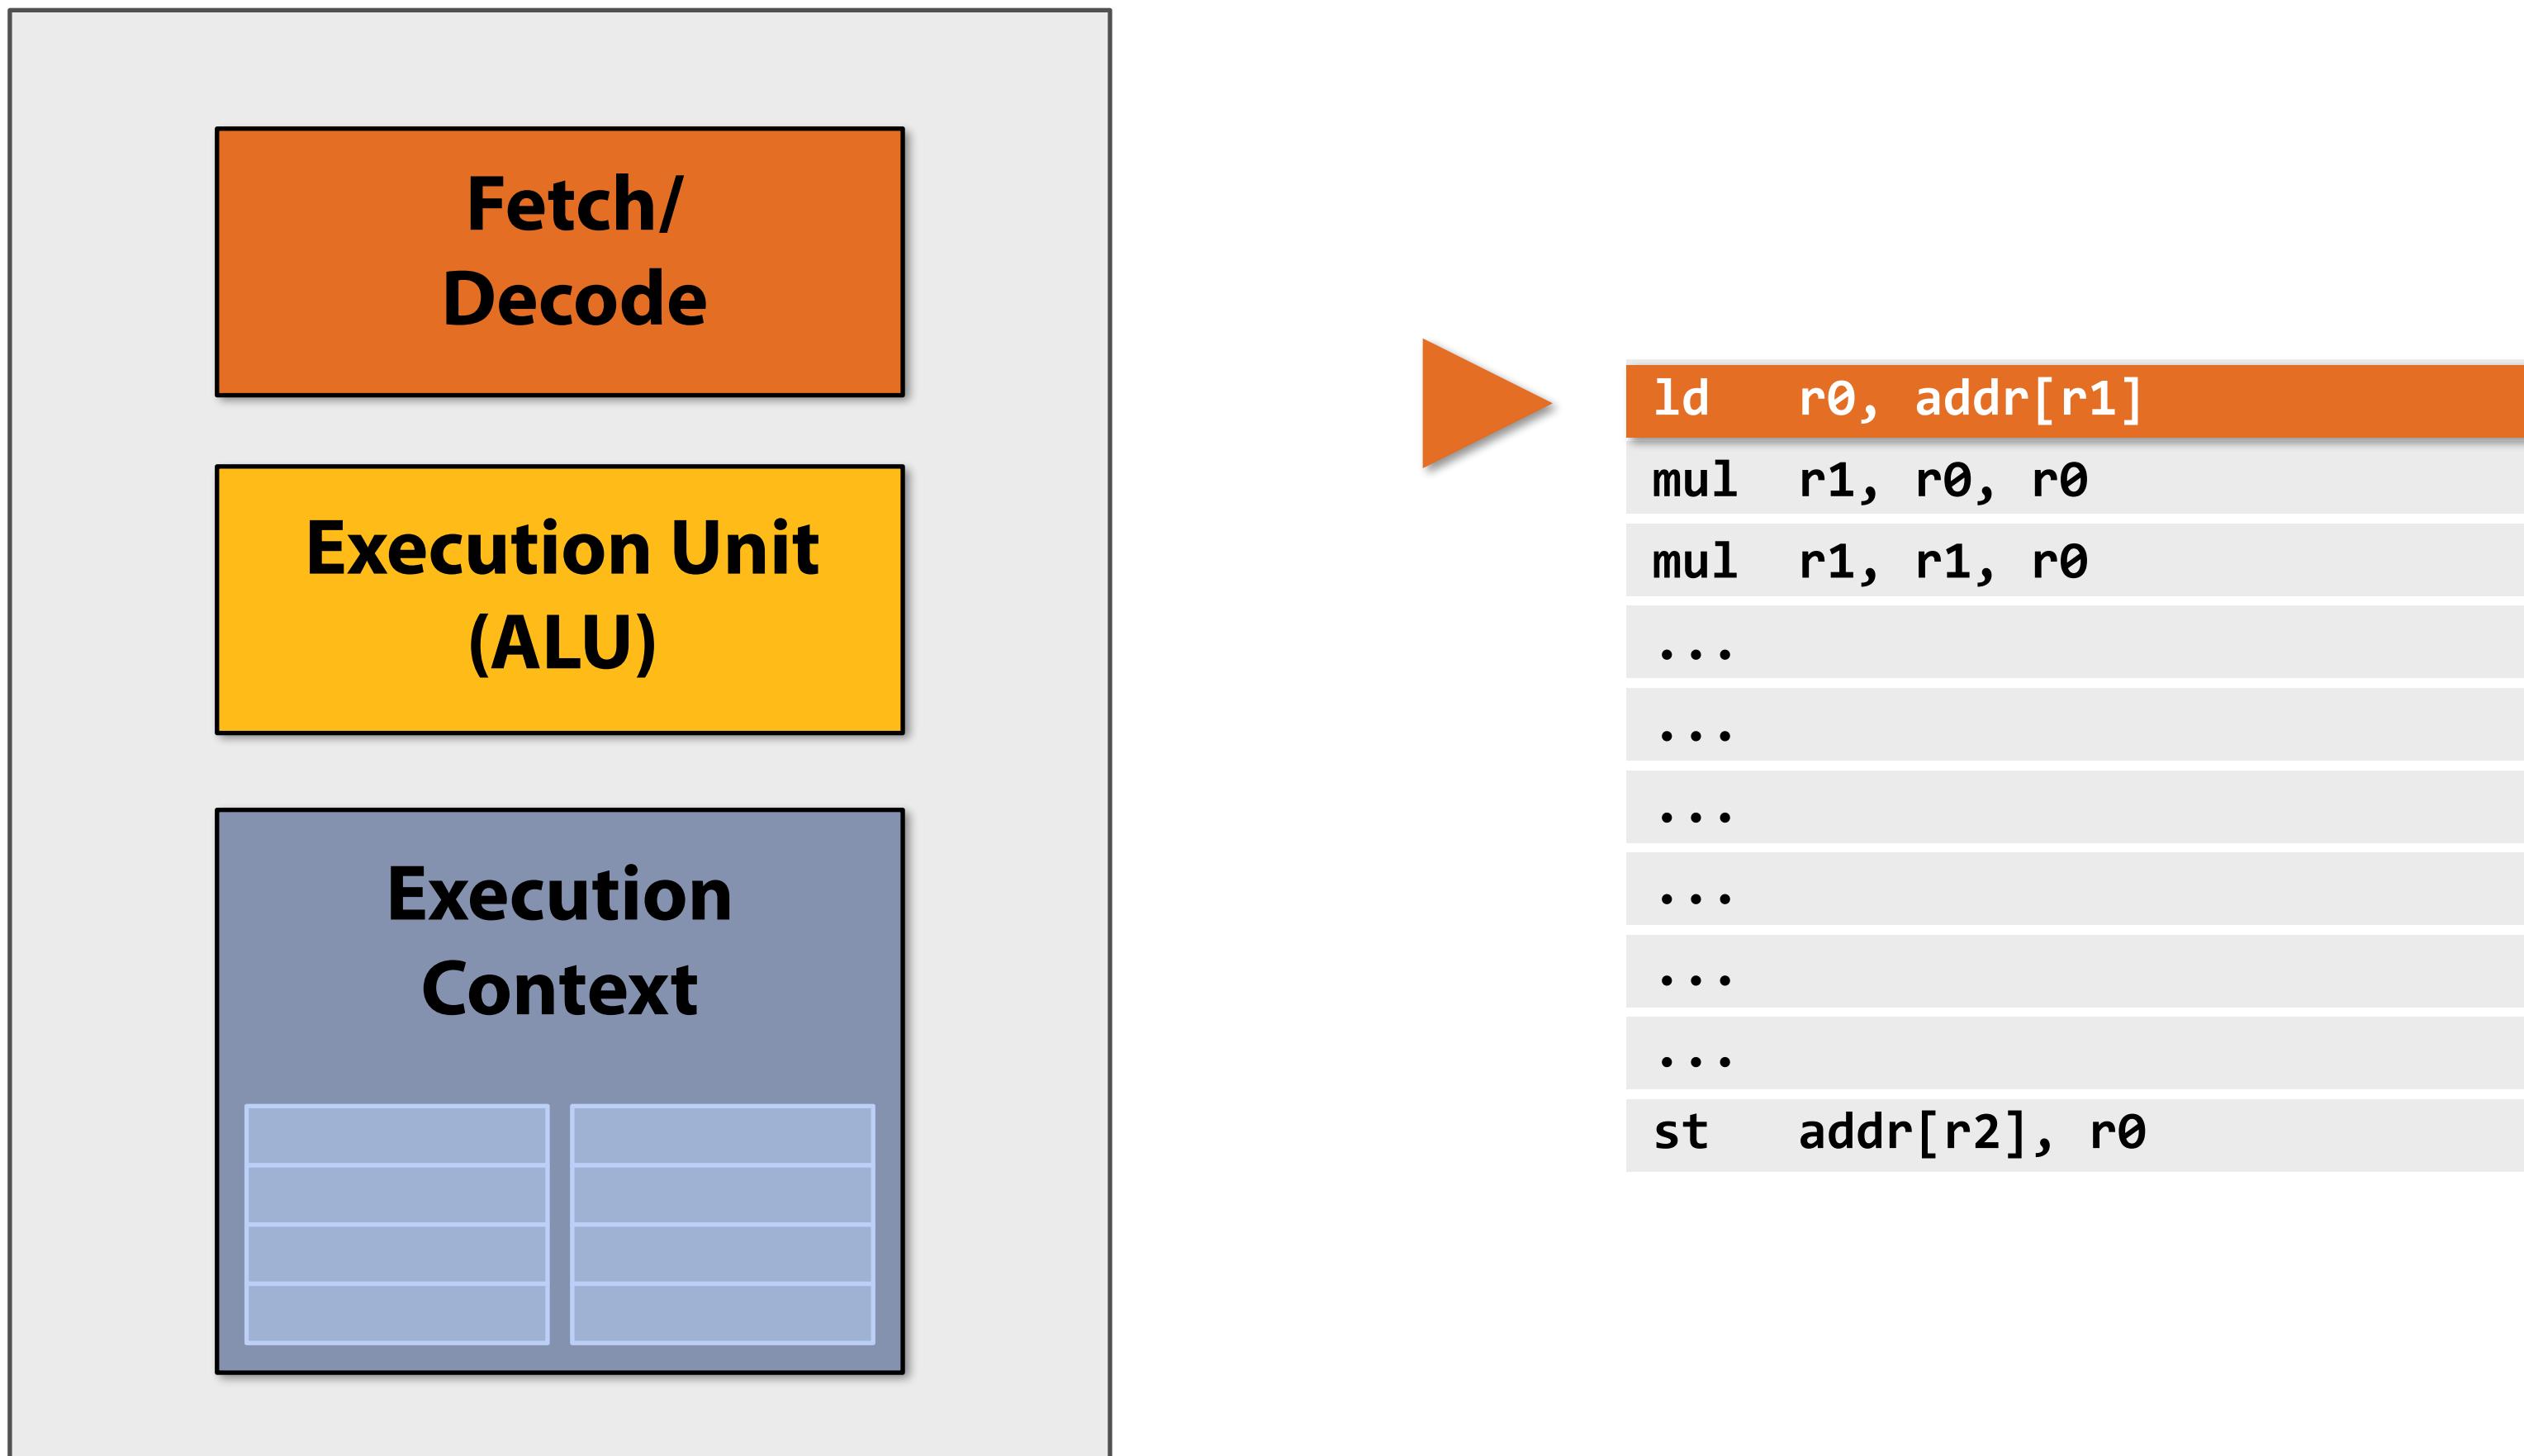
Execute program

My very simple processor: executes one instruction per clock



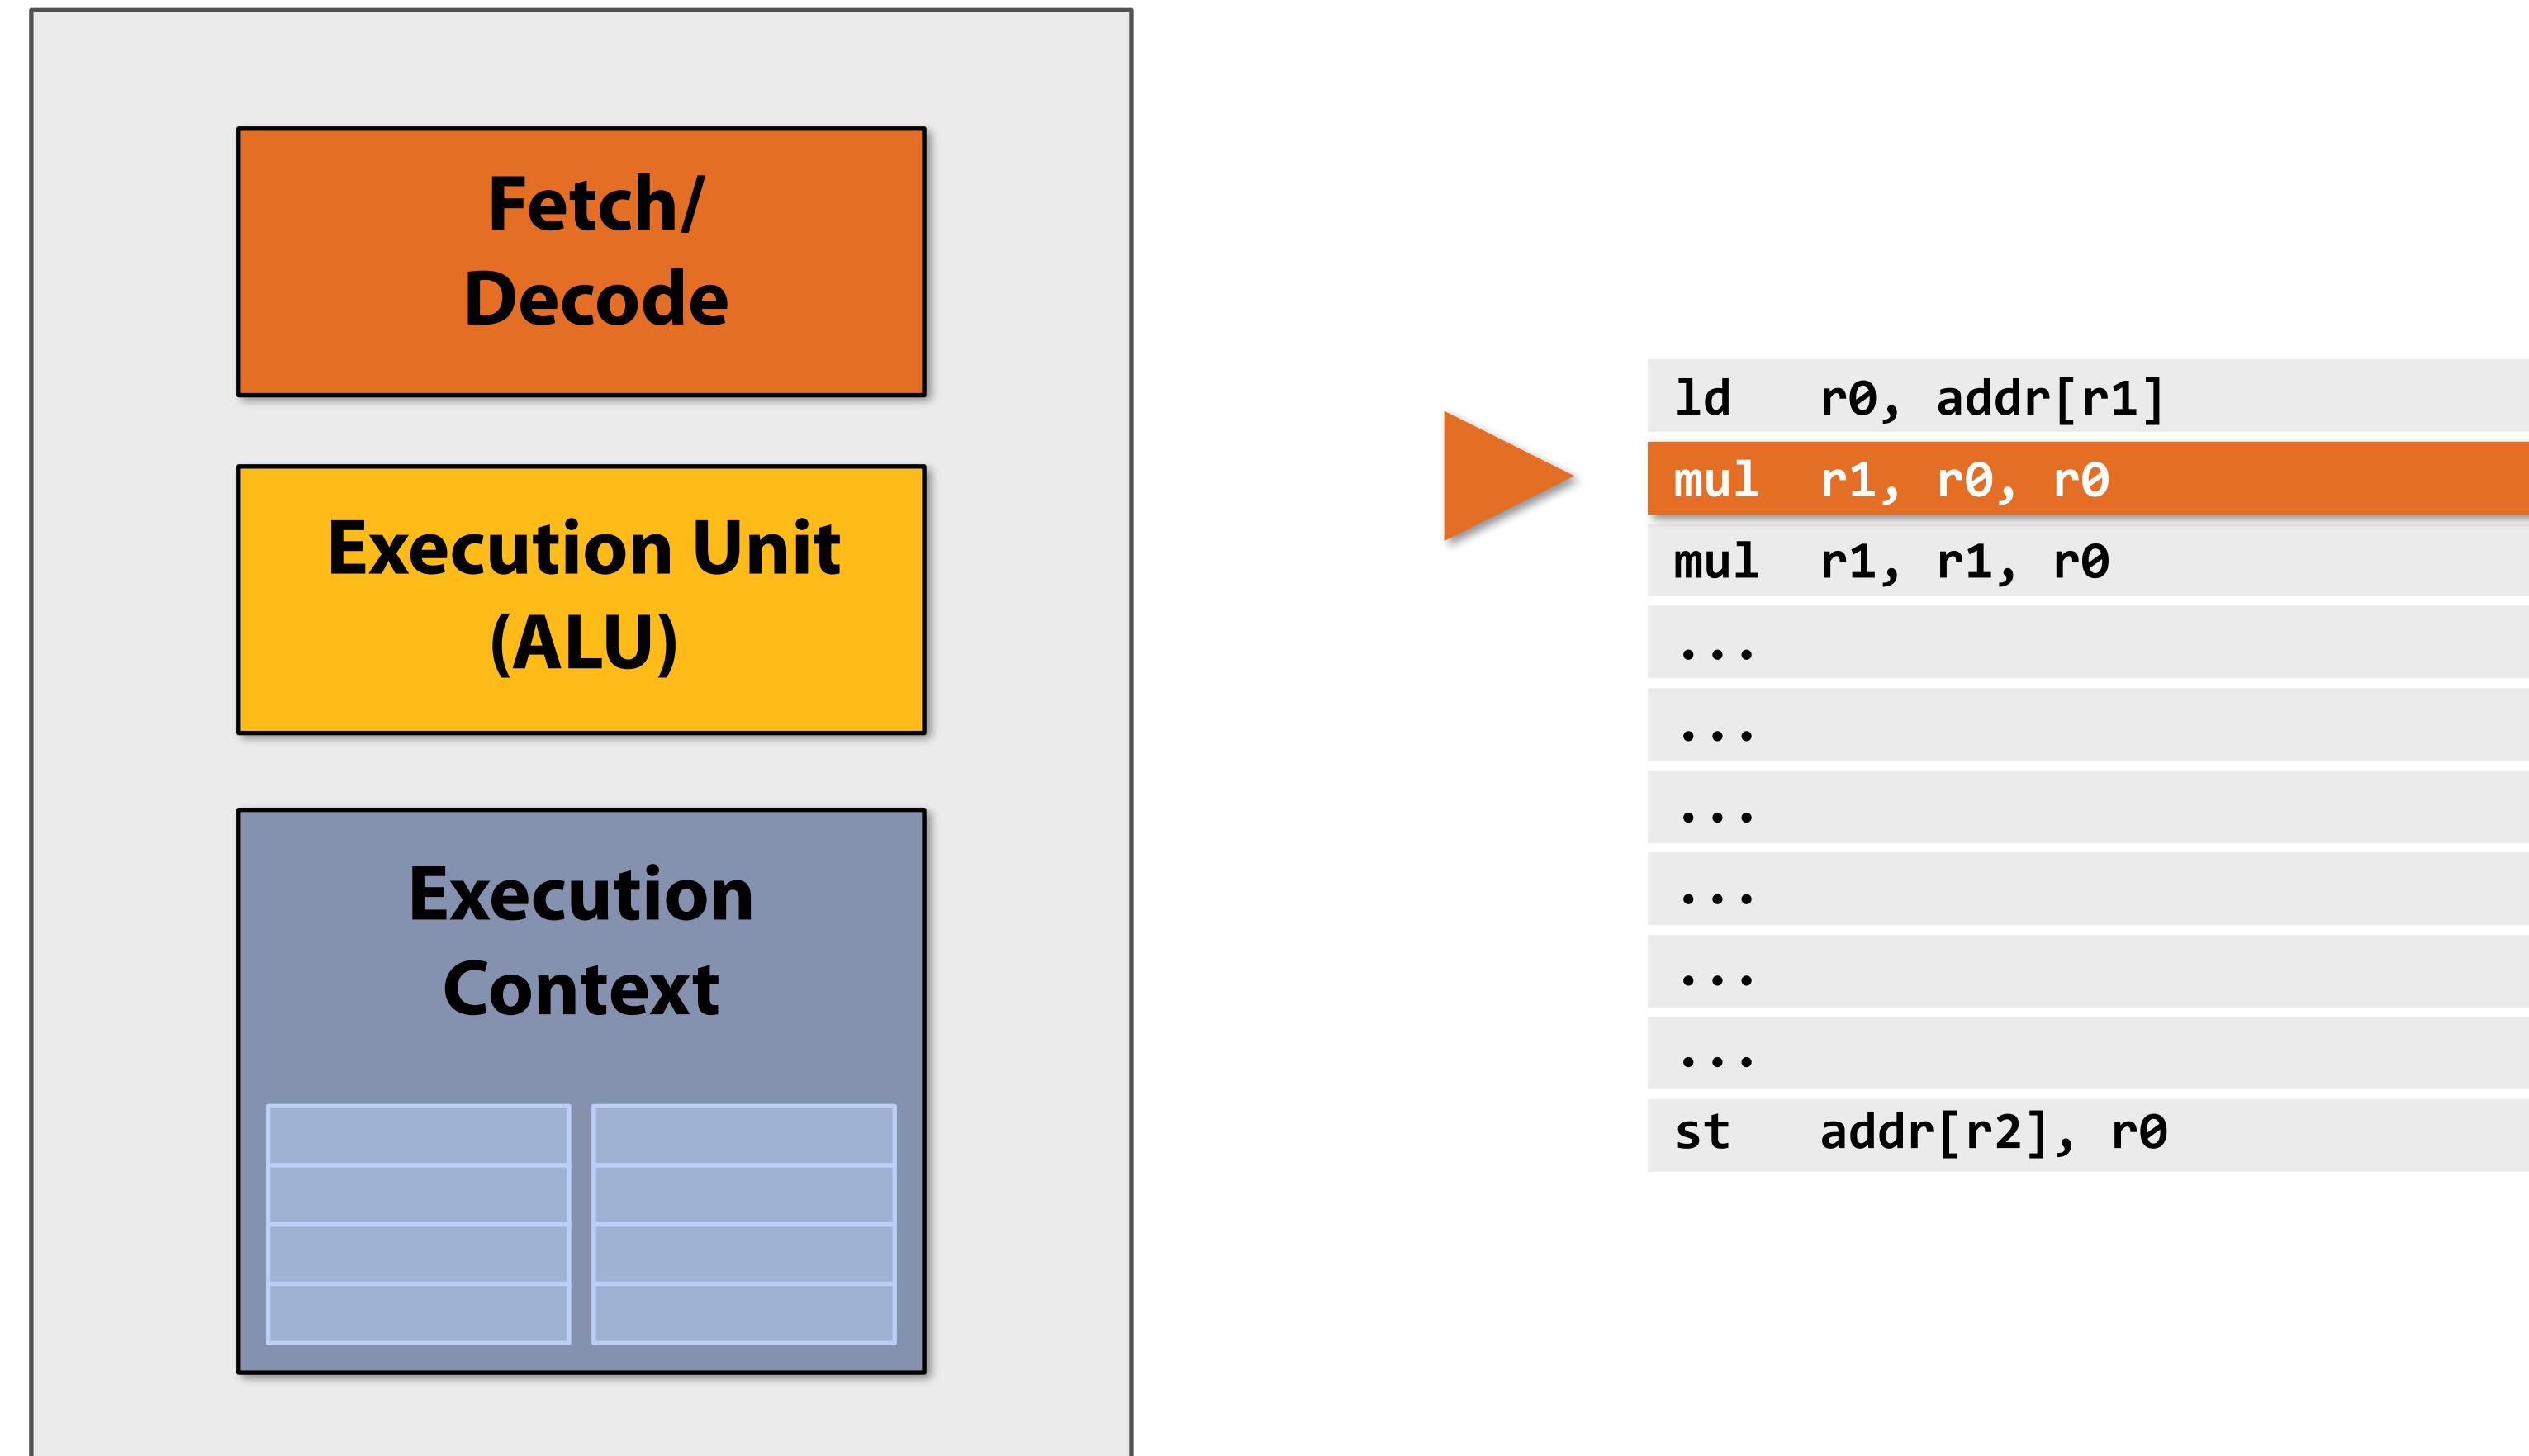
Execute program

My very simple processor: executes one instruction per clock



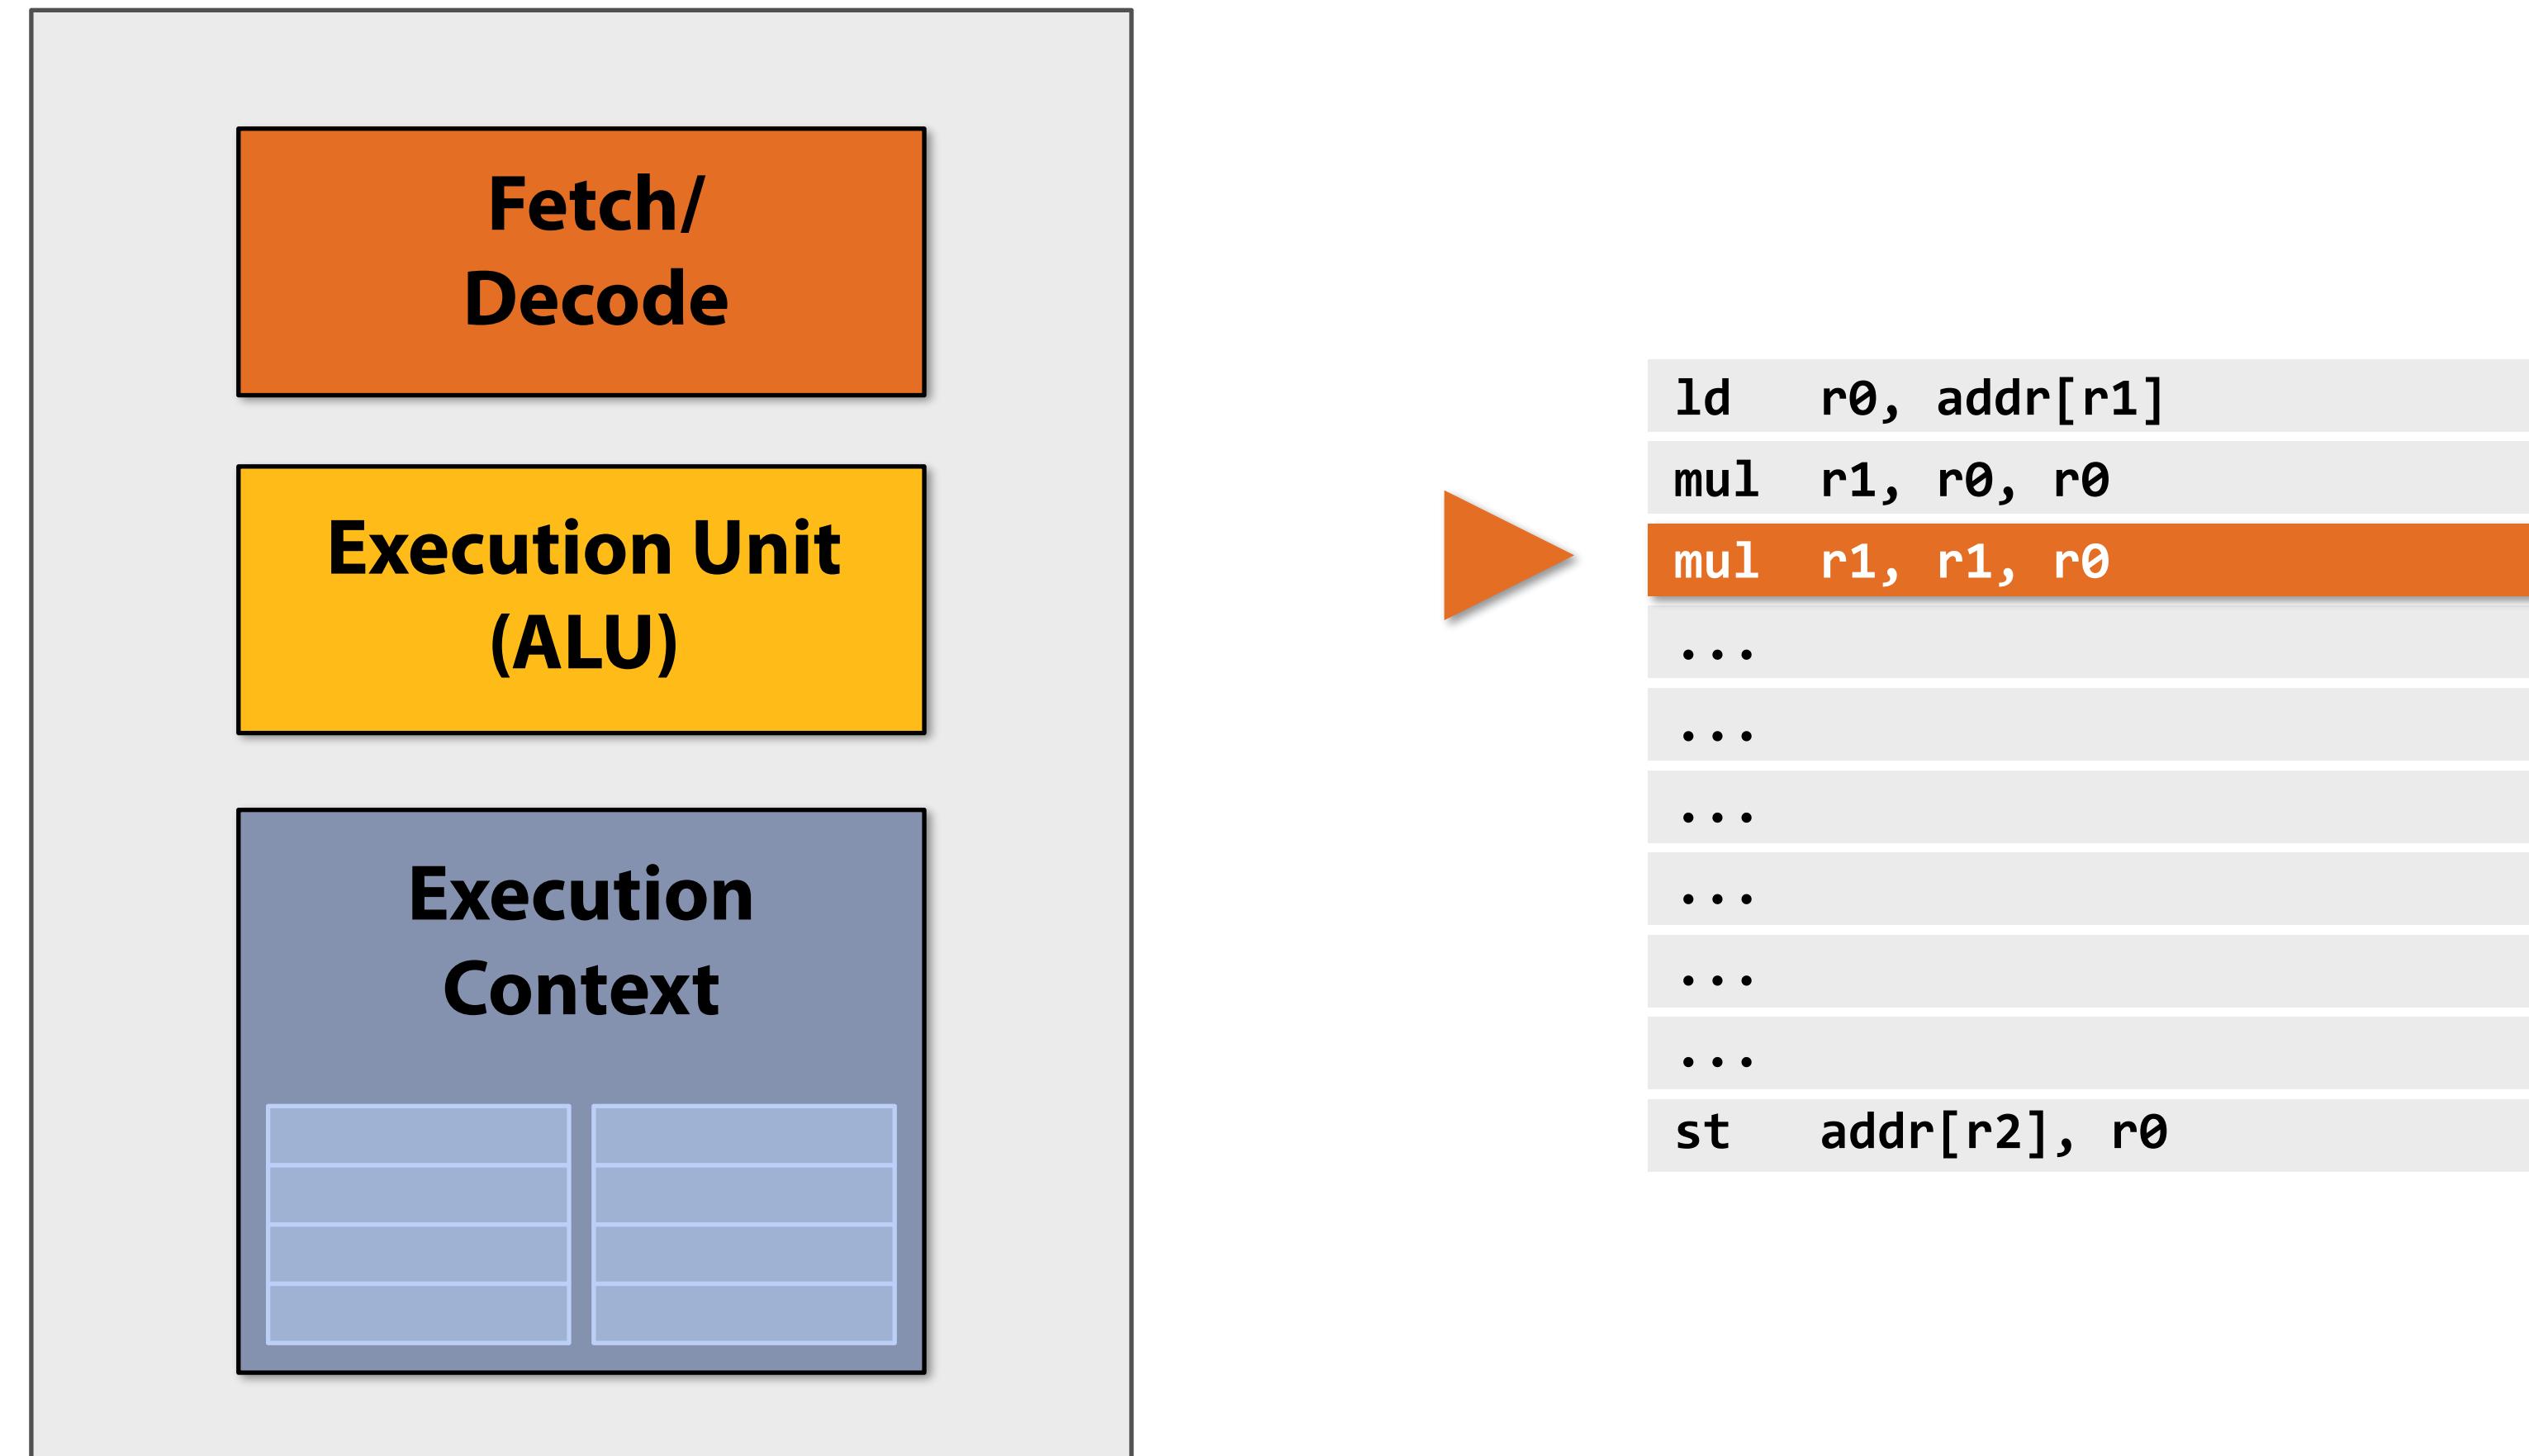
Execute program

My very simple processor: executes one instruction per clock



Execute program

My very simple processor: executes one instruction per clock



Review of how computers work...

What is a computer program? (from a processor's perspective)

It is a list of instructions to execute!

What is an instruction?

It describes an operation for a processor to perform.

Executing an instruction typically modifies the computer's state.

What do I mean when I talk about a computer's "state"?

The values of program data, which are stored in a processor's registers or in memory.

Lets consider a very simple piece of code

$$a = x*x + y*y + z*z$$

Consider the following five instruction program:

Assume register $R0 = x, R1 = y, R2 = z$

```
1 mul R0, R0, R0
2 mul R1, R1, R1
3 mul R2, R2, R2
4 add R0, R0, R1
5 add R3, R0, R2
```

$R3$ now stores value of program variable ‘ a ’

This program has five instructions, so it will take five clocks to execute, correct?

Can we do better?

What if up to two instructions can be performed at once?

$$a = x*x + y*y + z*z$$

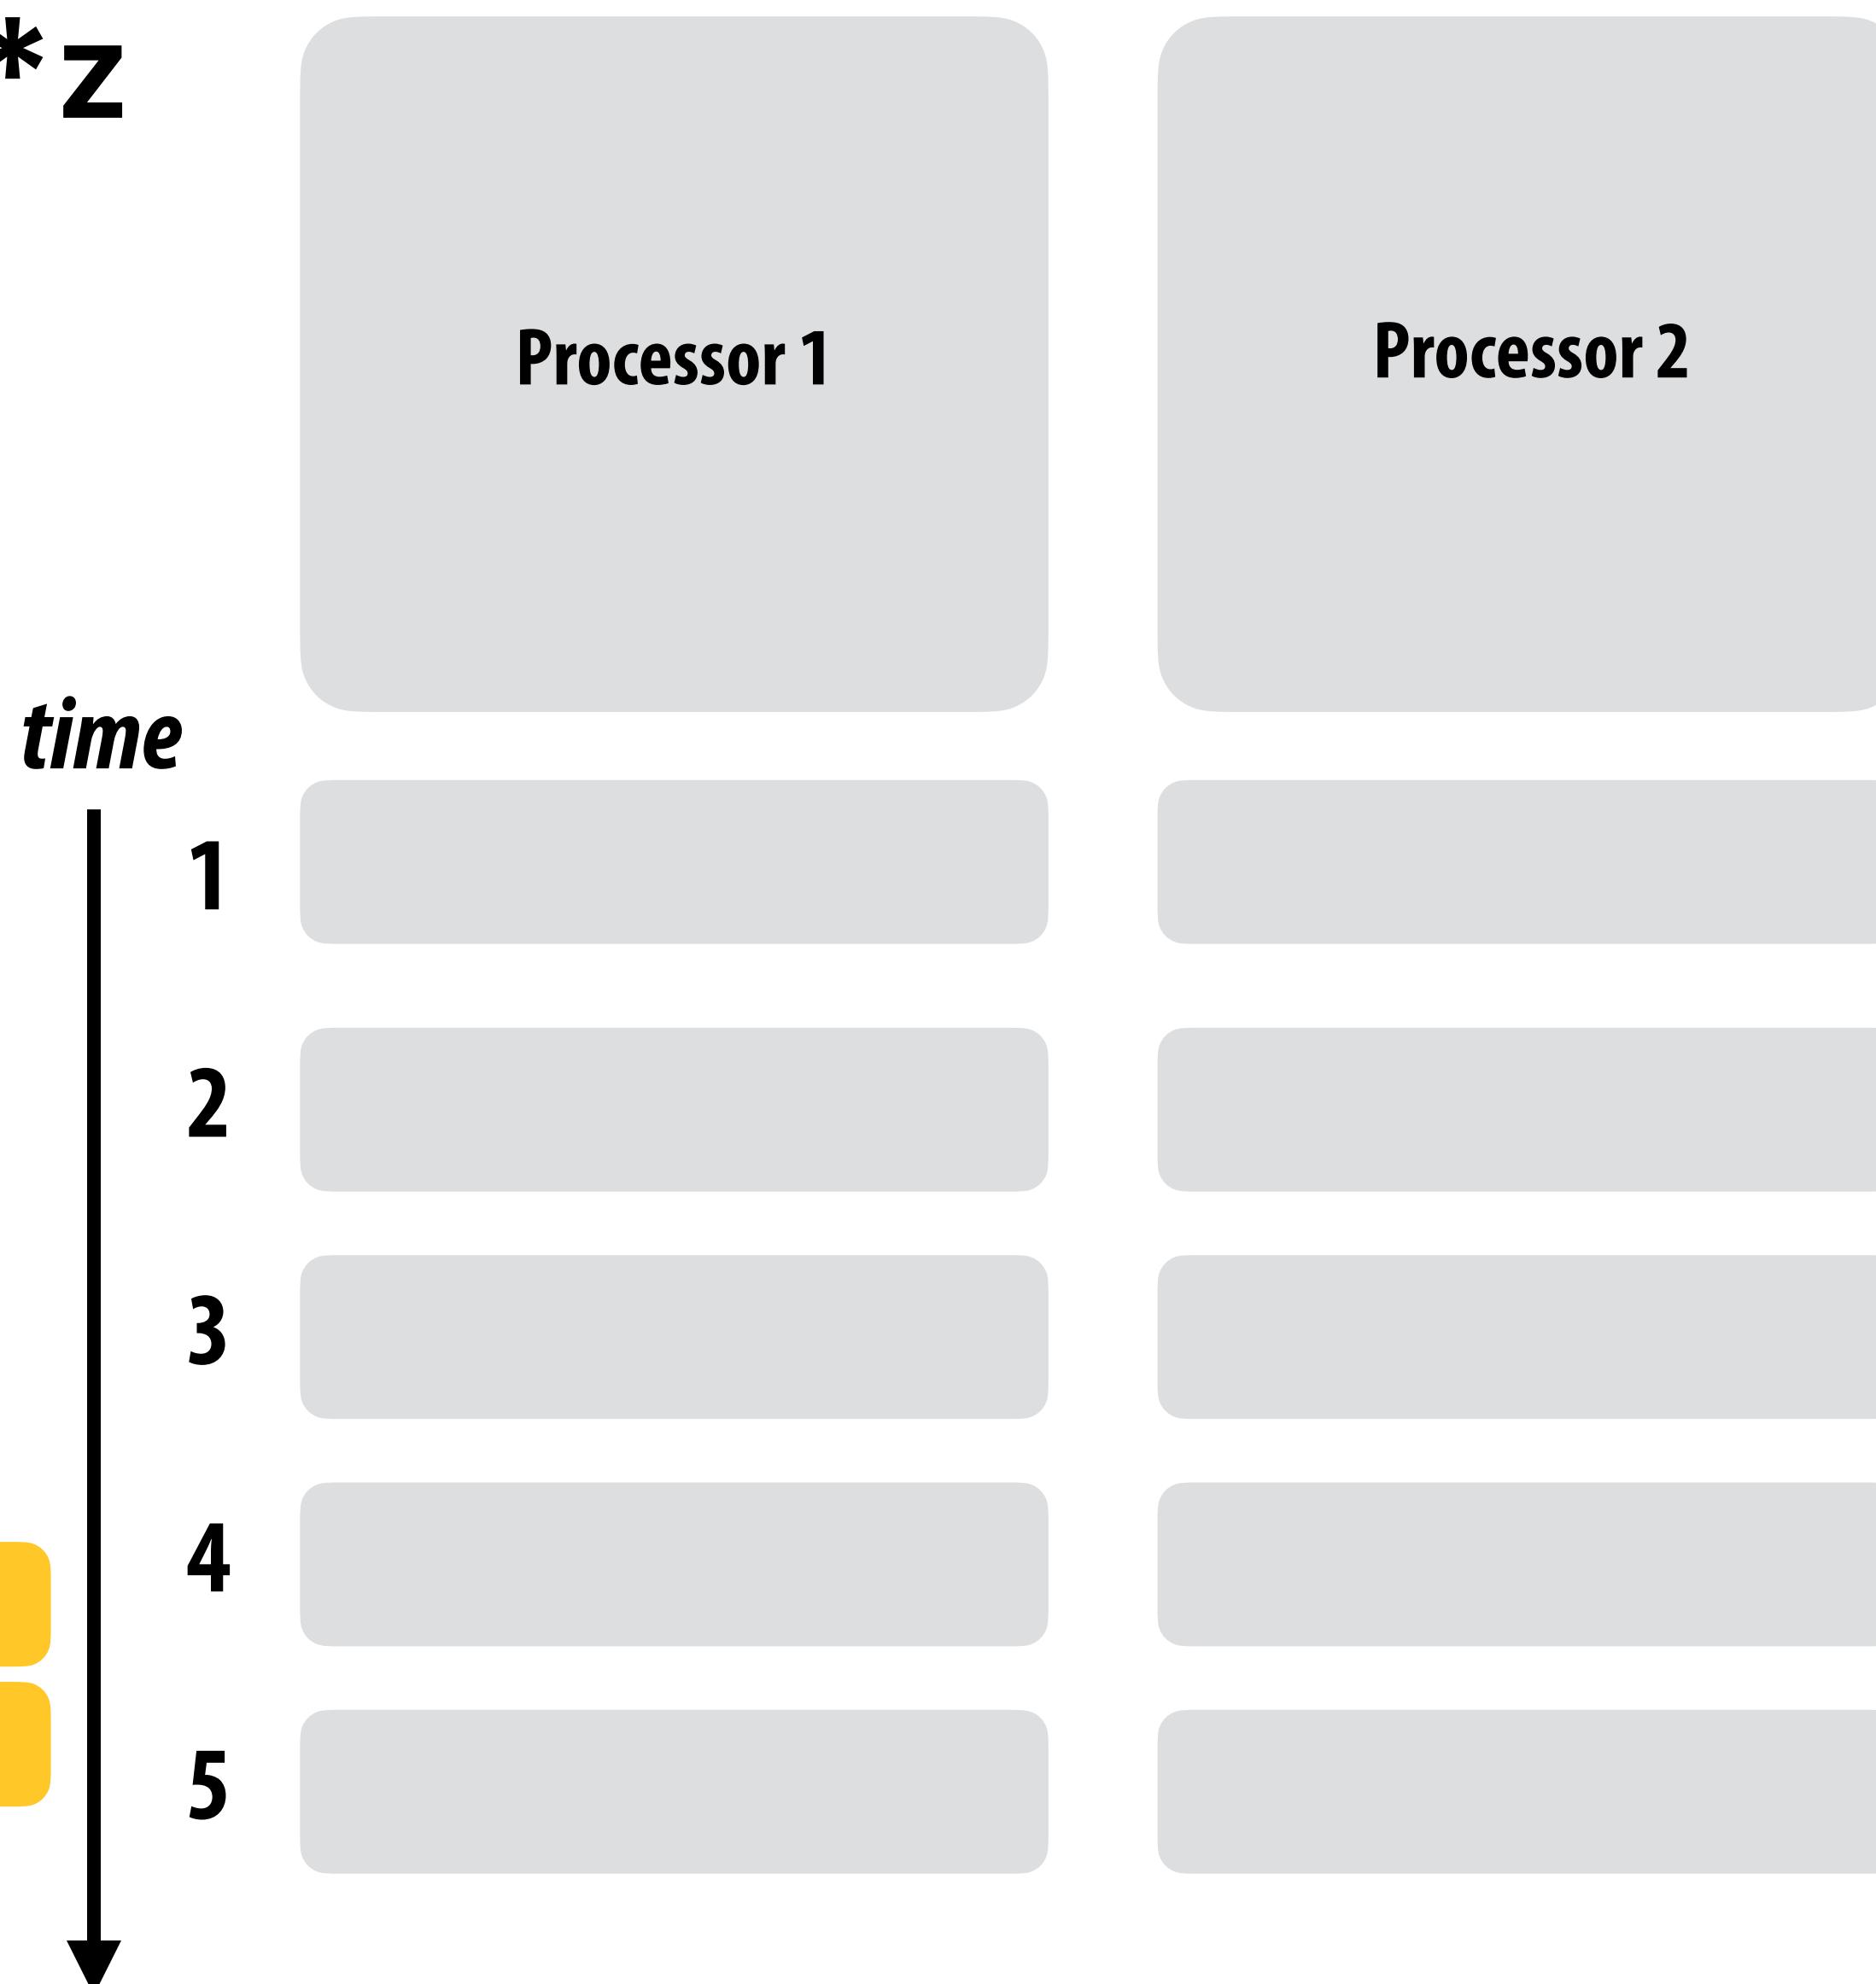
Assume register

$$R0 = x, R1 = y, R2 = z$$

```
1 mul R0, R0, R0
2 mul R1, R1, R1
3 mul R2, R2, R2
4 add R0, R0, R1
5 add R3, R0, R2
```

*R3 now stores value of
program variable 'a'*

- 1. mul R0, R0, R0
- 2. mul R1, R1, R1
- 3. mul R2, R2, R2
- 4. add R0, R0, R1
- 5. add R3, R0, R2



What if up to two instructions can be performed at once?

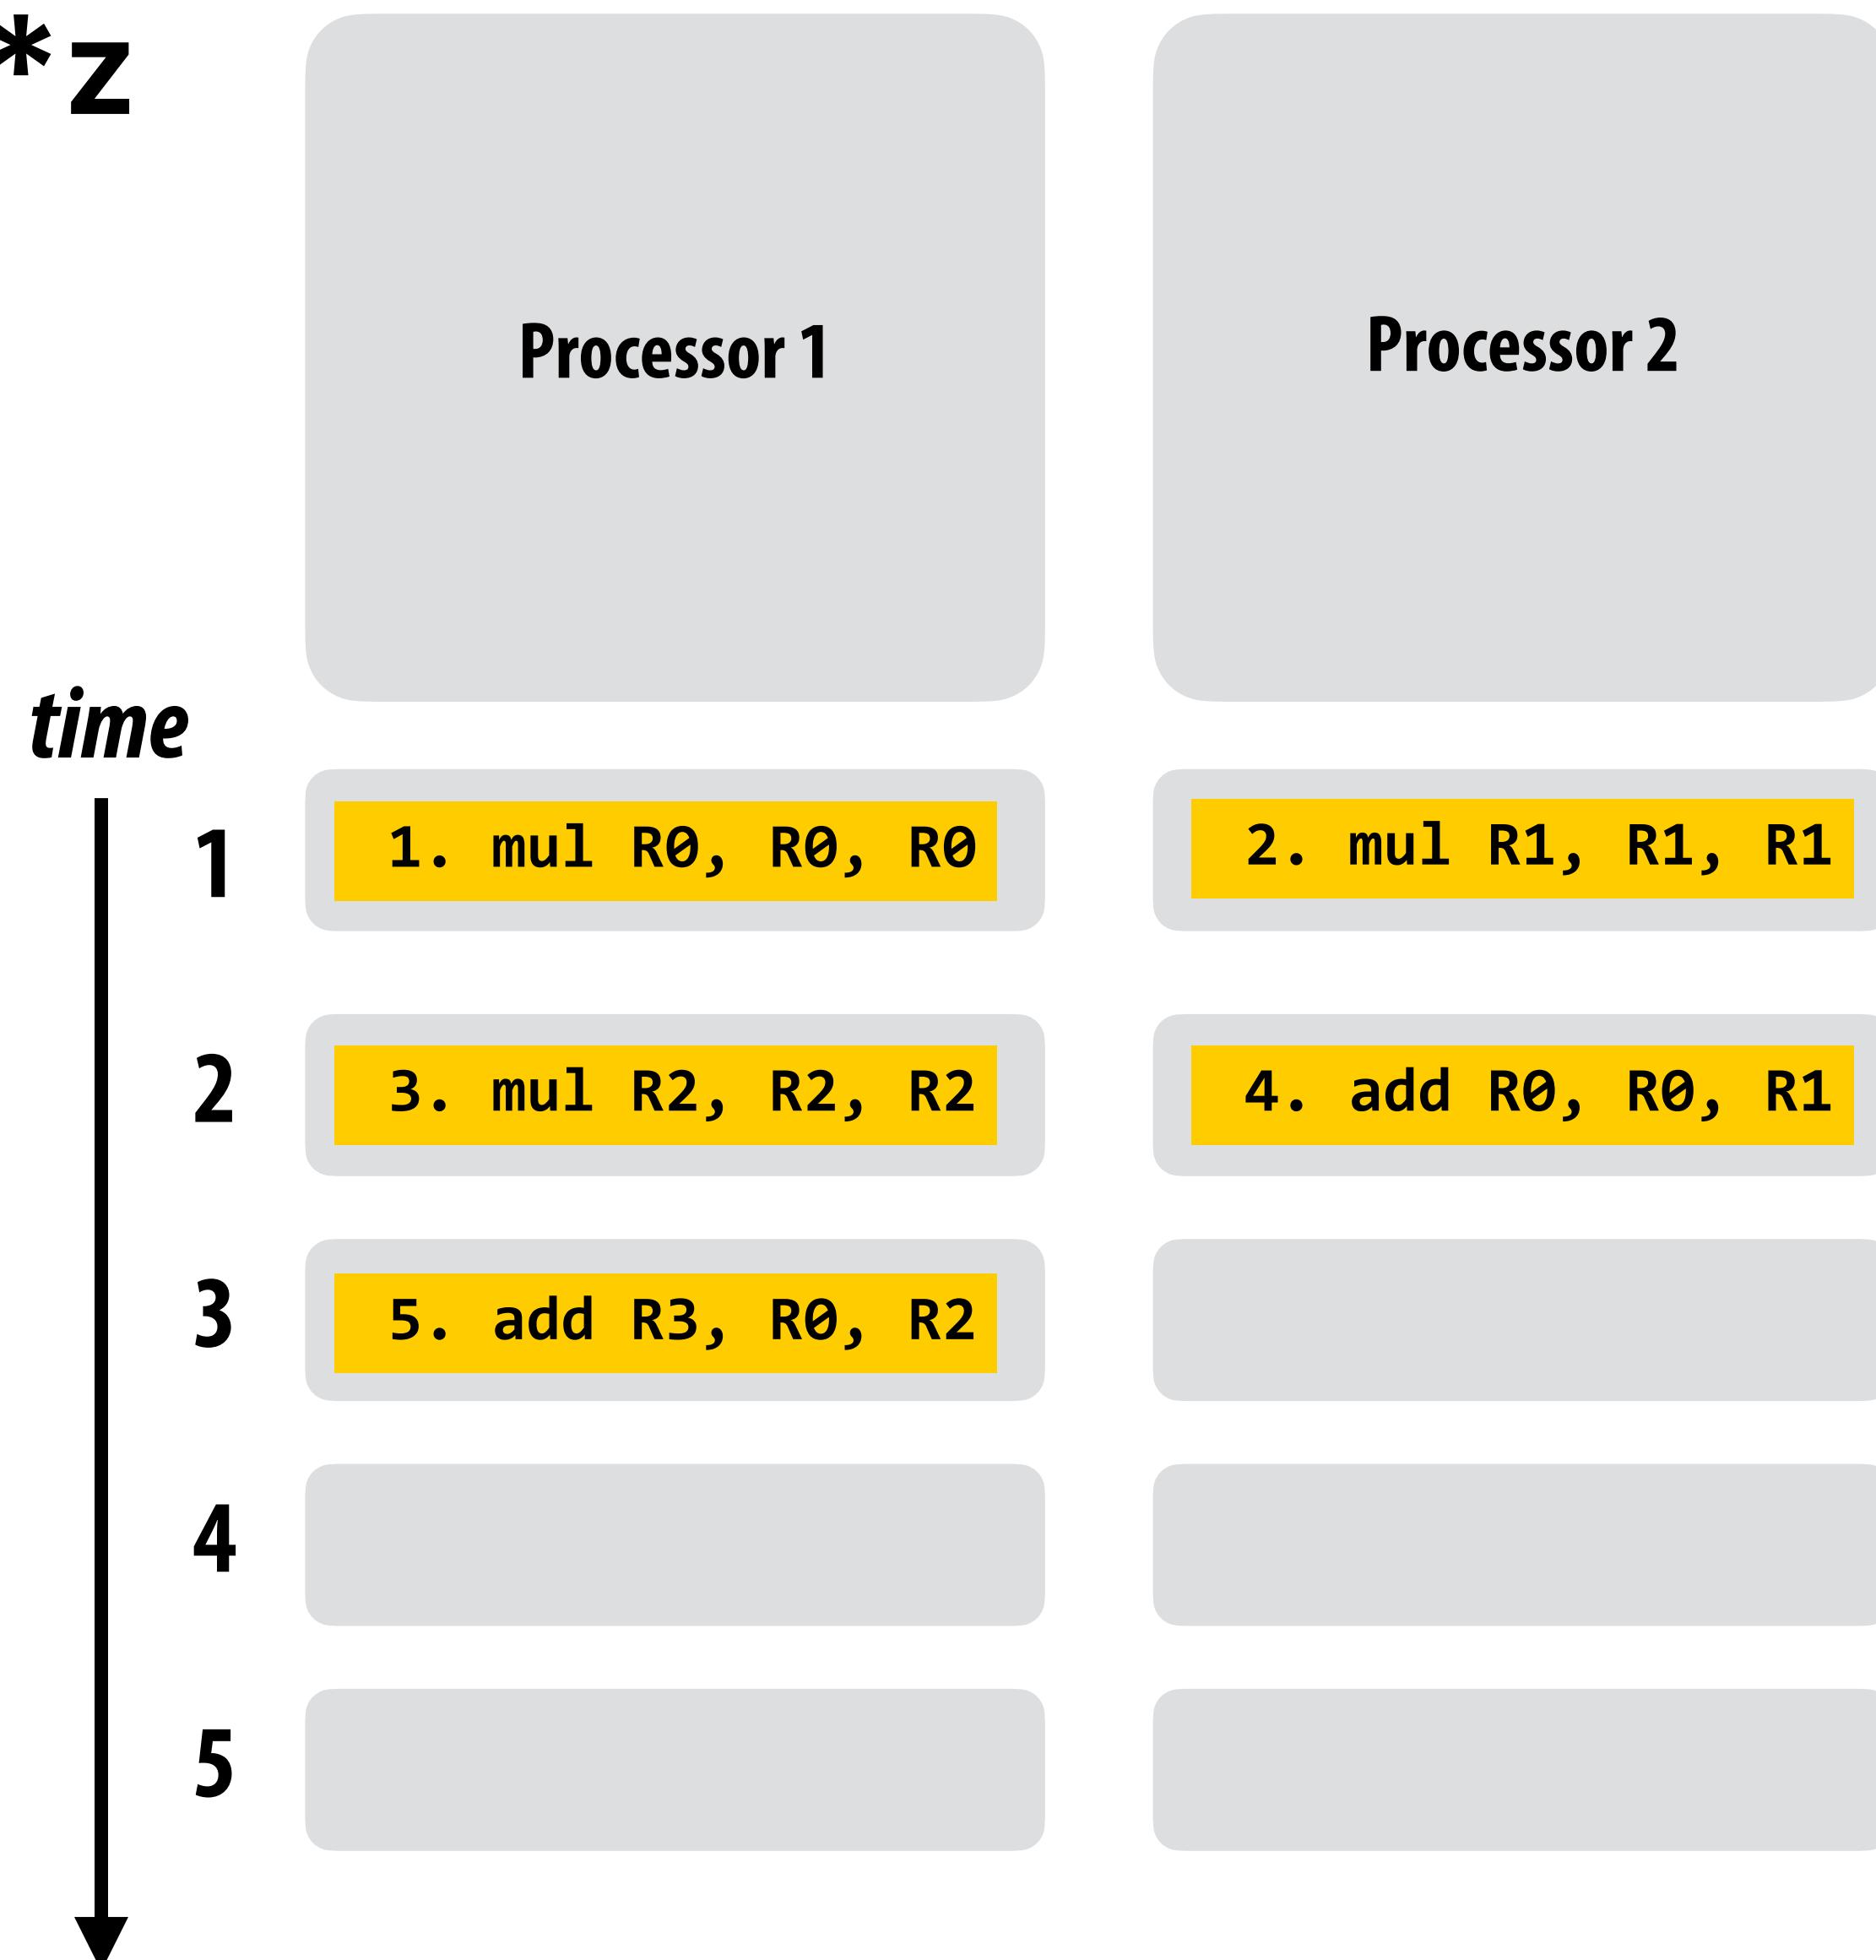
$$a = x*x + y*y + z*z$$

Assume register

$$R0 = x, R1 = y, R2 = z$$

```
1 mul R0, R0, R0
2 mul R1, R1, R1
3 mul R2, R2, R2
4 add R0, R0, R1
5 add R3, R0, R2
```

*R3 now stores value of
program variable 'a'*



QUESTION:

**What does it mean for our parallel to scheduling to
that “respects program order”?**

Hint: What is expected of the output.

What about three instructions at once?

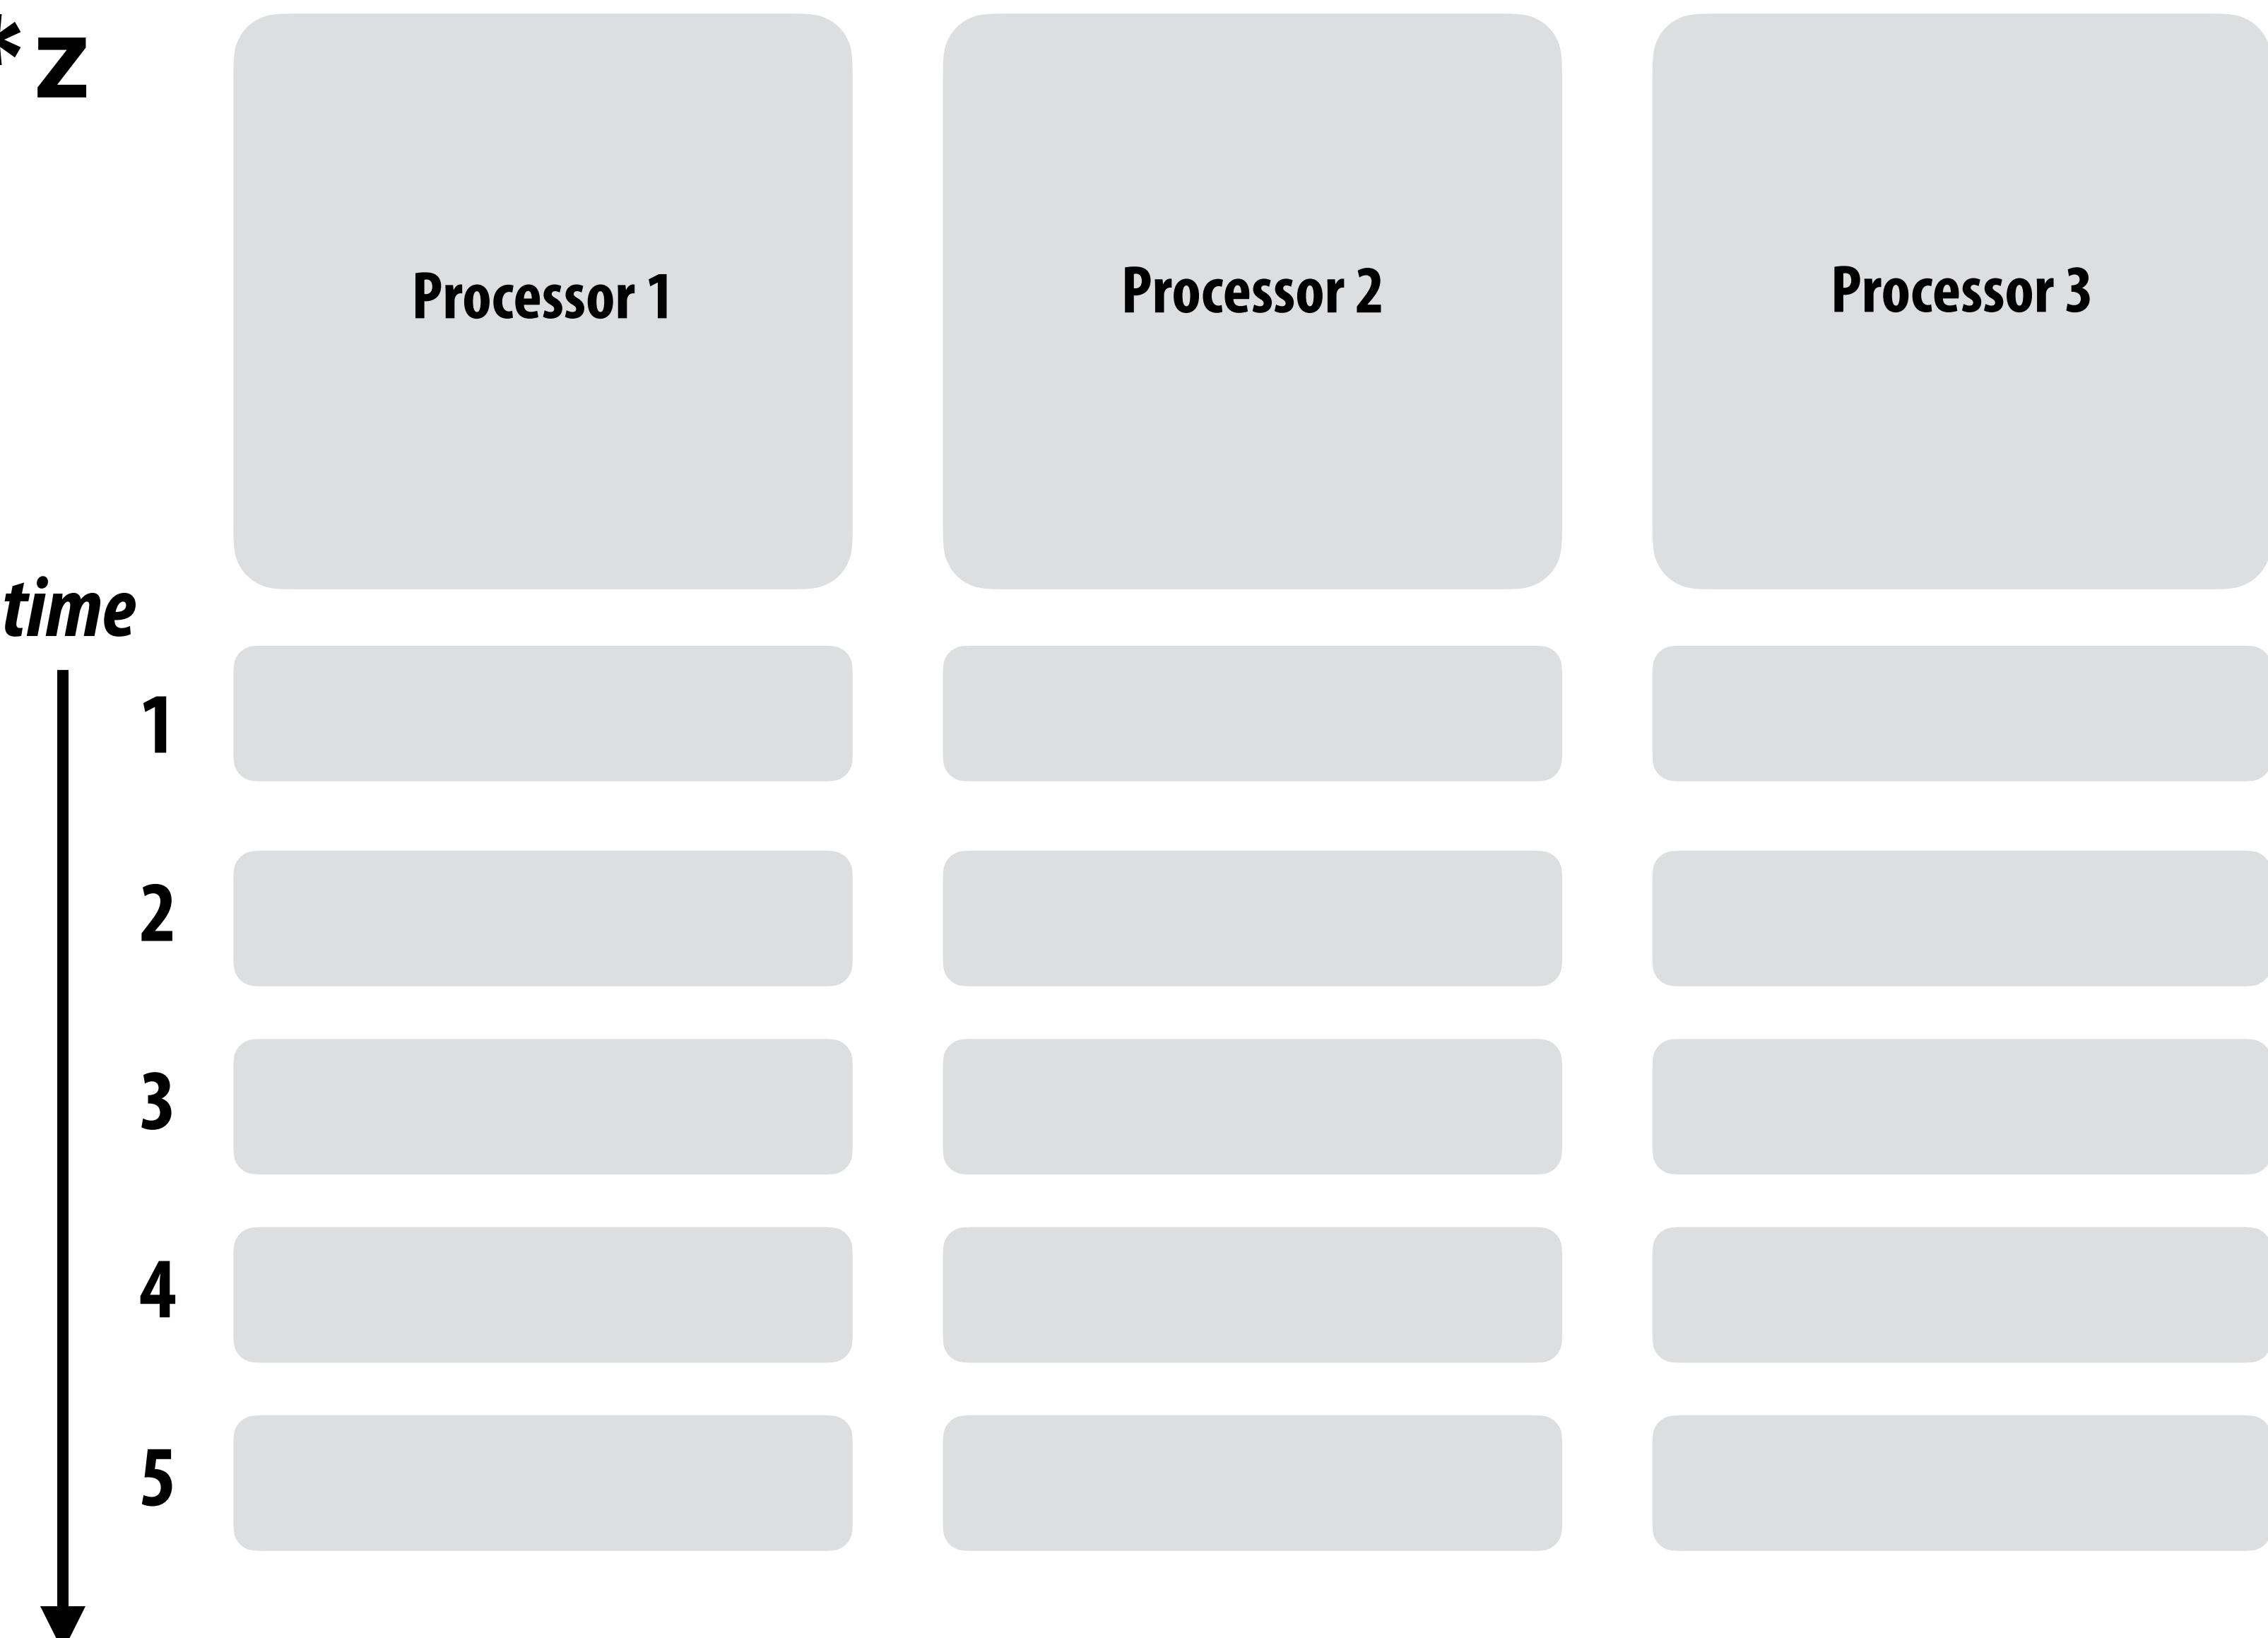
$$a = x*x + y*y + z*z$$

Assume register

$$R0 = x, R1 = y, R2 = z$$

```
1 mul R0, R0, R0
2 mul R1, R1, R1
3 mul R2, R2, R2
4 add R0, R0, R1
5 add R3, R0, R2
```

*R3 now stores value of
program variable 'a'*



What about three instructions at once?

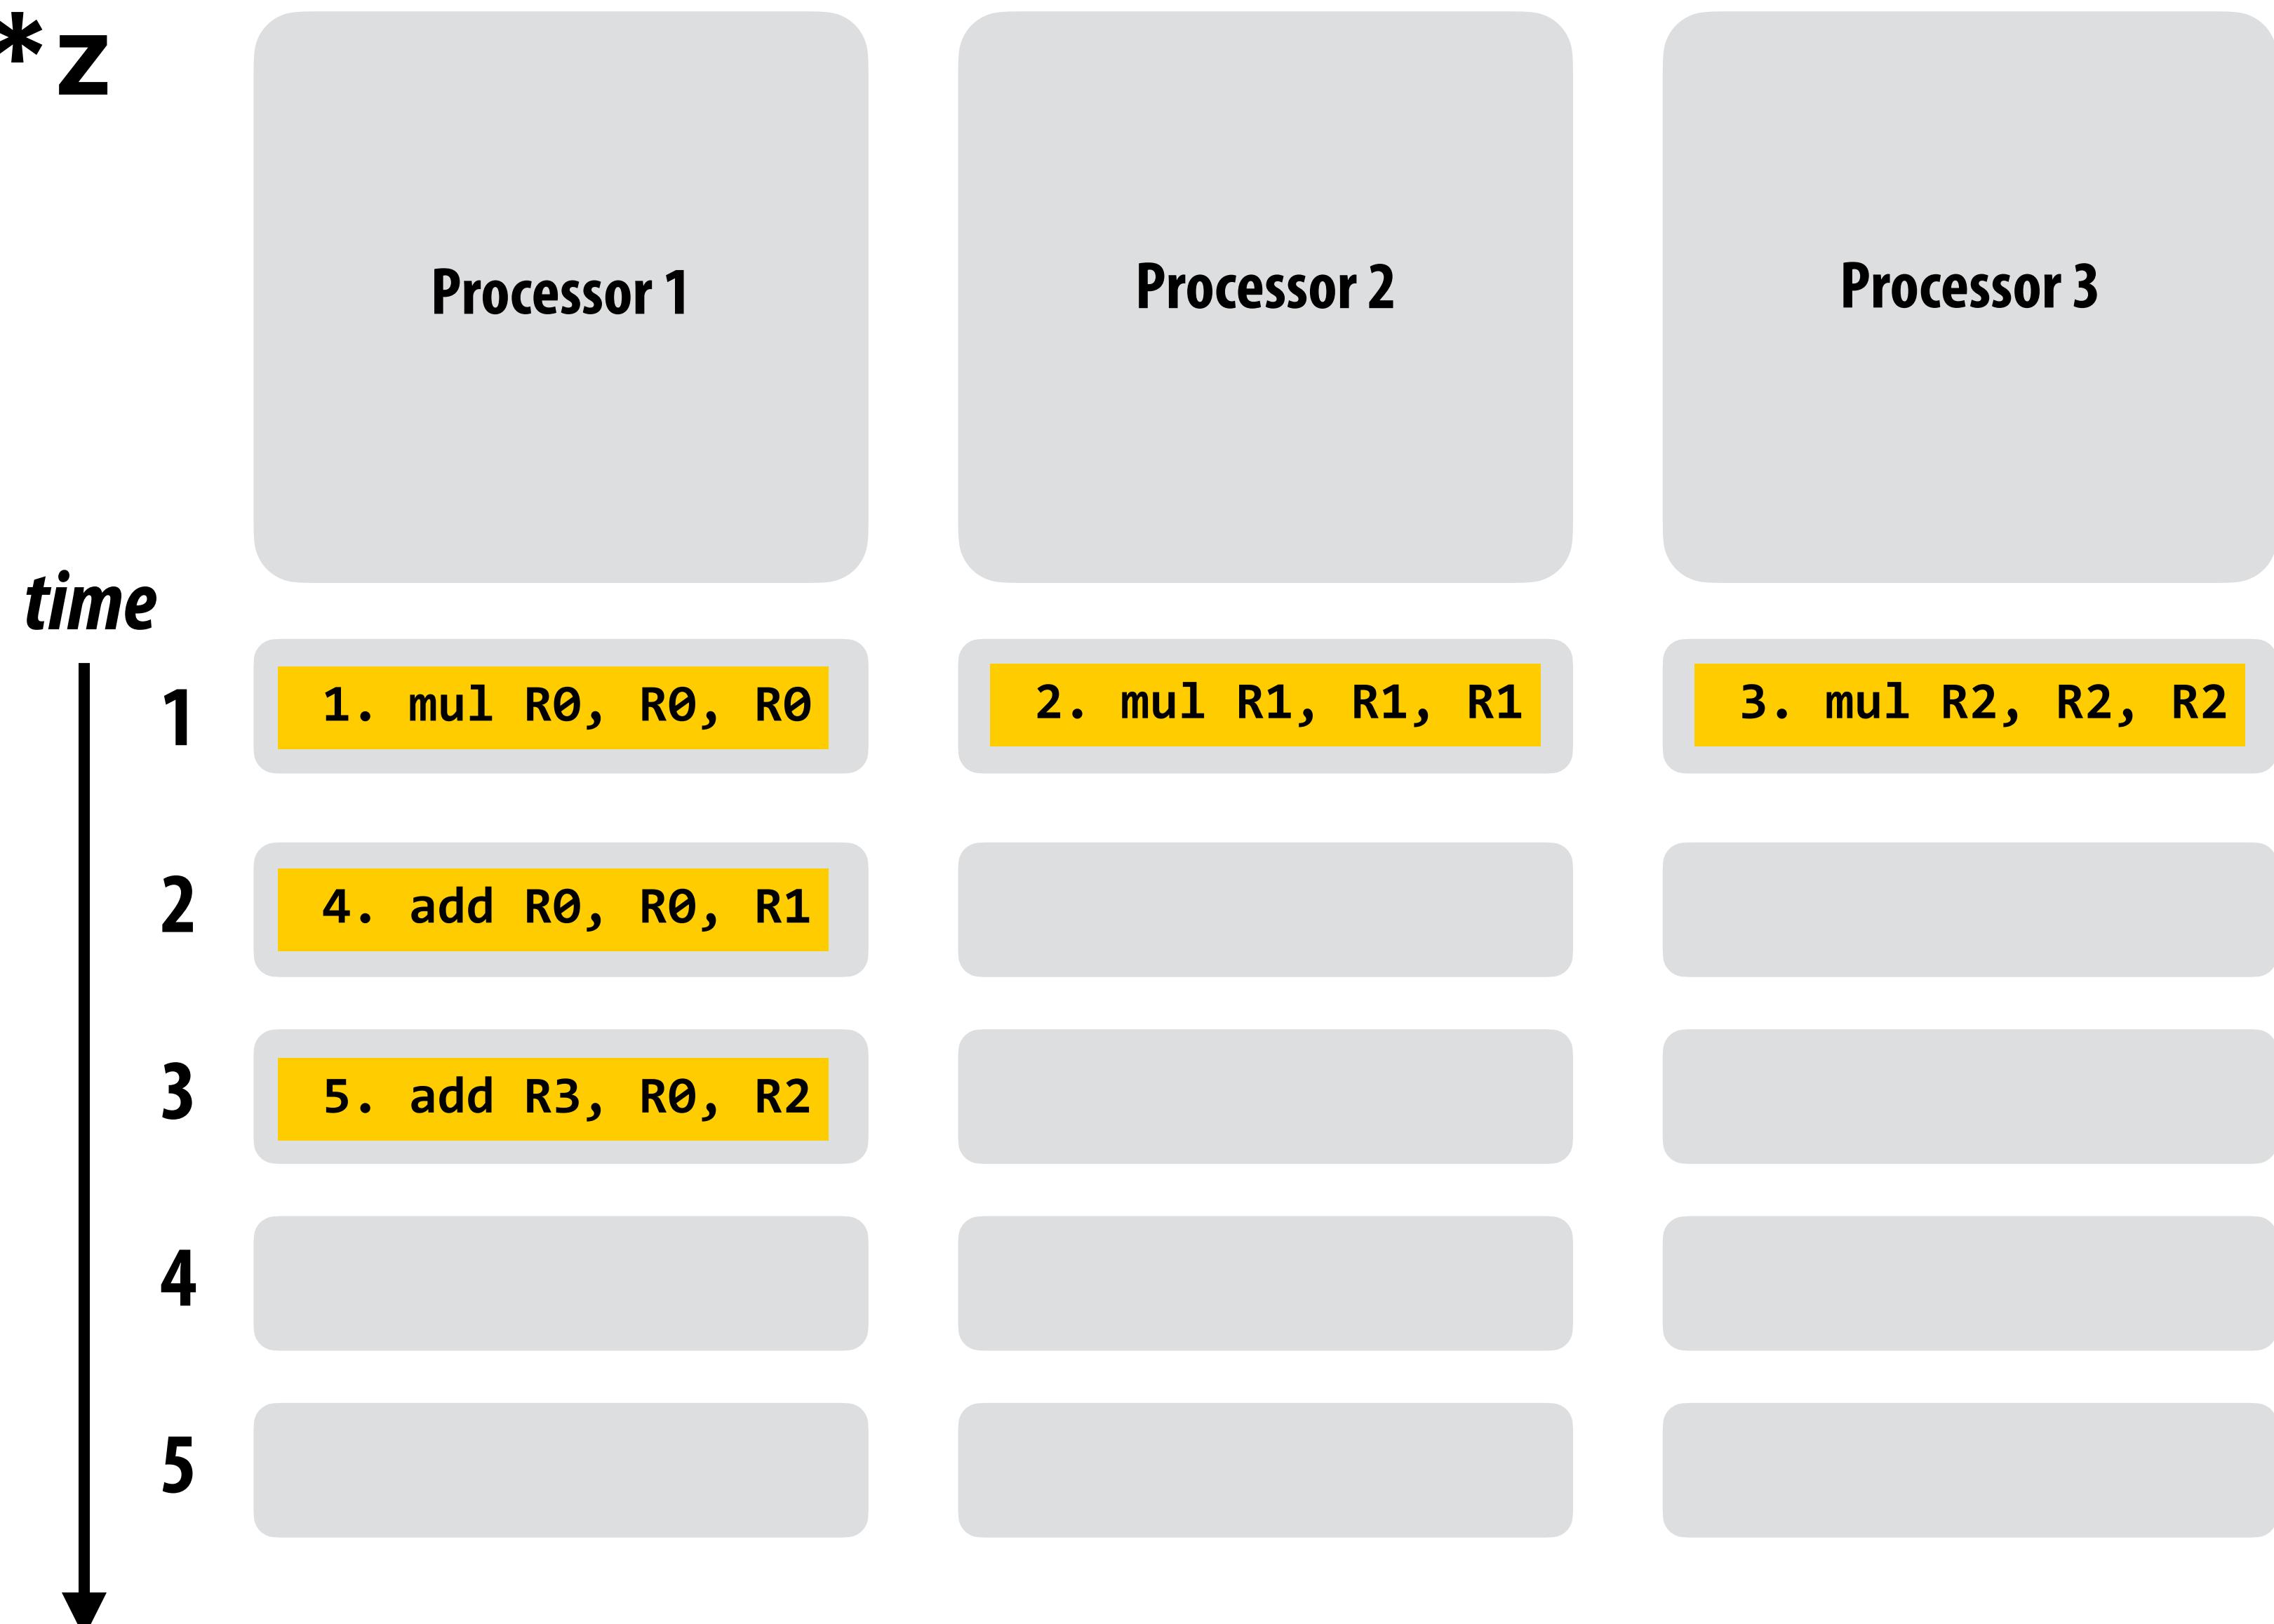
$$a = x*x + y*y + z*z$$

Assume register

$R0 = x, R1 = y, R2 = z$

```
1 mul R0, R0, R0
2 mul R1, R1, R1
3 mul R2, R2, R2
4 add R0, R0, R1
5 add R3, R0, R2
```

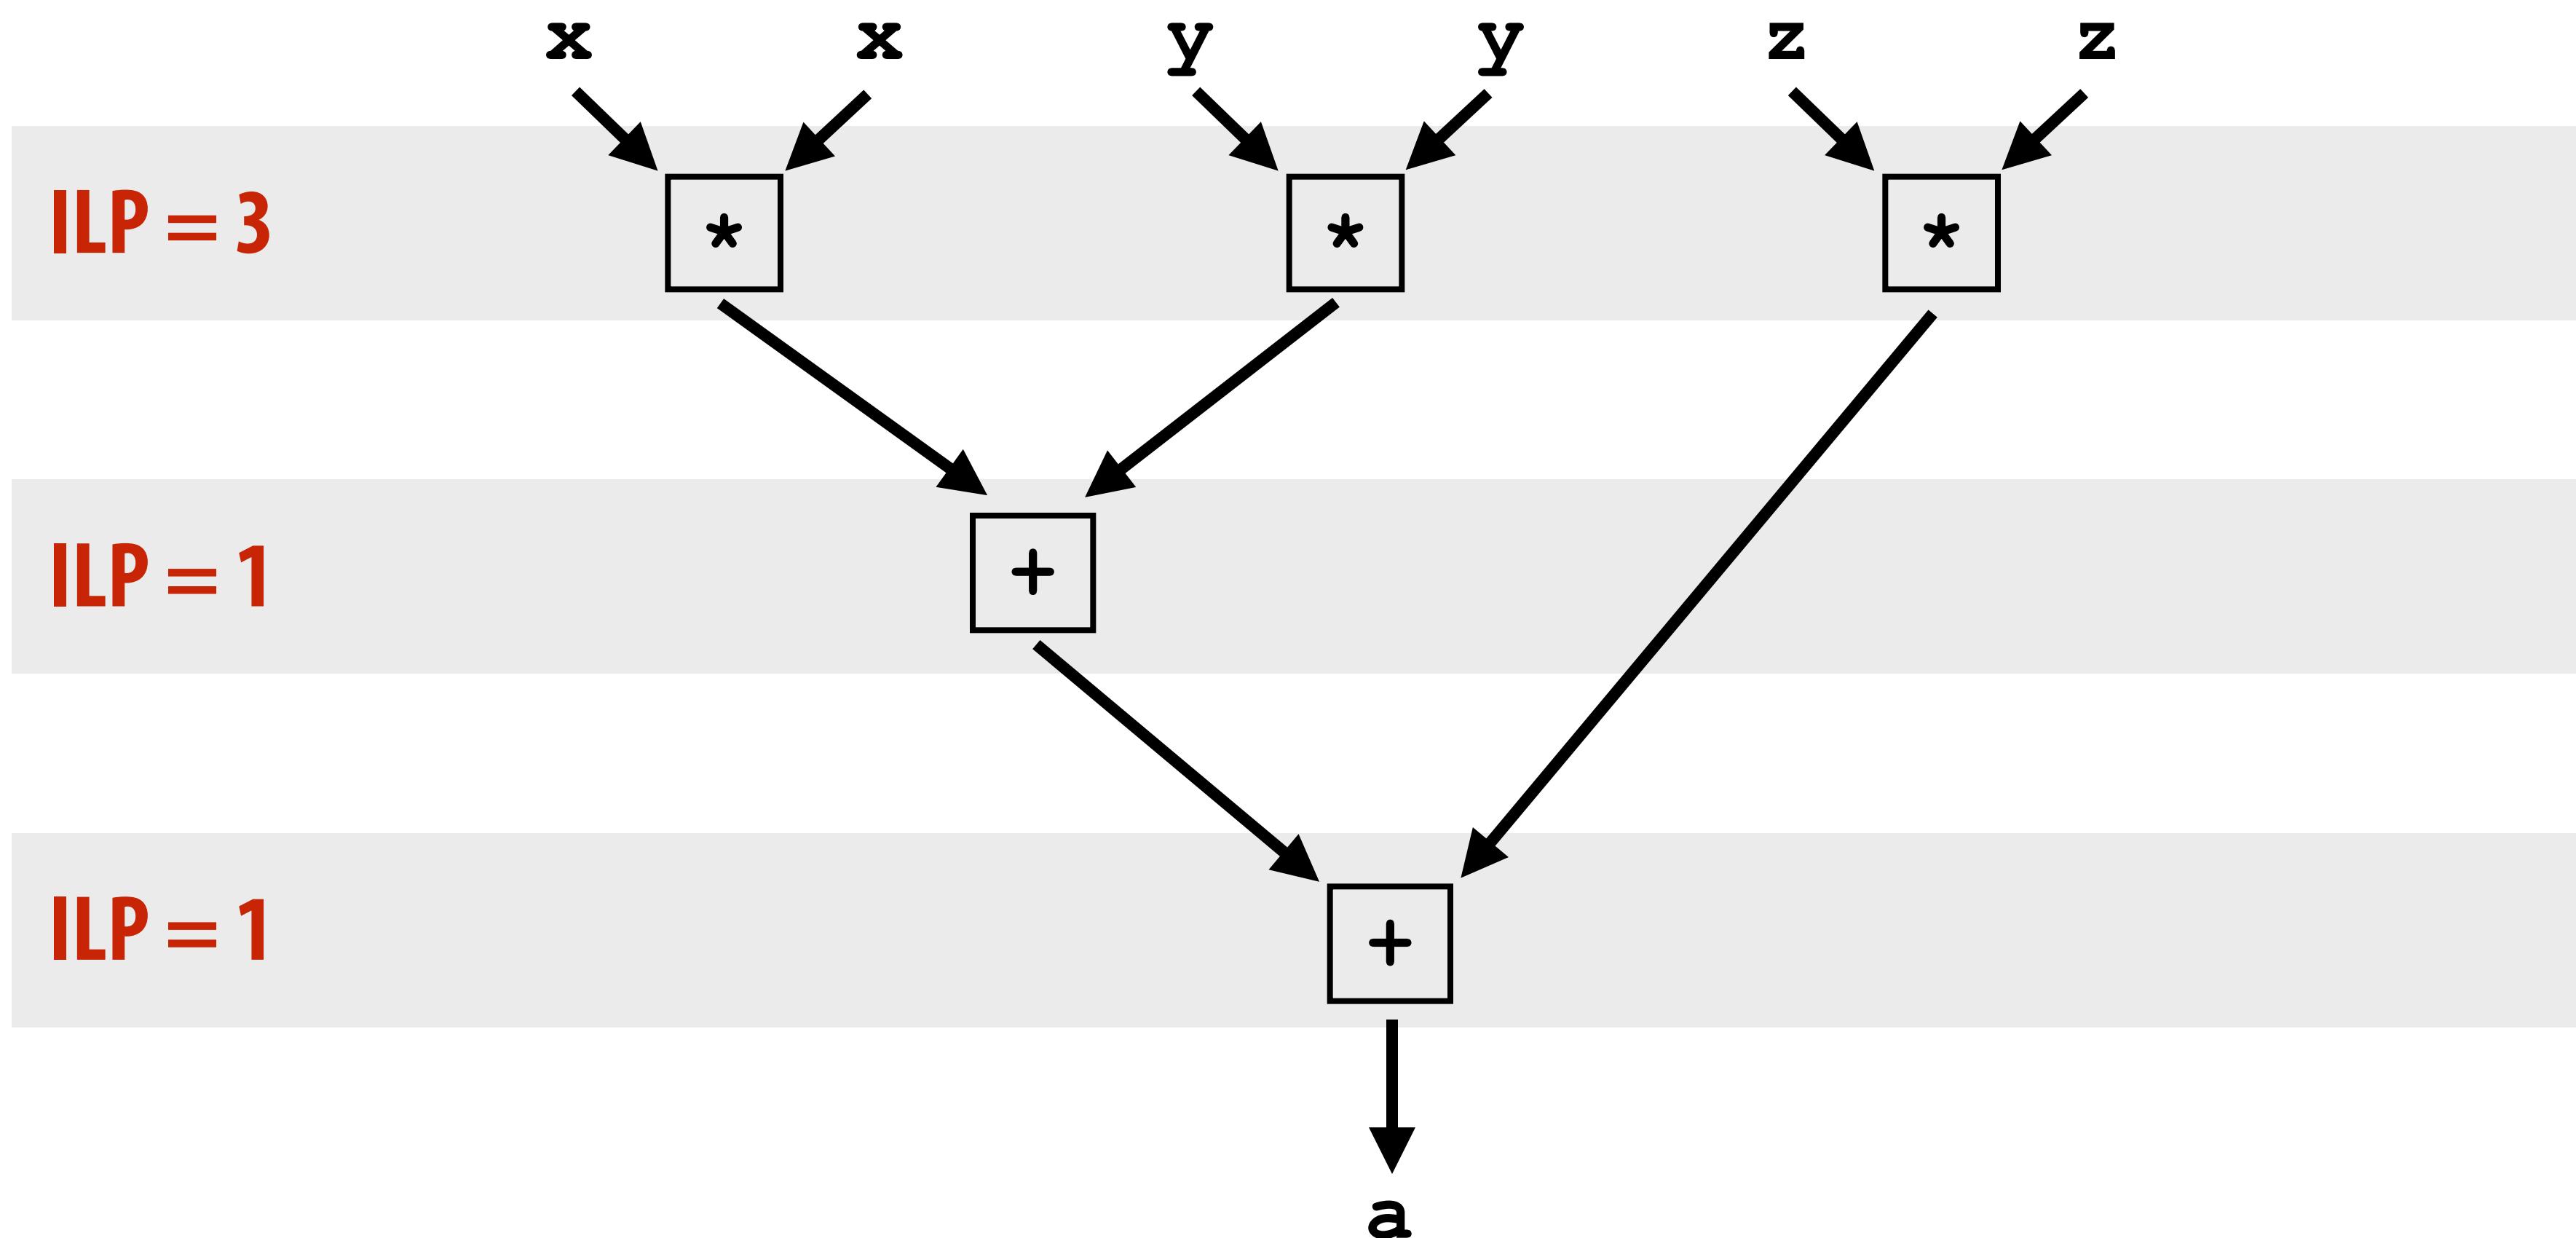
*R3 now stores value of
program variable 'a'*



Instruction level parallelism (ILP) example

- ILP = 3

$$a = x*x + y*y + z*z$$



Superscalar processor execution

$$a = x*x + y*y + z*z$$

Assume register

$R0 = x, R1 = y, R2 = z$

```
1 mul R0, R0, R0
2 mul R1, R1, R1
3 mul R2, R2, R2
4 add R0, R0, R1
5 add R3, R0, R2
```

Idea #1:

Superscalar execution: processor automatically finds*
independent instructions in an instruction sequence and
executes them in parallel on multiple execution units!

**In this example: instructions 1, 2, and 3 can be executed in parallel without impacting program correctness
(on a superscalar processor that determines that the lack of dependencies exists)**

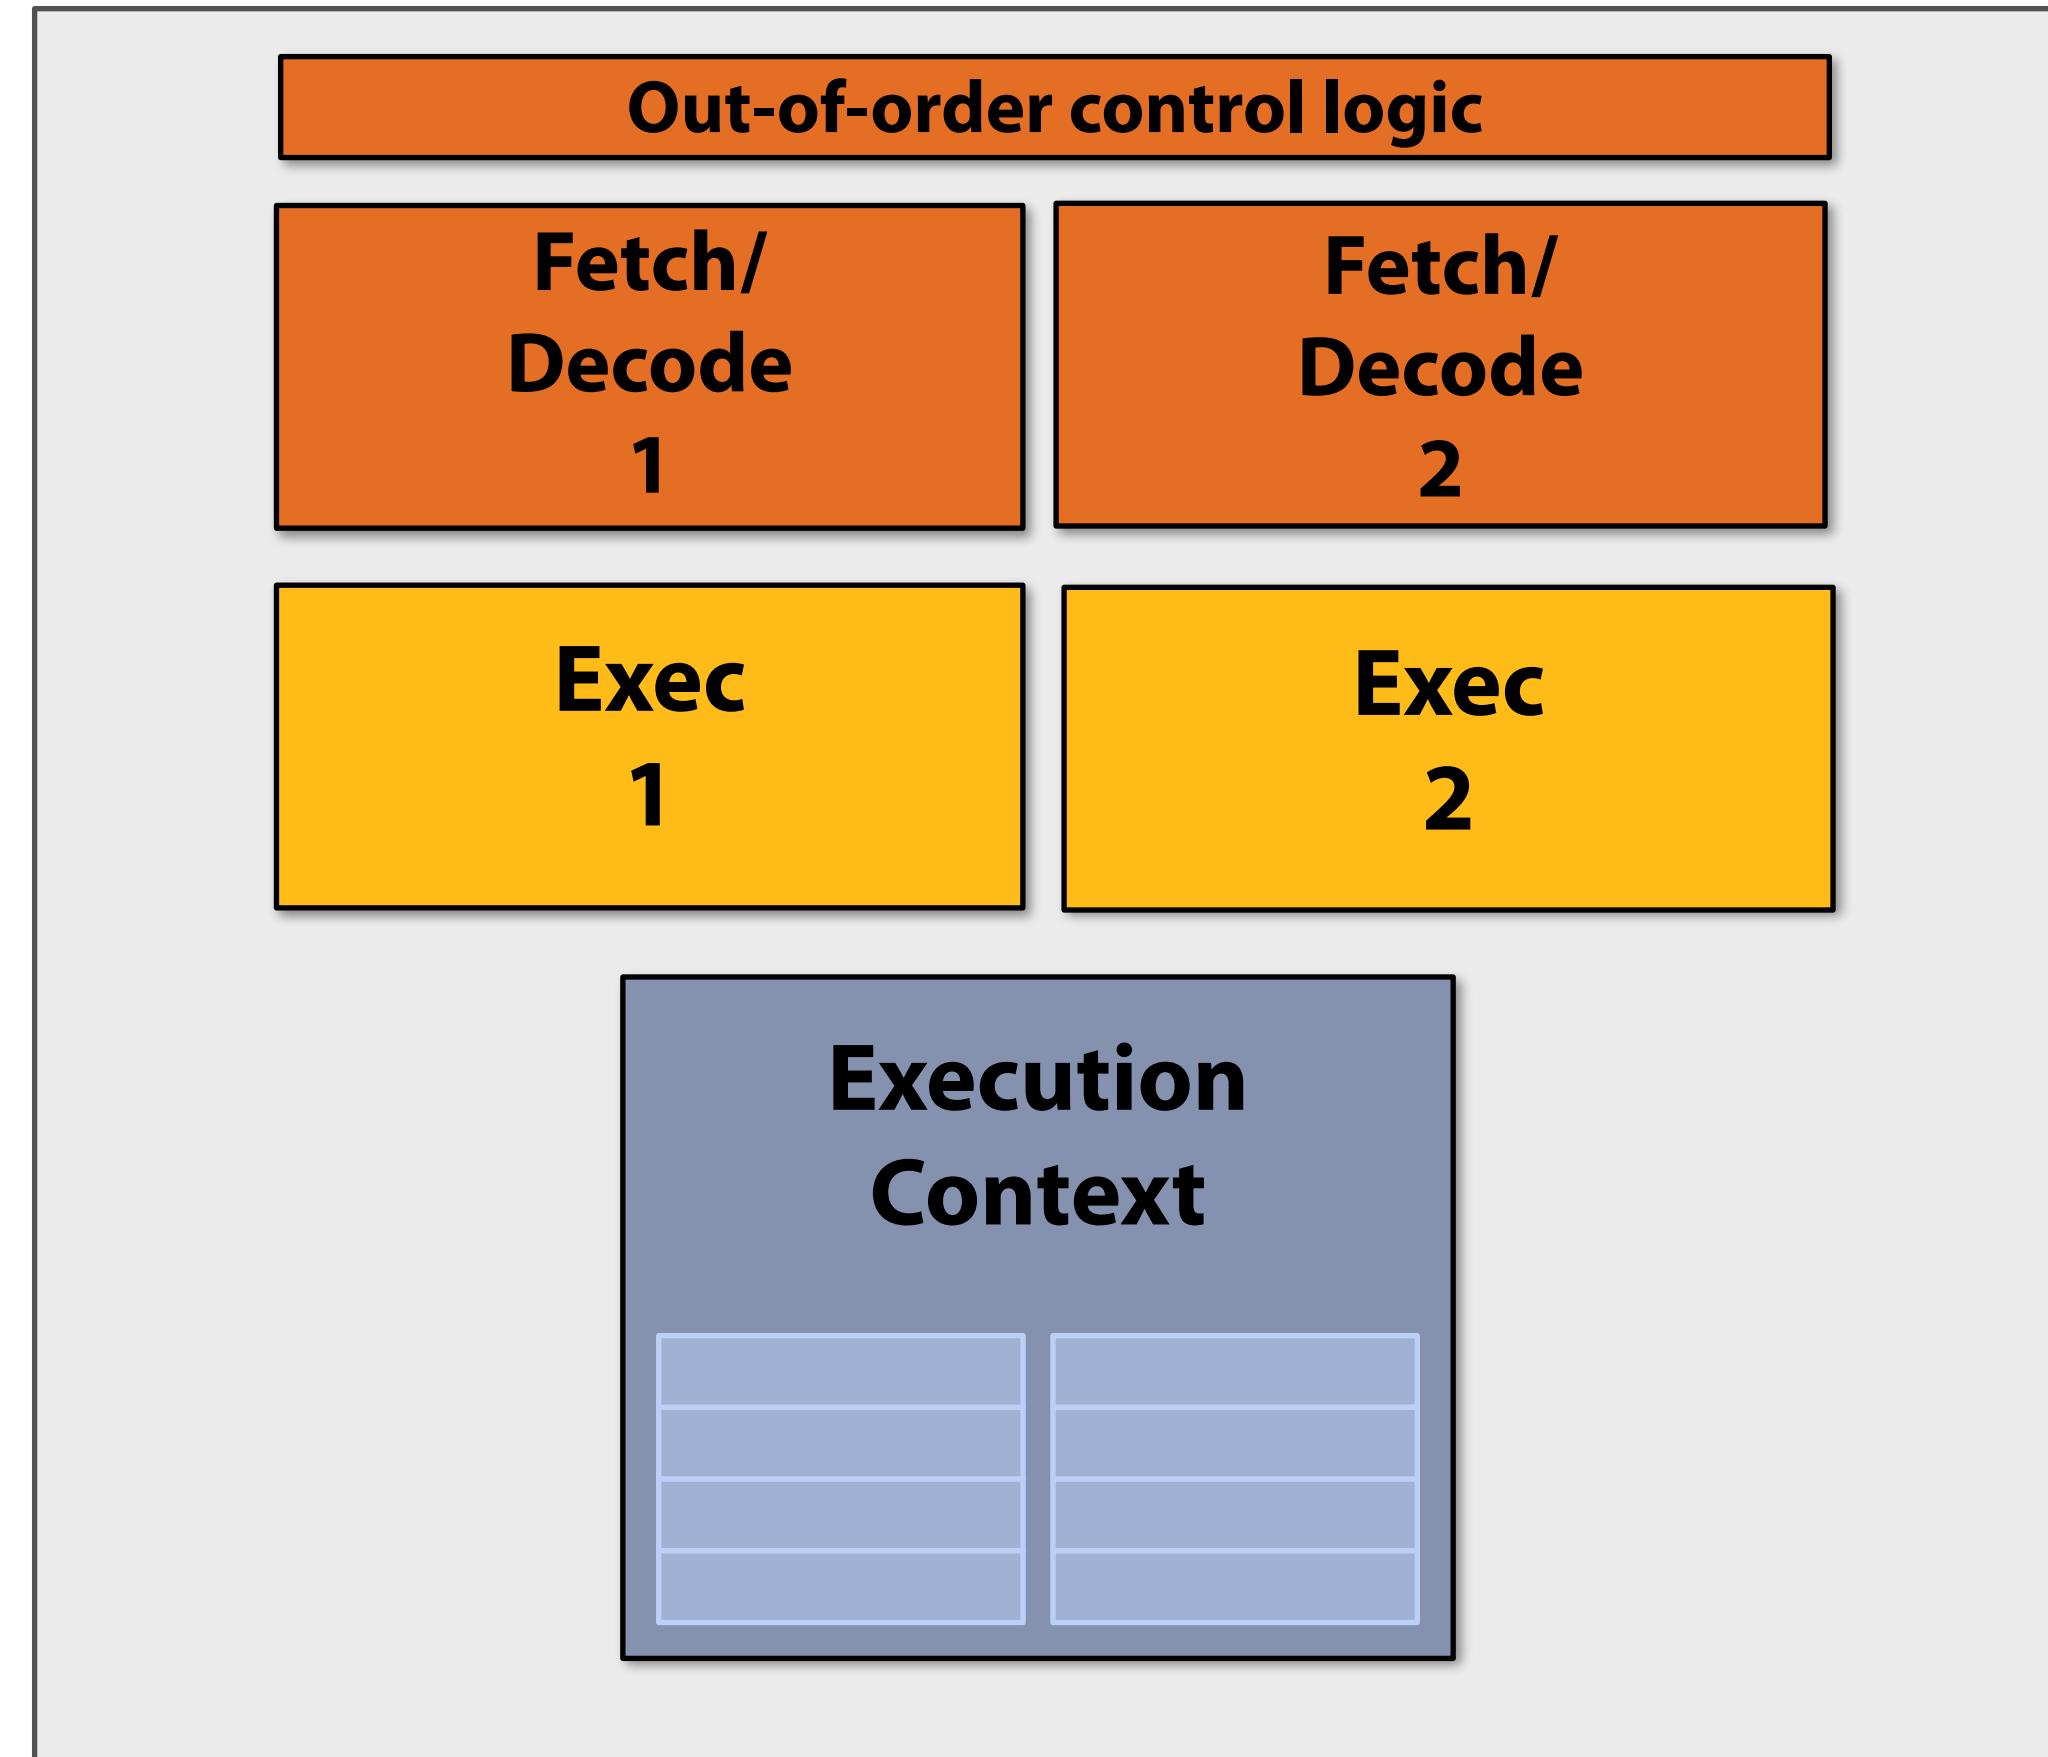
But instruction 4 must be executed after instructions 1 and 2

And instruction 5 must be executed after instruction 4

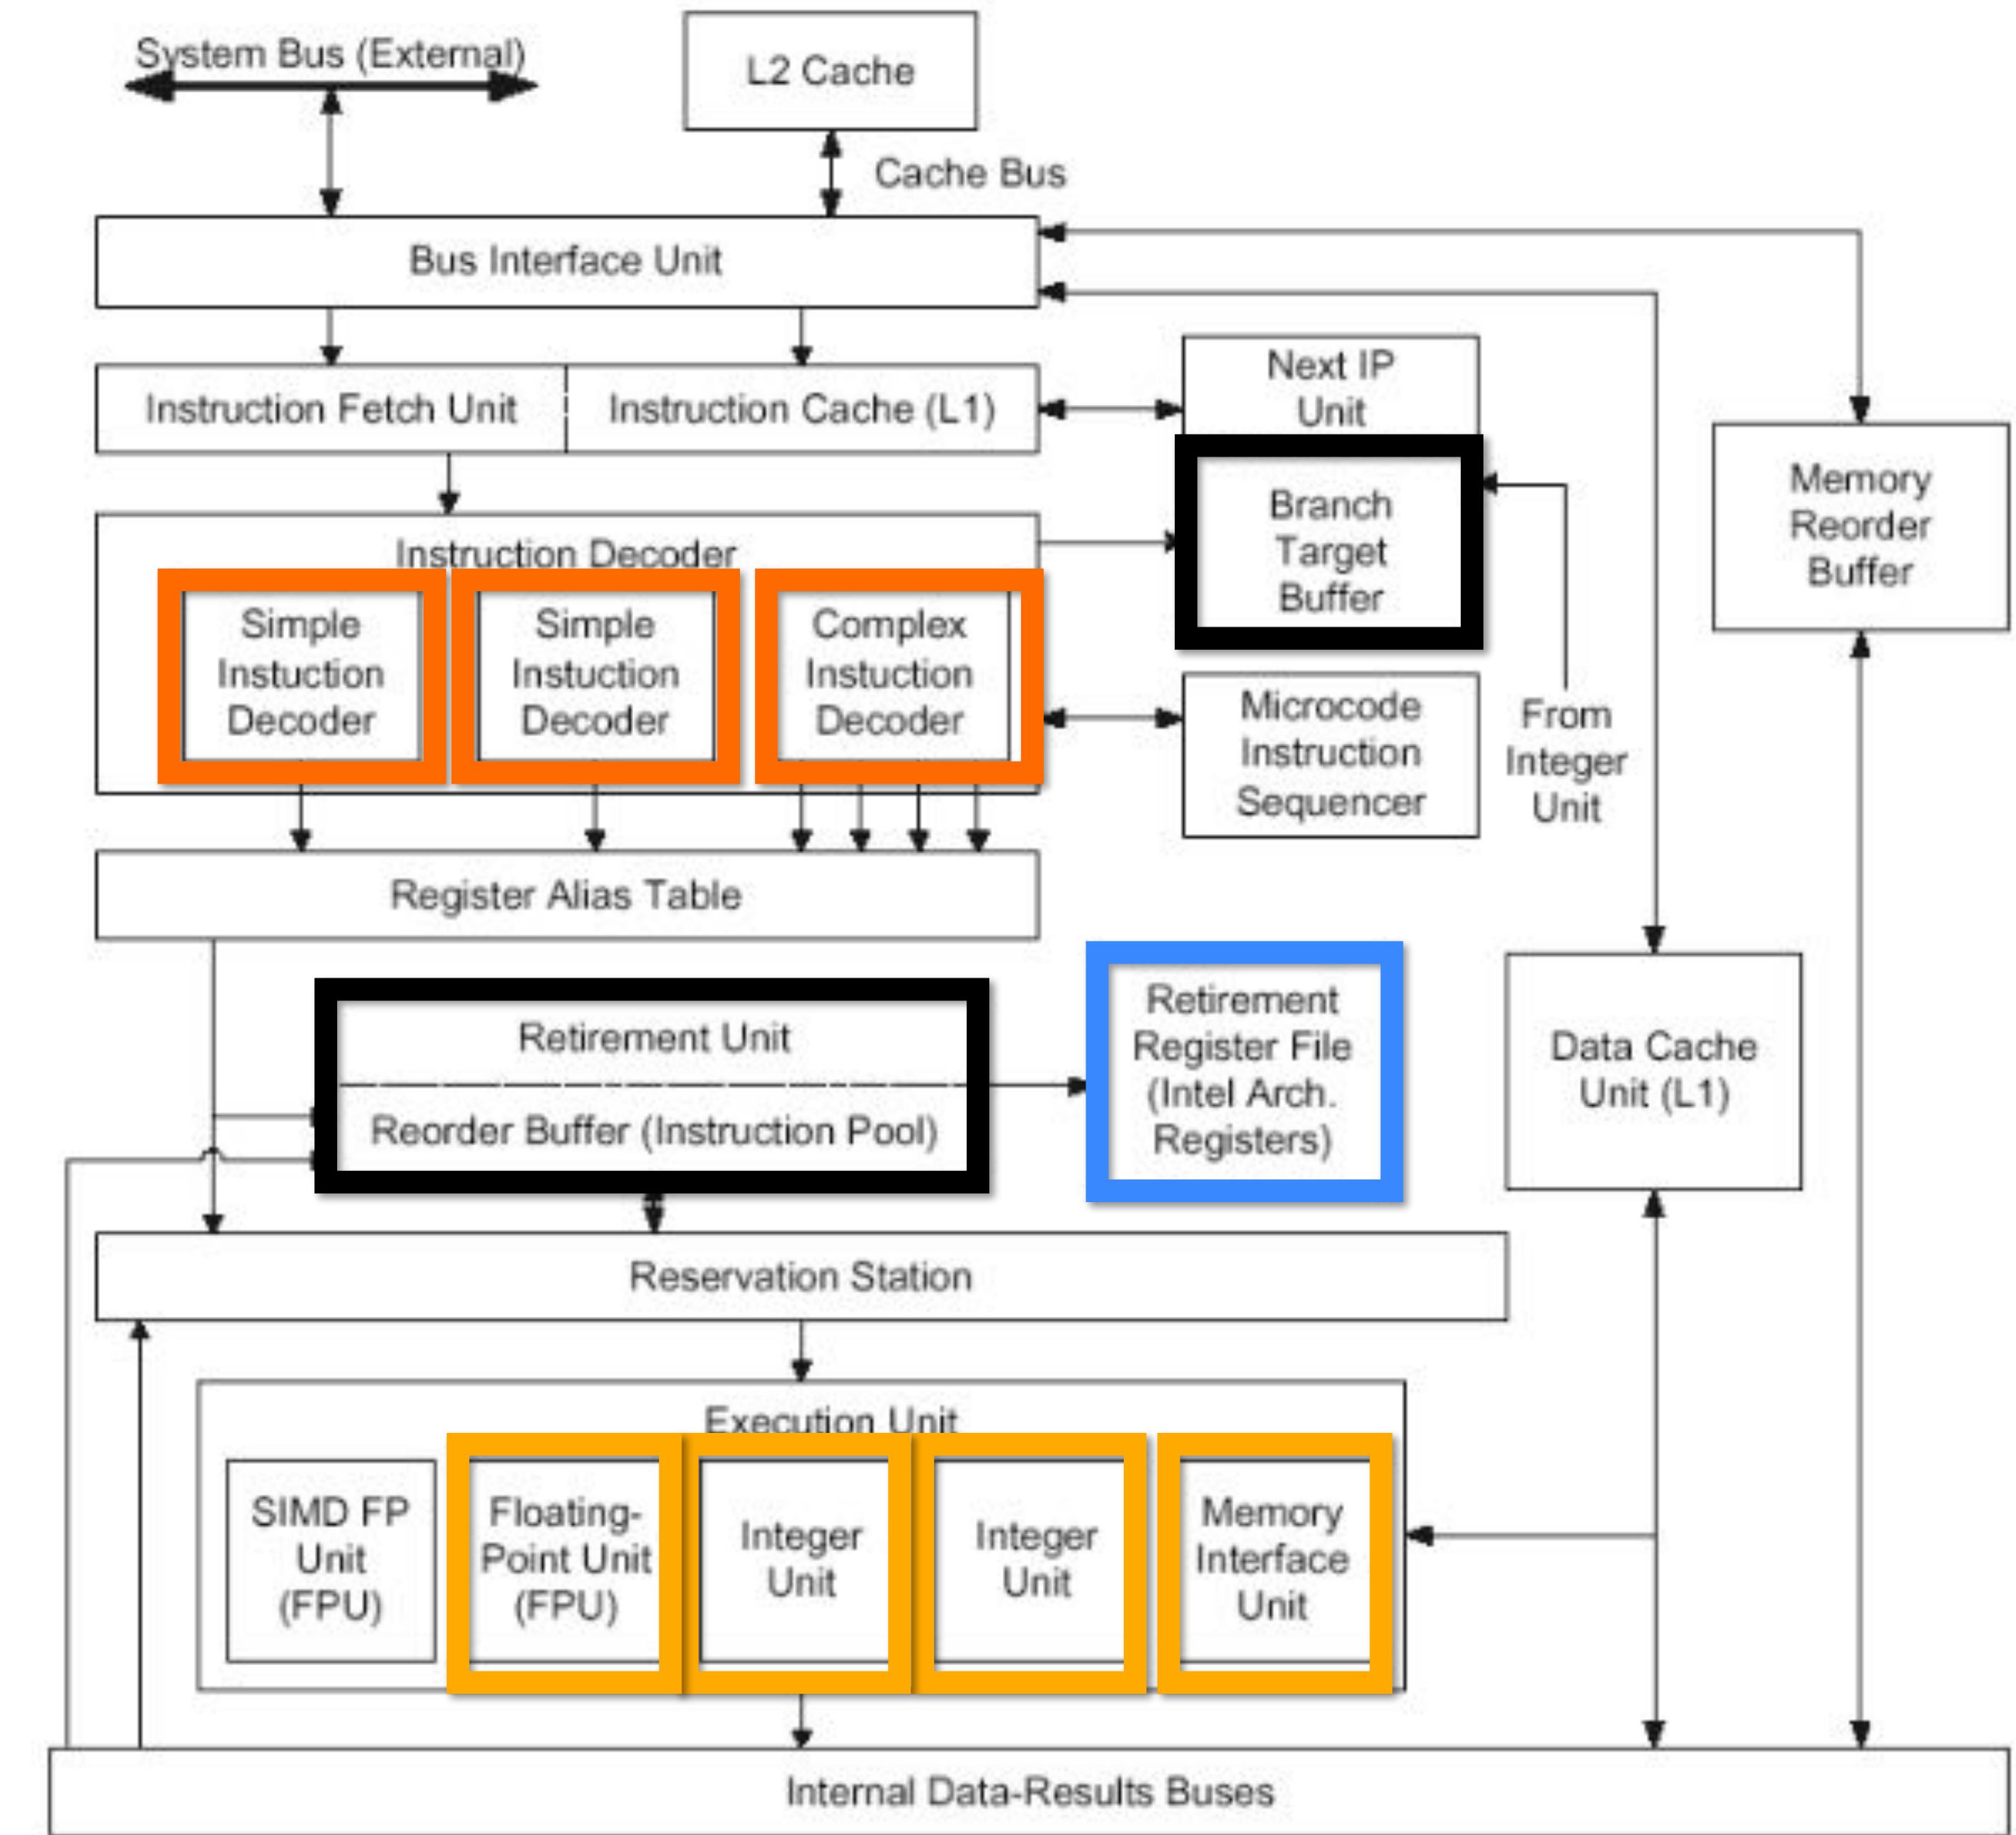
* Or the compiler finds independent instructions at compile time and explicitly encodes dependencies in the compiled binary.

Superscalar processor

This processor can decode and execute up to two instructions per clock



Aside: Old Intel Pentium 4 CPU

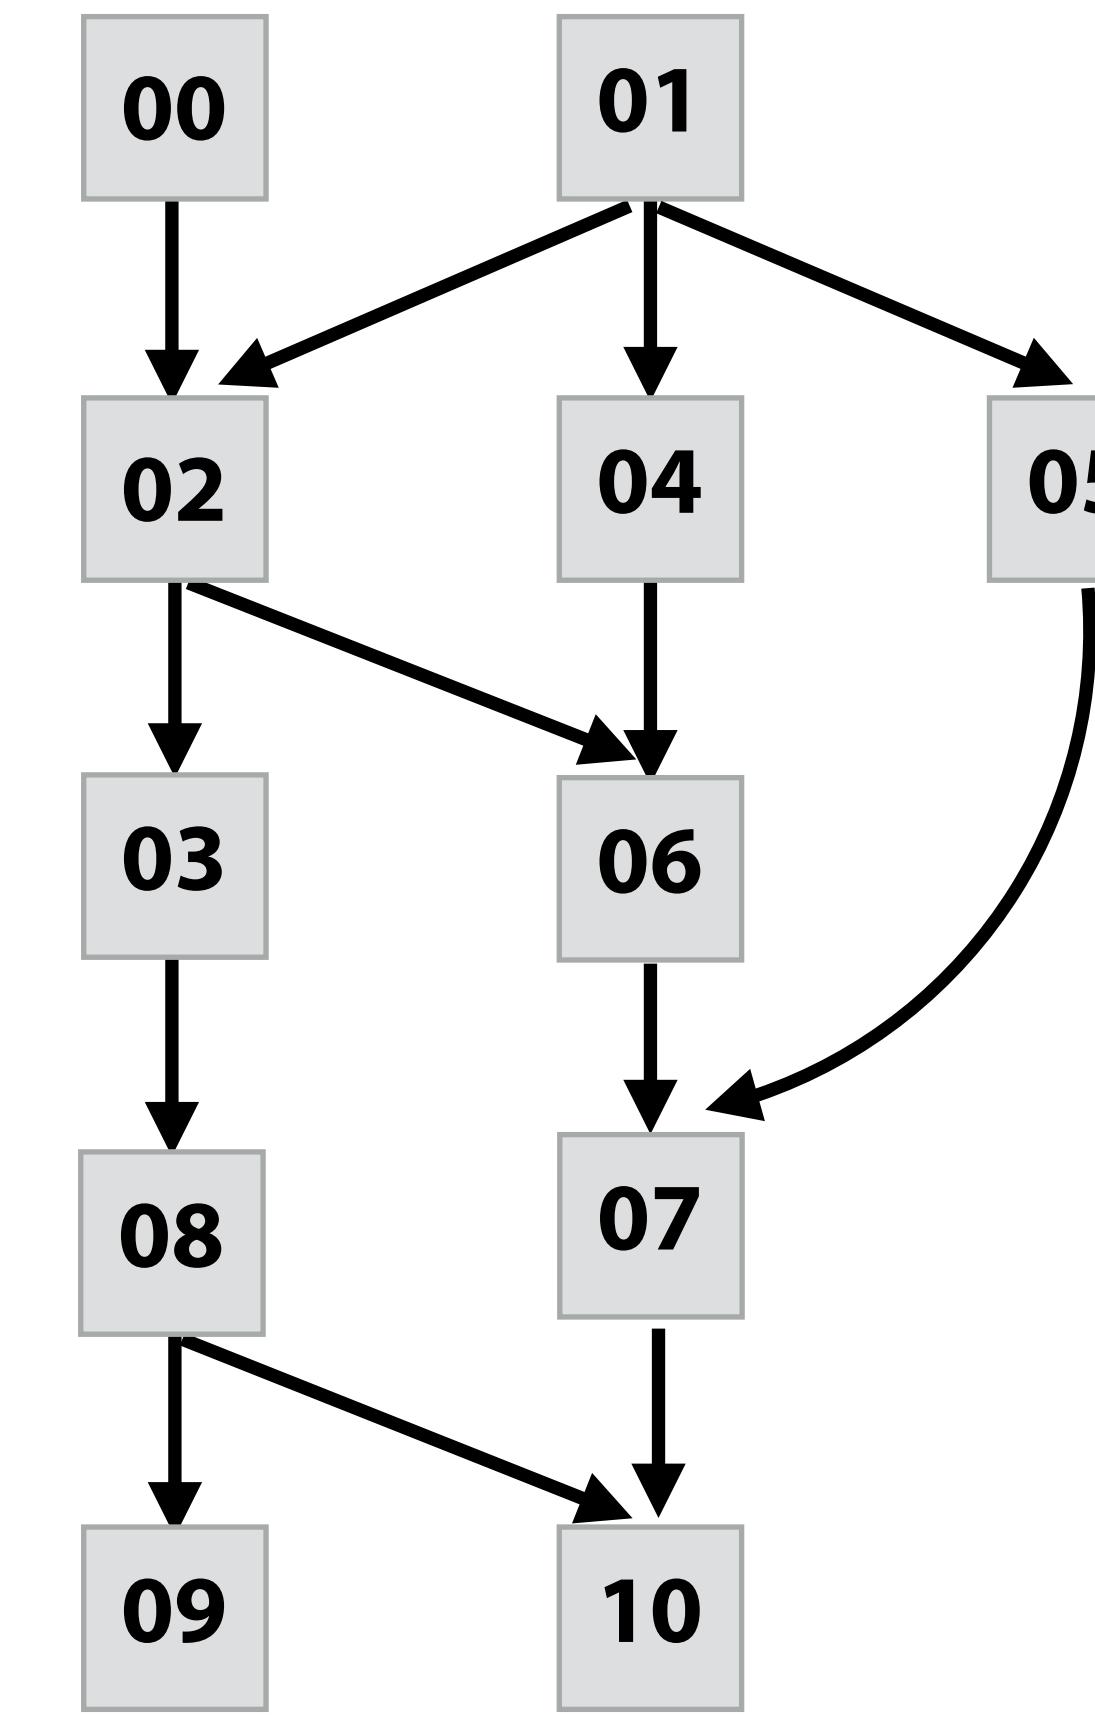


A more complex example

Program (sequence of instructions)

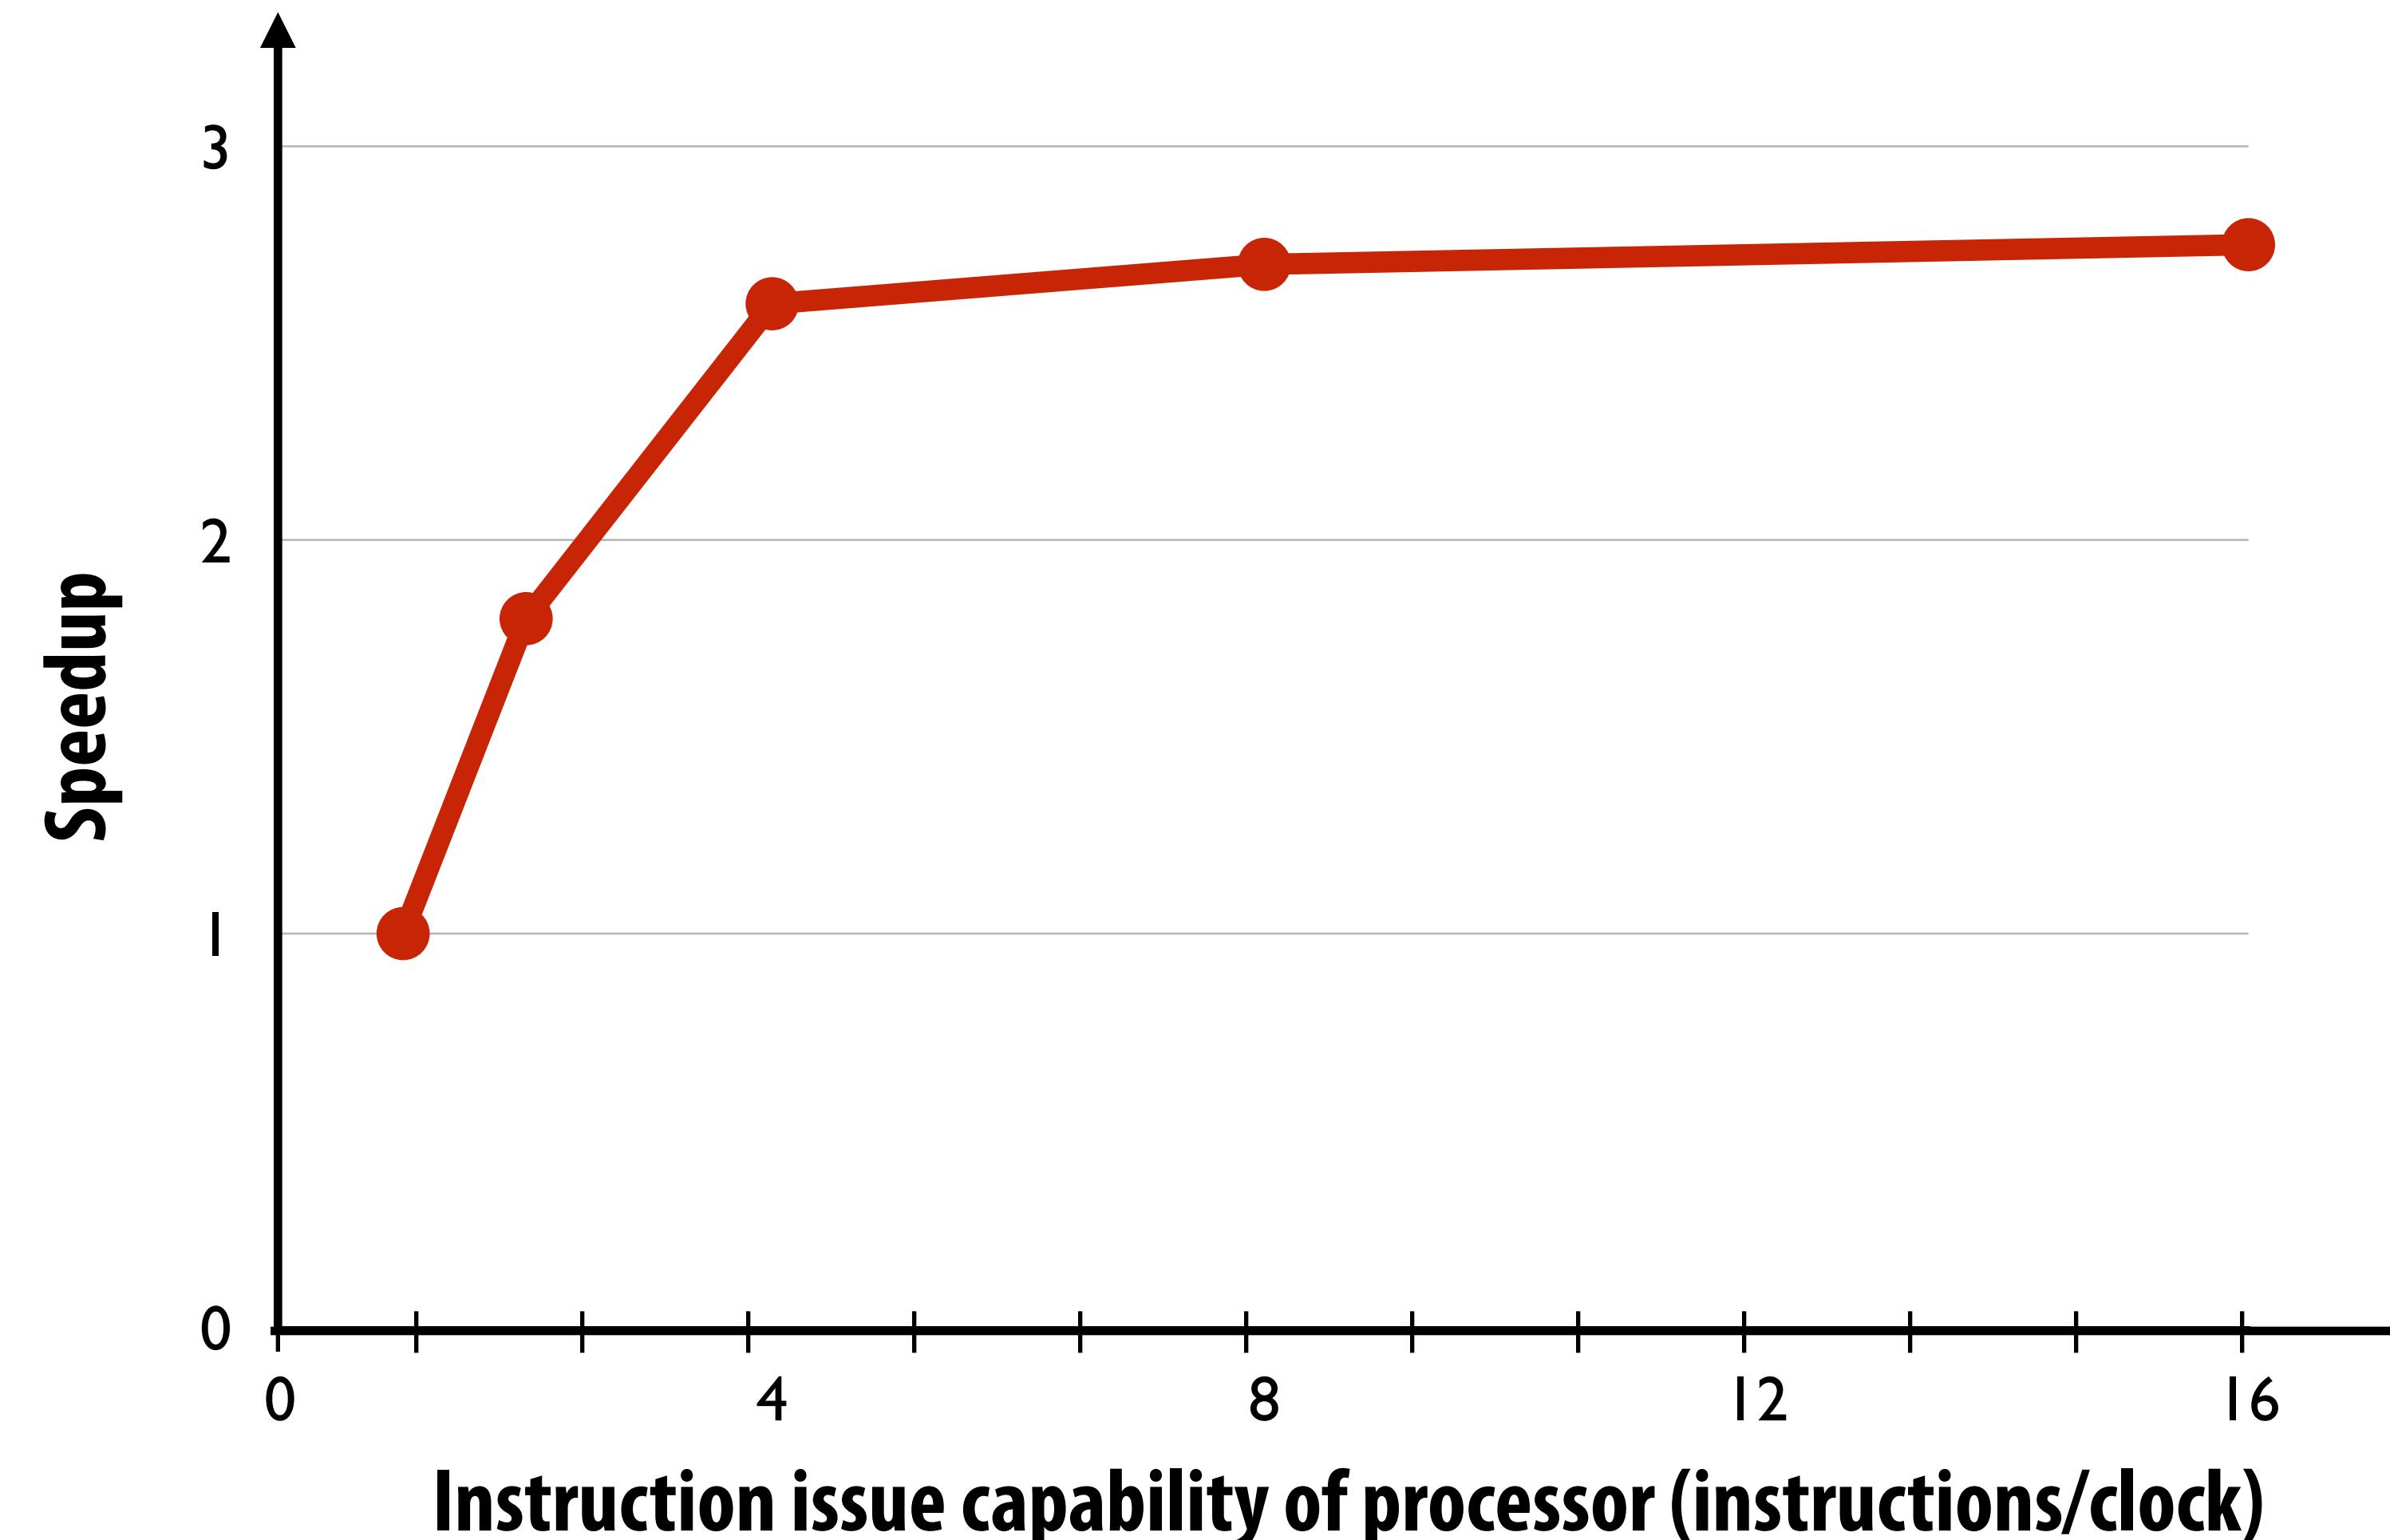
PC	Instruction	Computed value
00	a = 2	
01	b = 4	
02	tmp2 = a + b // 6	
03	tmp3 = tmp2 + a // 8	
04	tmp4 = b + b // 8	
05	tmp5 = b * b // 16	
06	tmp6 = tmp2 + tmp4 // 14	
07	tmp7 = tmp5 + tmp6 // 30	
08	if (tmp3 > 7)	
09	print tmp3	
else		
10	print tmp7	

Instruction dependency graph



Diminishing returns of superscalar execution

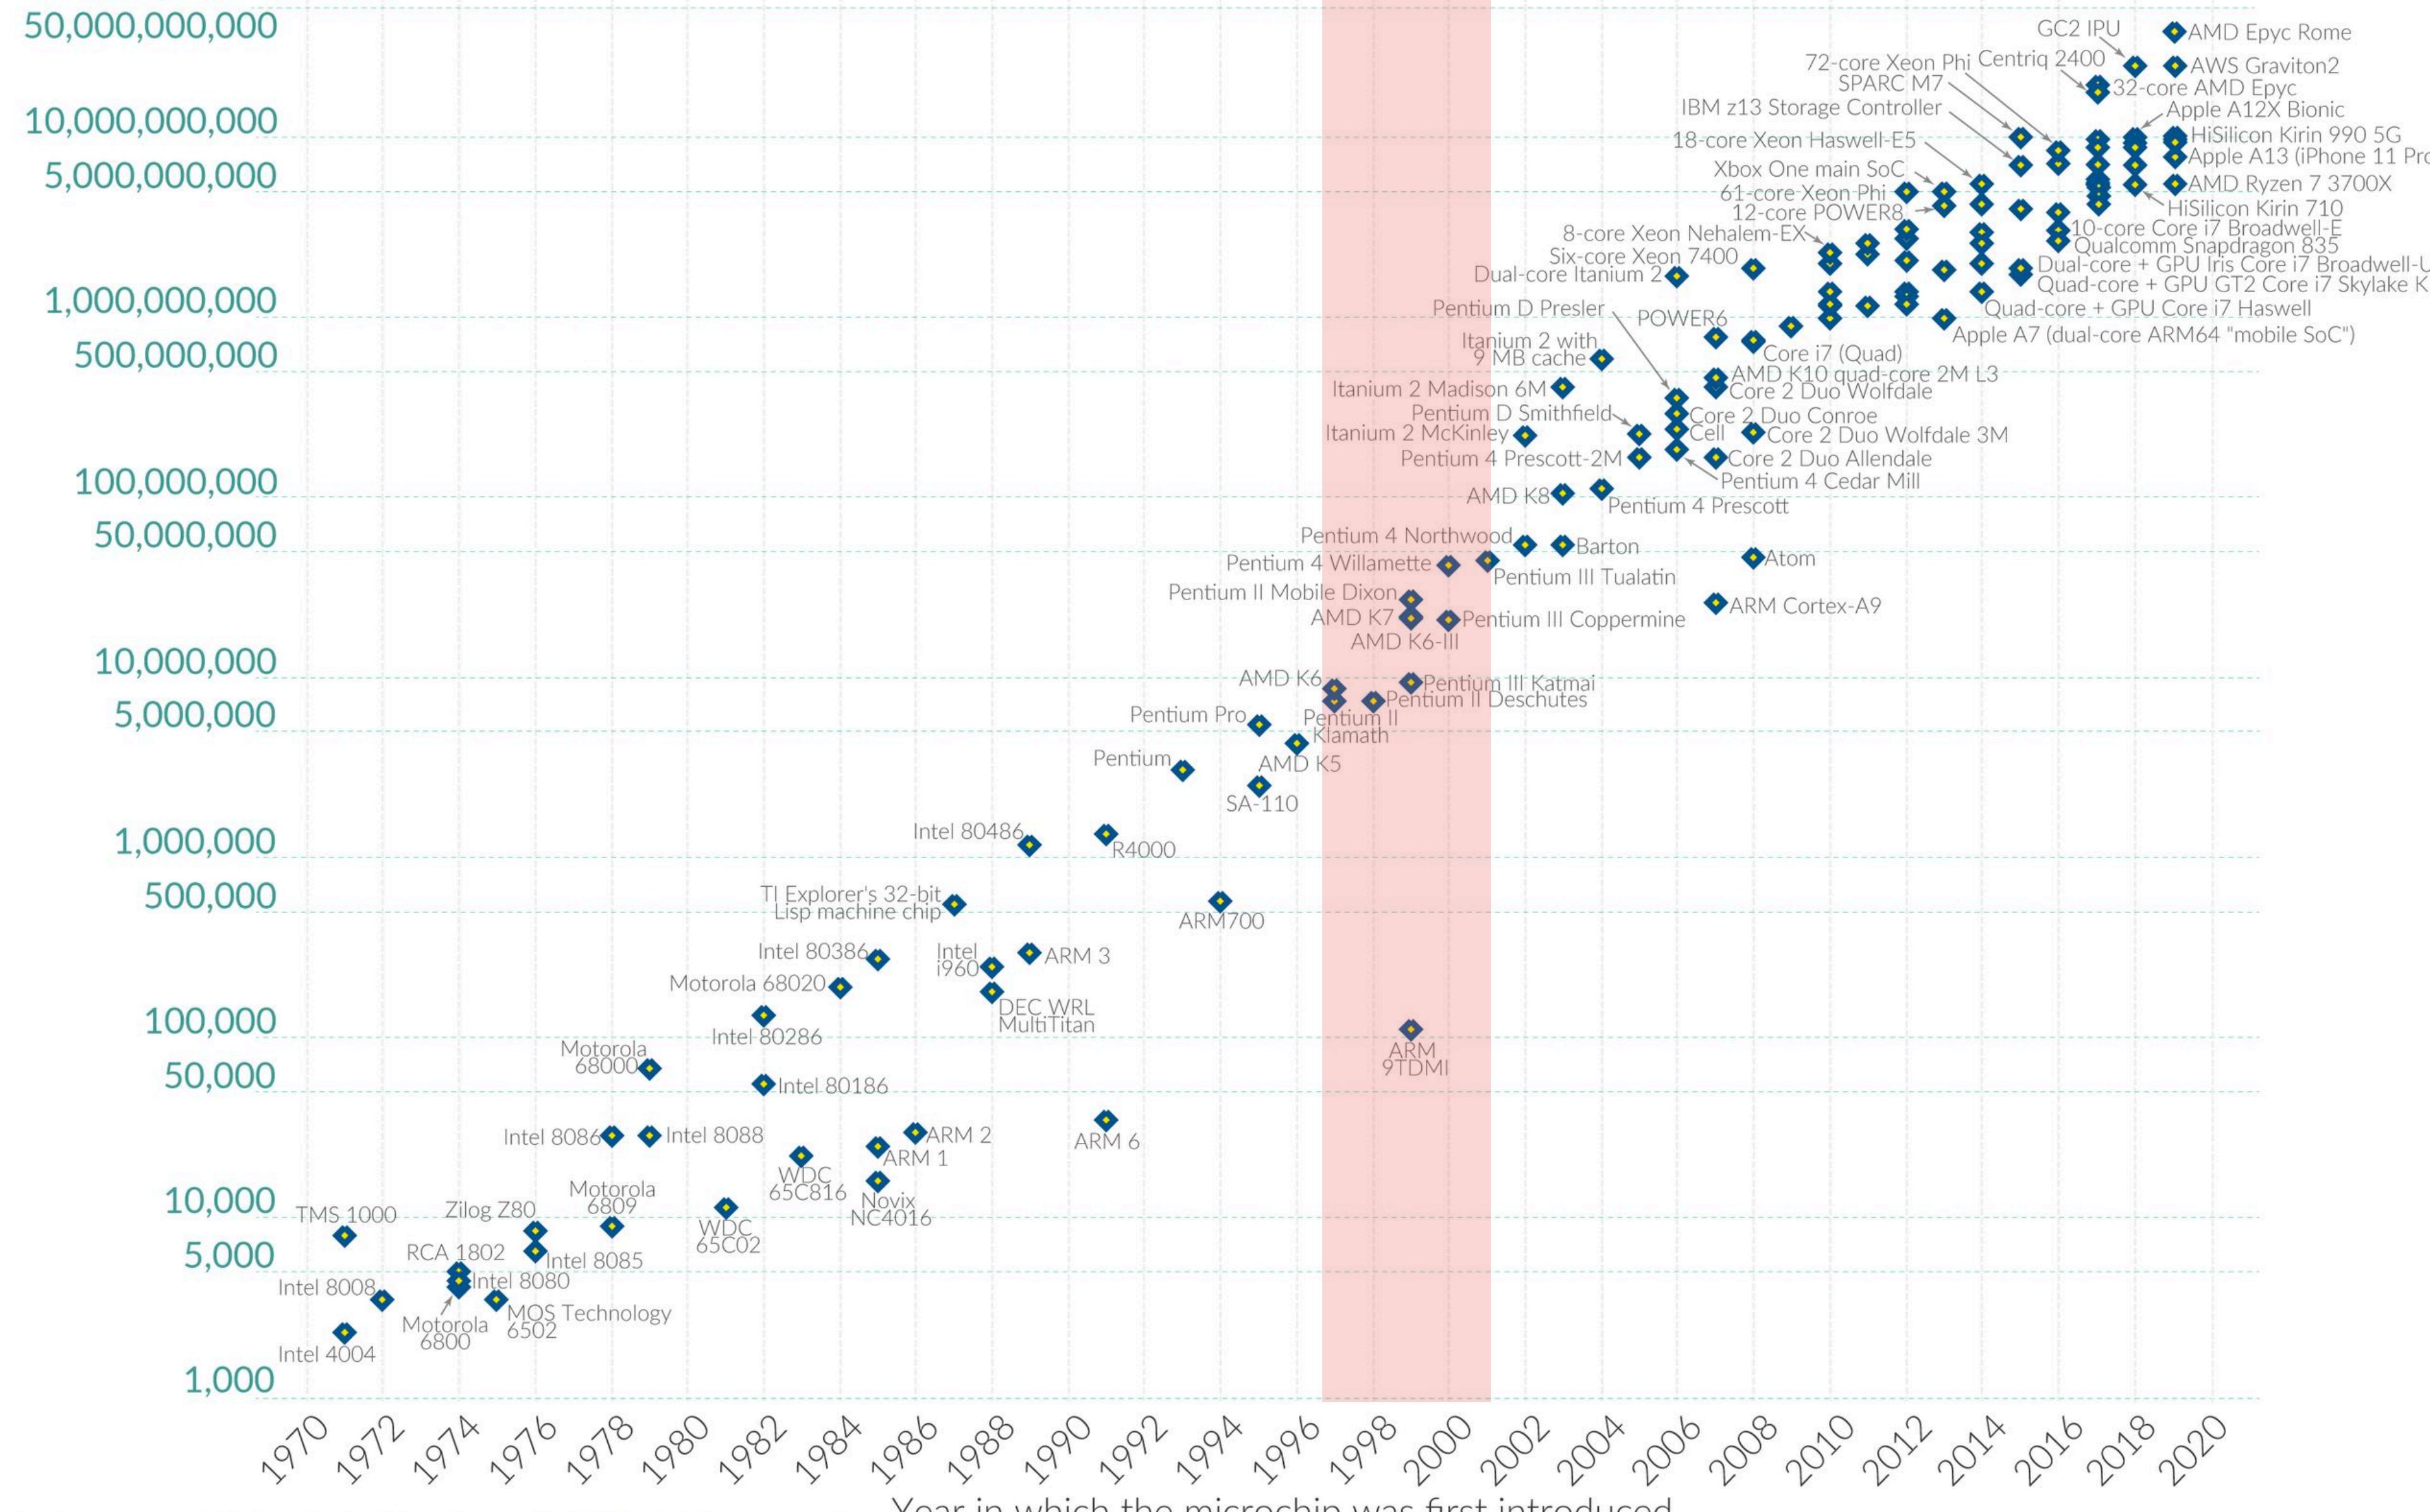
Most available ILP is exploited by a processor capable of issuing four instructions per clock
(Little performance benefit from building a processor that can issue more)



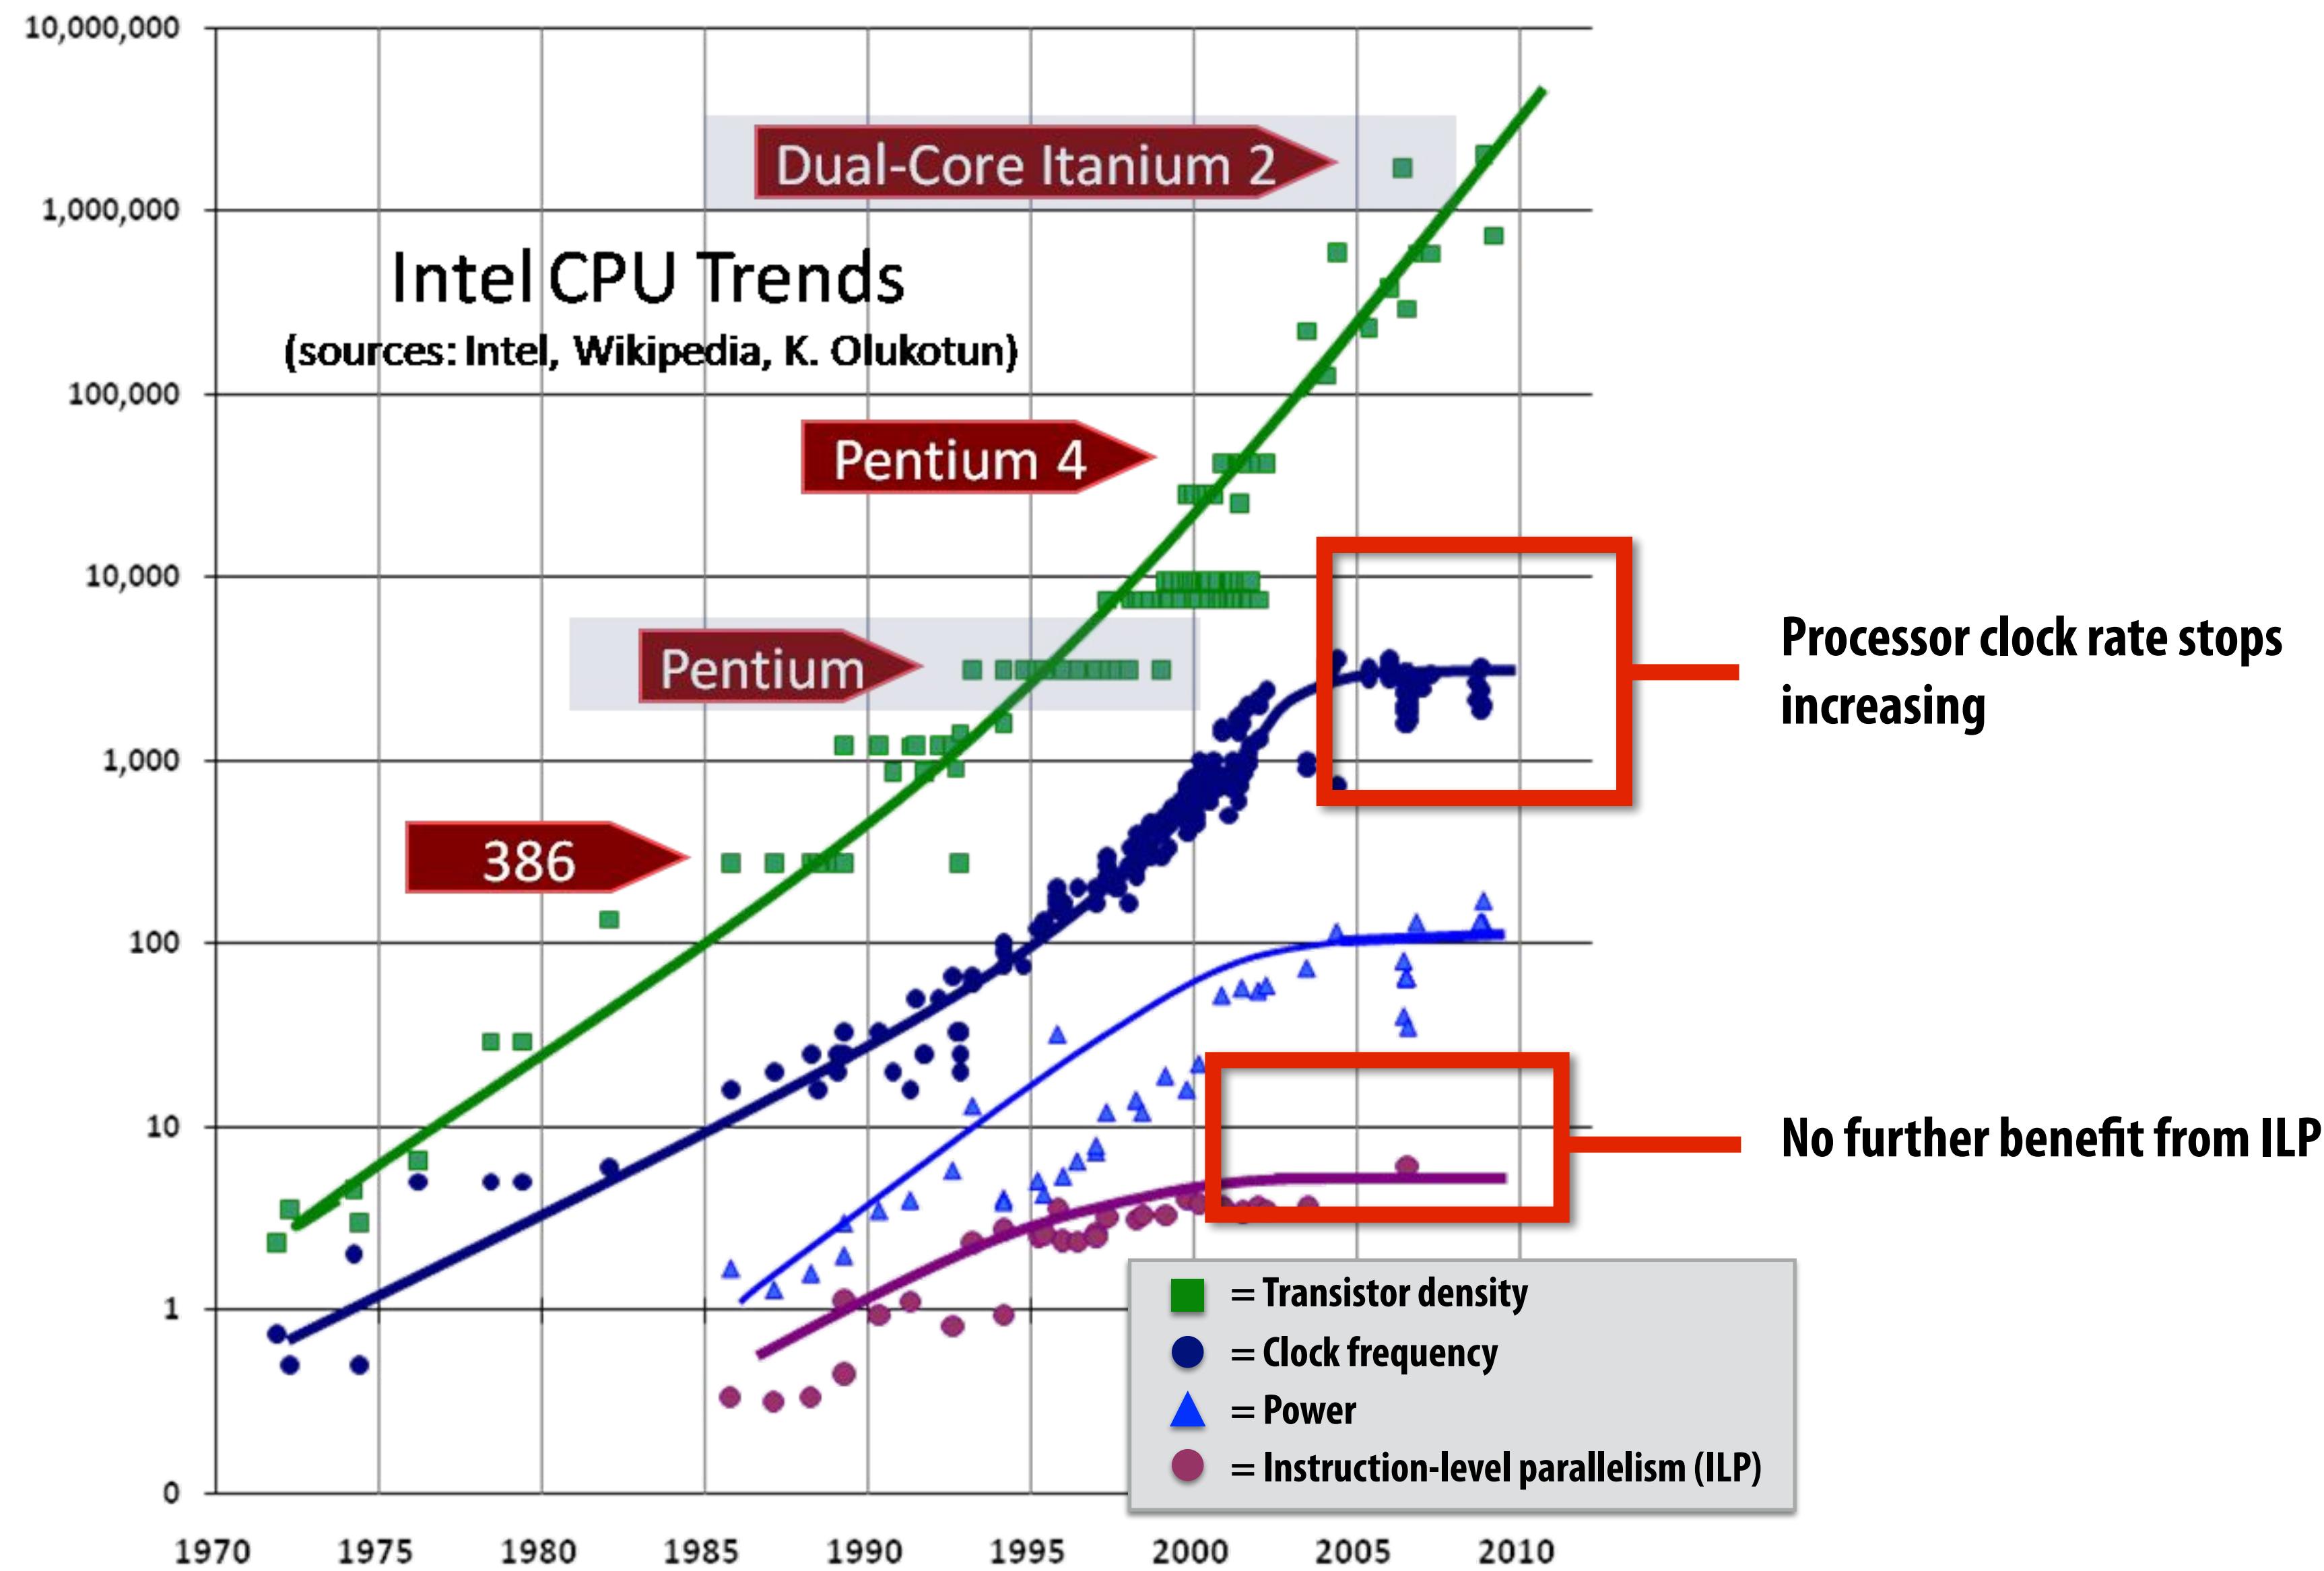
Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Transistor count



ILP tapped out + end of frequency scaling



The “power wall”

Power consumed by a transistor:

Dynamic power \propto capacitive load \times voltage² \times frequency

Static power: transistors burn power even when inactive due to leakage

High power = high heat

Power is a critical design constraint in modern processors

TDP

Apple M1 laptop:

13W

Intel Core i9 10900K (in desktop CPU):

95W

NVIDIA RTX 4090 GPU

450W

Mobile phone processor

1/2 - 2W

World's fastest supercomputer

megawatts

Standard microwave oven

900W

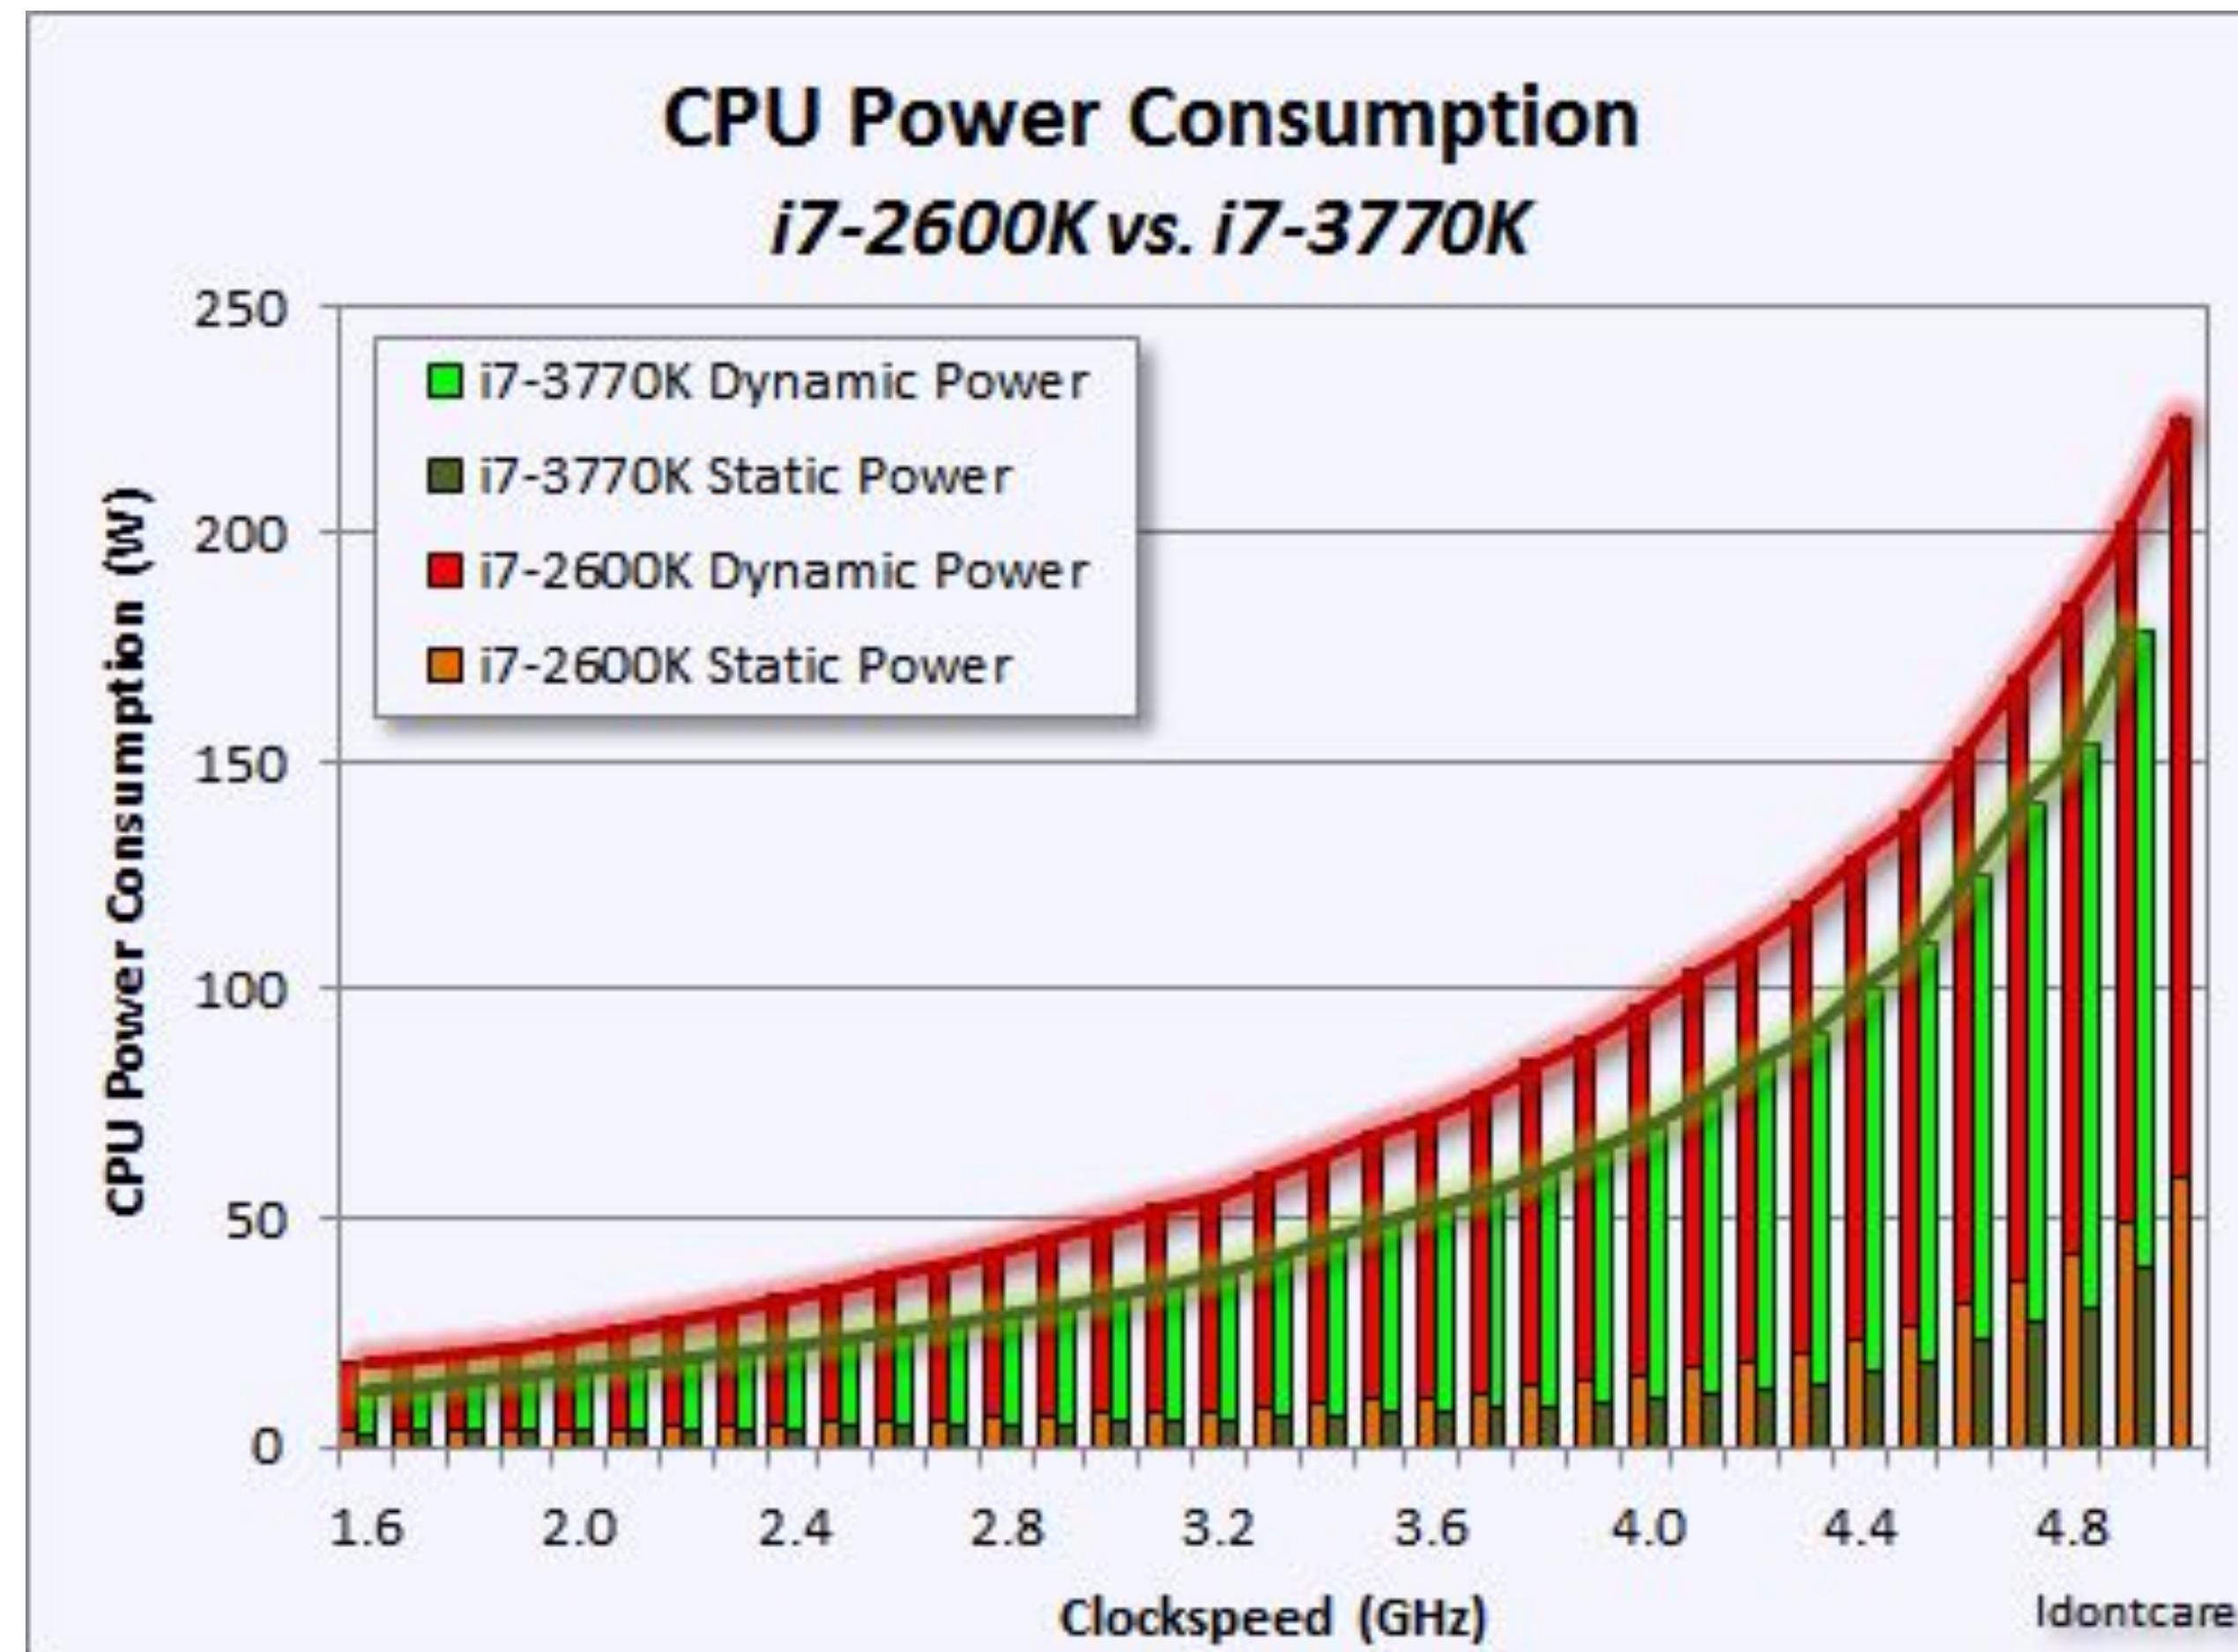


Power draw as a function of clock frequency

Dynamic power \propto capacitive load \times voltage² \times frequency

Static power: transistors burn power even when inactive due to leakage

Maximum allowed frequency determined by processor's core voltage



Single-core performance scaling

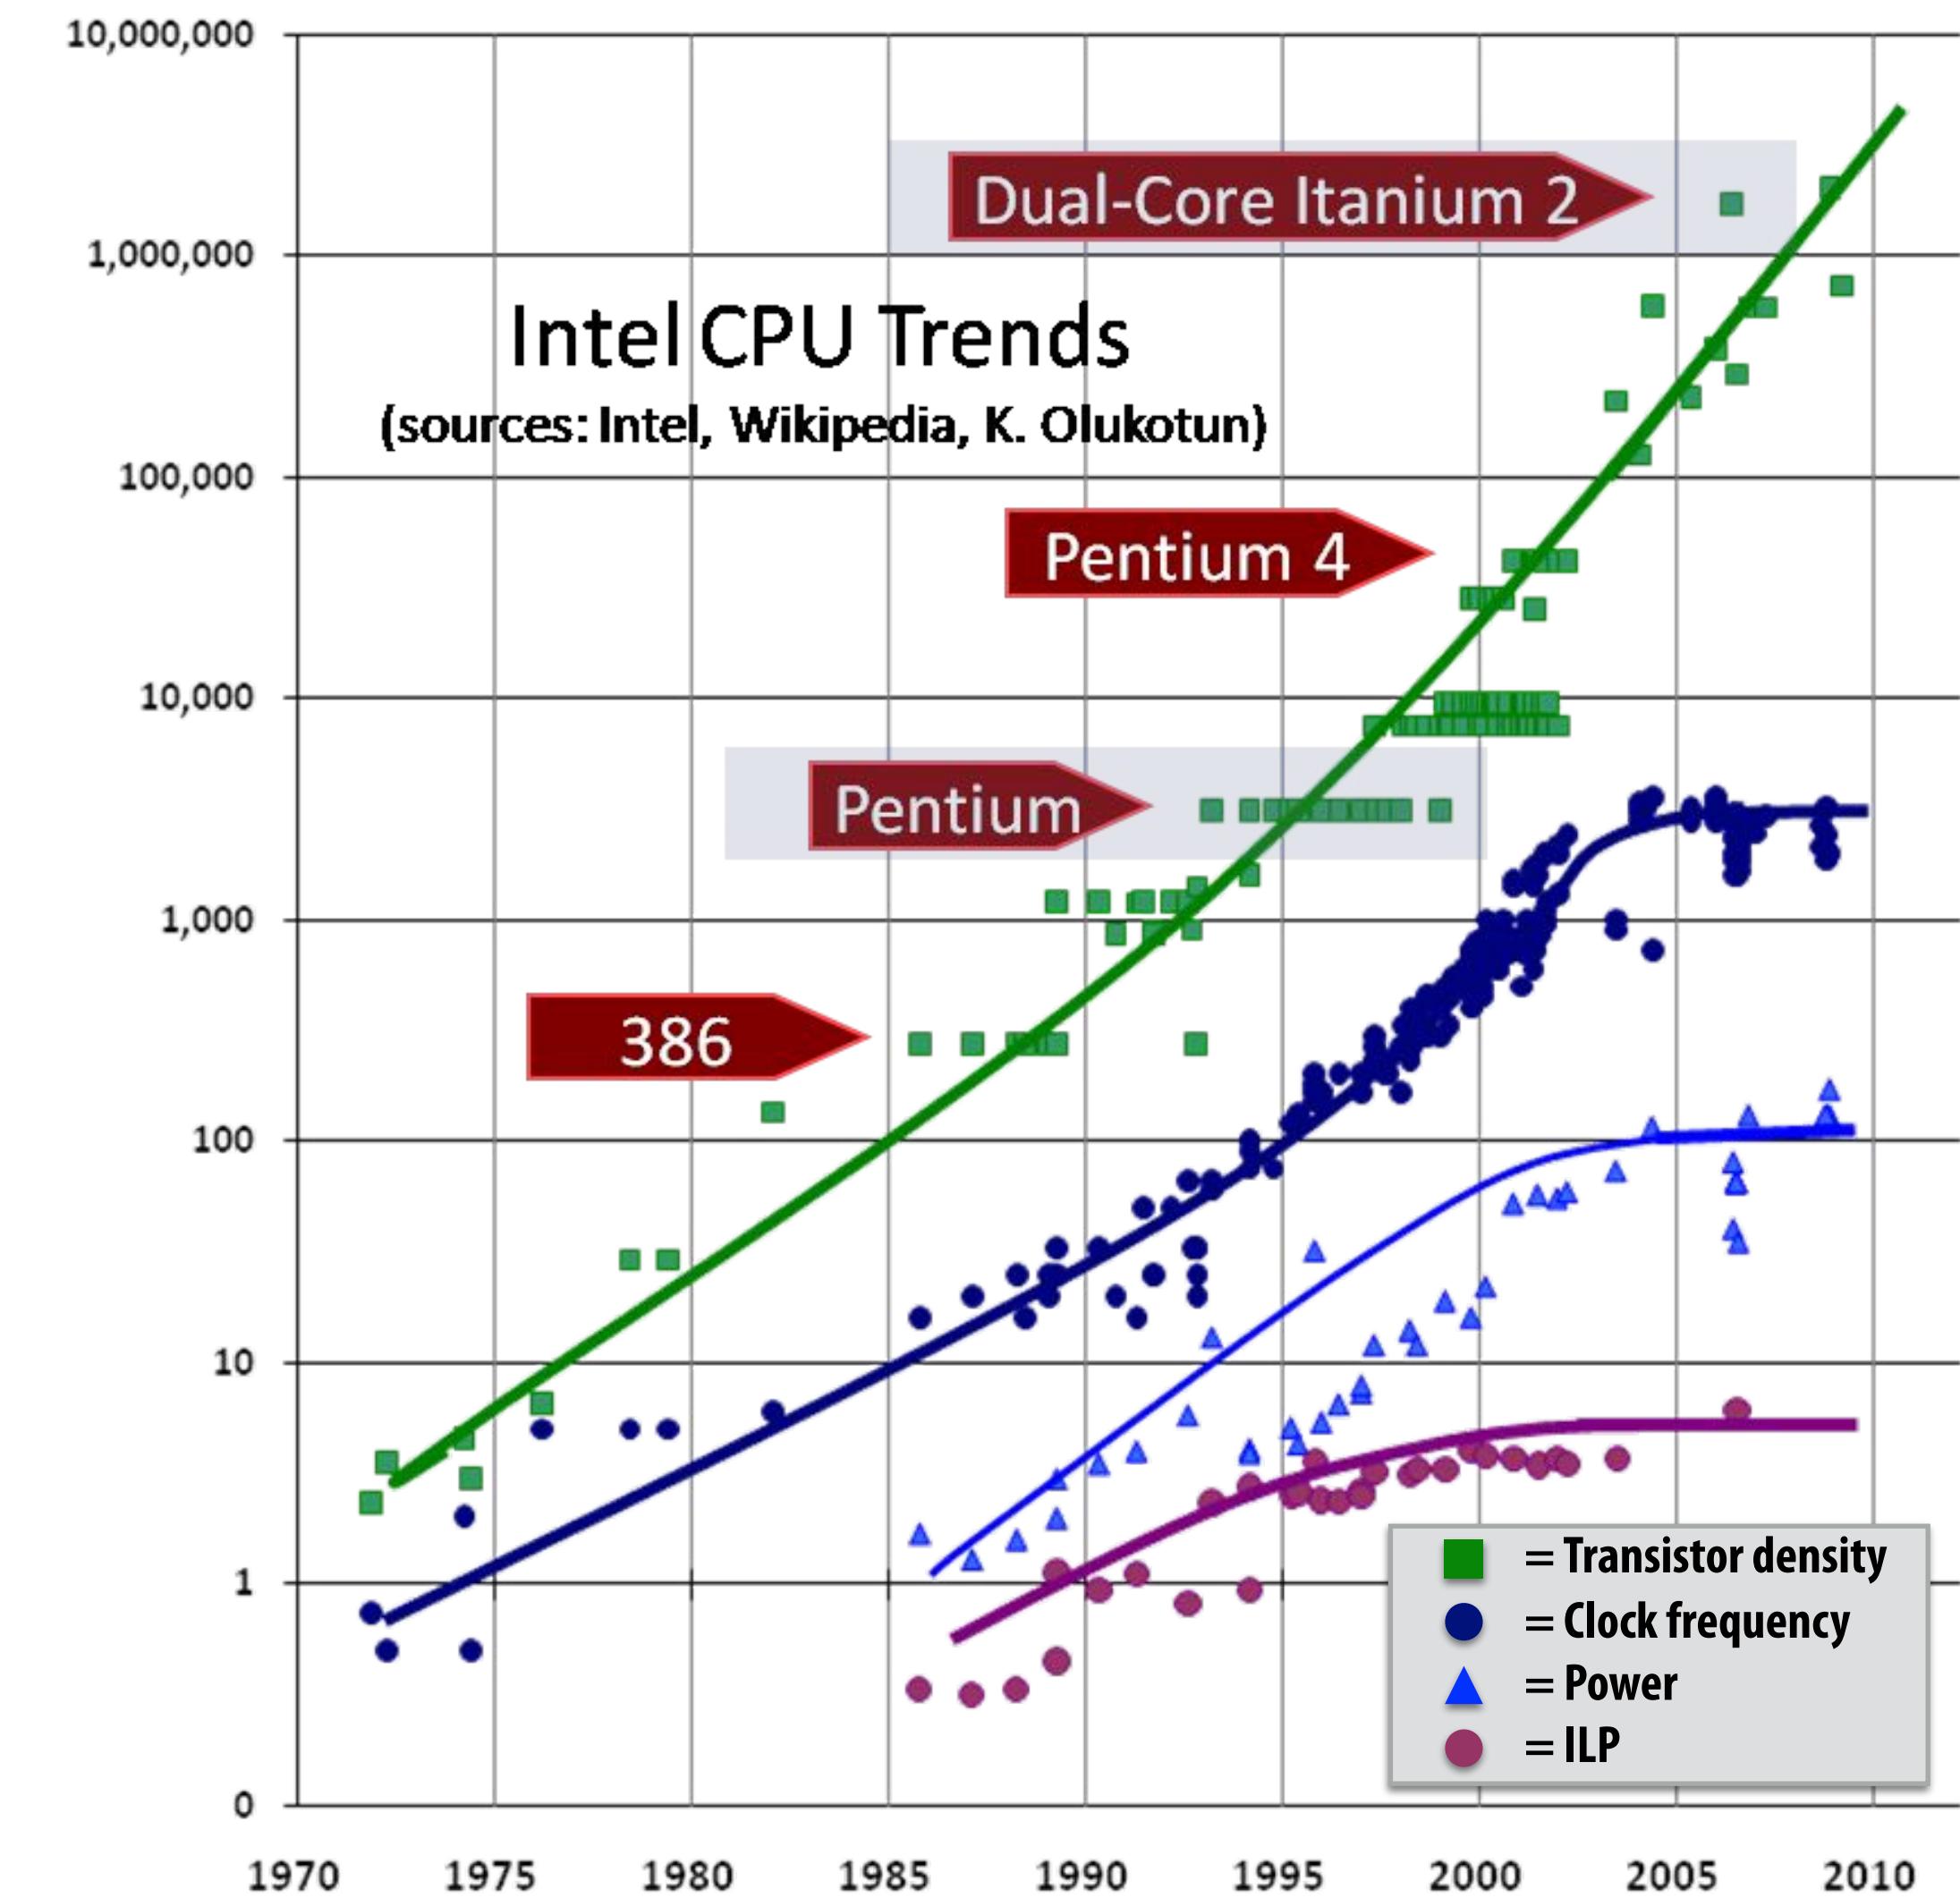
The rate of single-instruction stream performance scaling has decreased (almost to zero)

1. Frequency scaling limited by power
2. ILP scaling tapped out

Architects are now building faster processors by adding more execution units that run in parallel

(Or units that are specialized for a specific task: like graphics, or audio/video playback)

Software must be written to be parallel to see performance gains. No more free lunch for software developers!



Example: multi-core CPU

Intel “Comet Lake” 10th Generation Core i9 10-core CPU (2020)

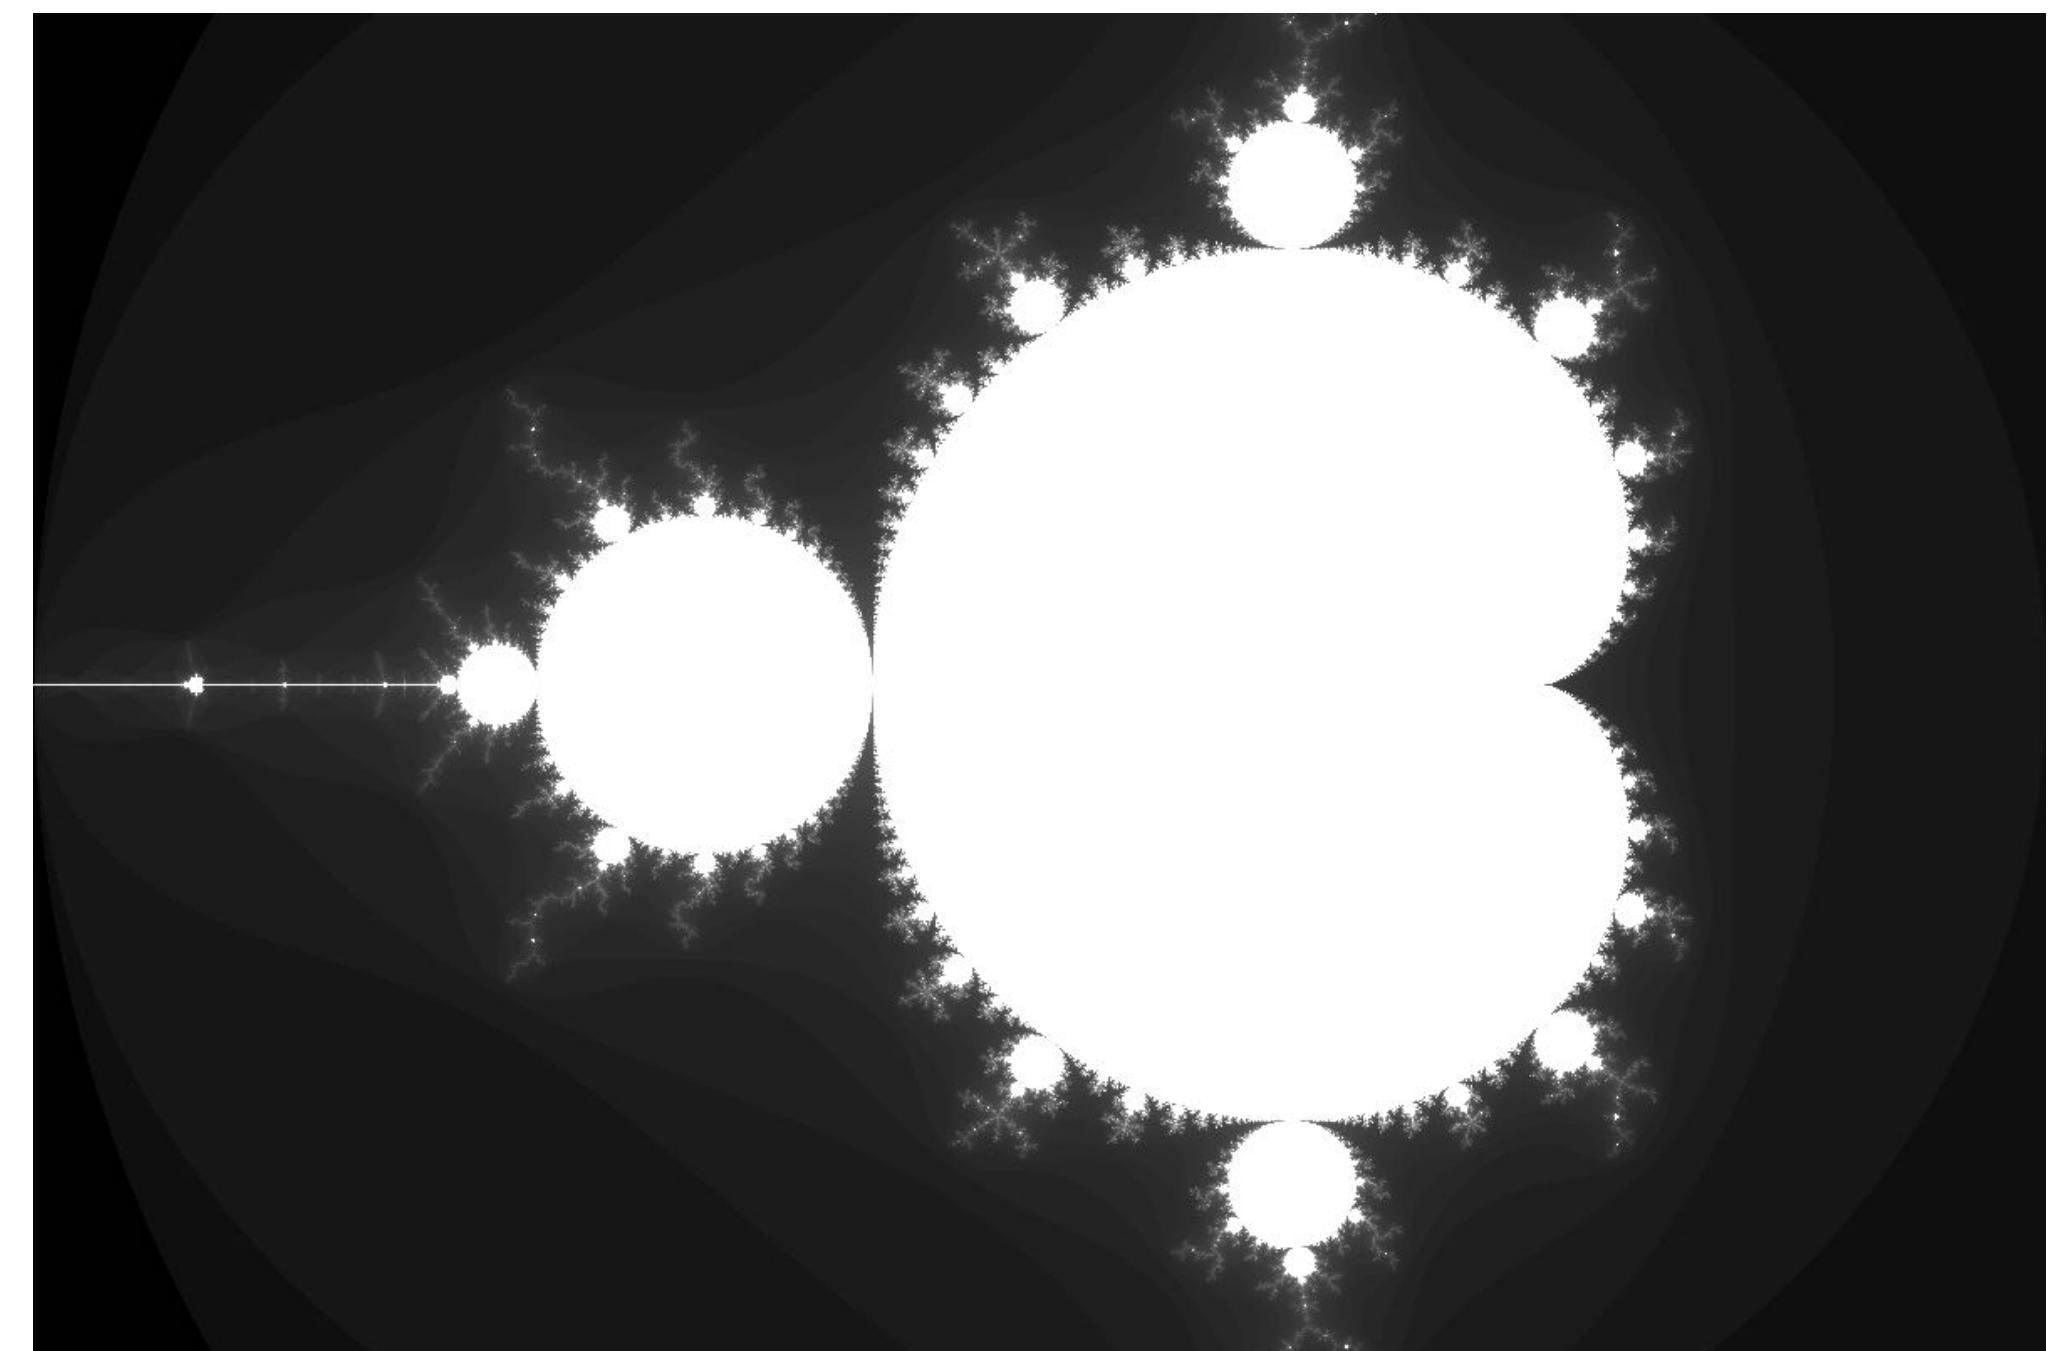


One thing you will learn in this course

- How to write code that efficiently uses the resources in a modern multi-core CPU
- Example: assignment 1 (coming up!)
 - Running on a quad-core Intel CPU
 - Four CPU cores
 - AVX SIMD vector instructions + hyper-threading
 - Baseline: single-threaded C program compiled with -O3
 - Parallelized program that uses all parallel execution resources on this CPU...

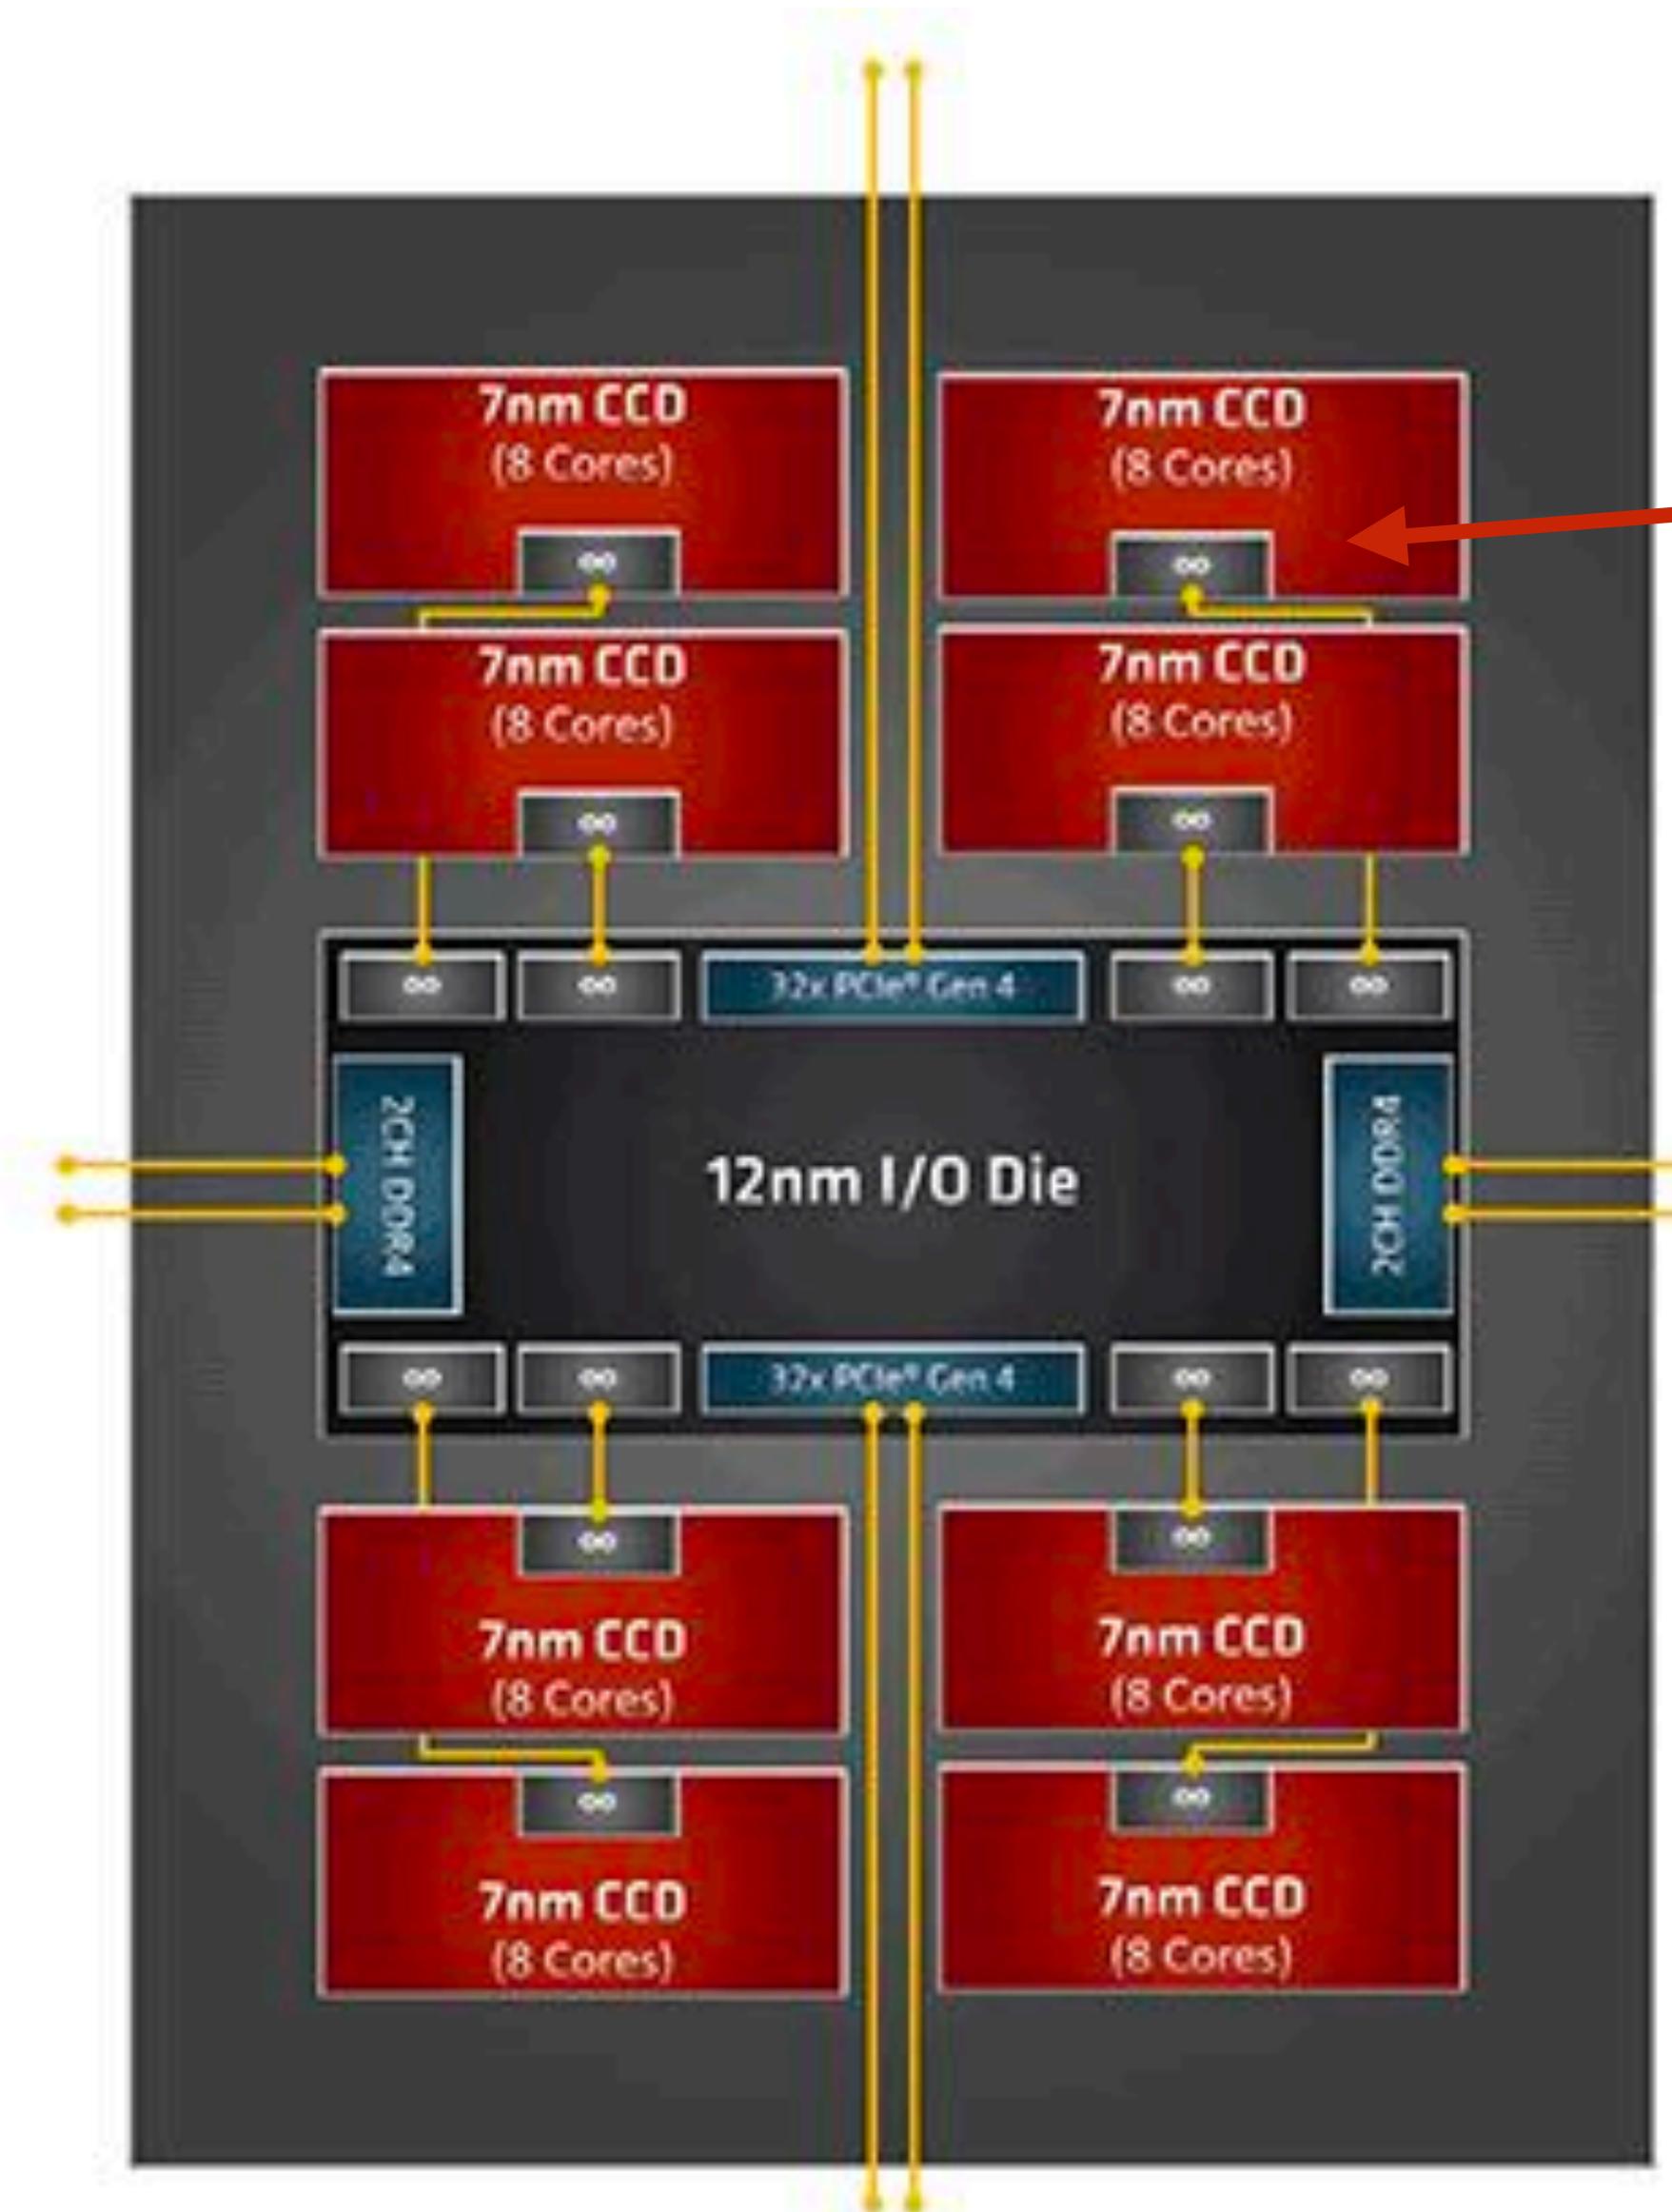
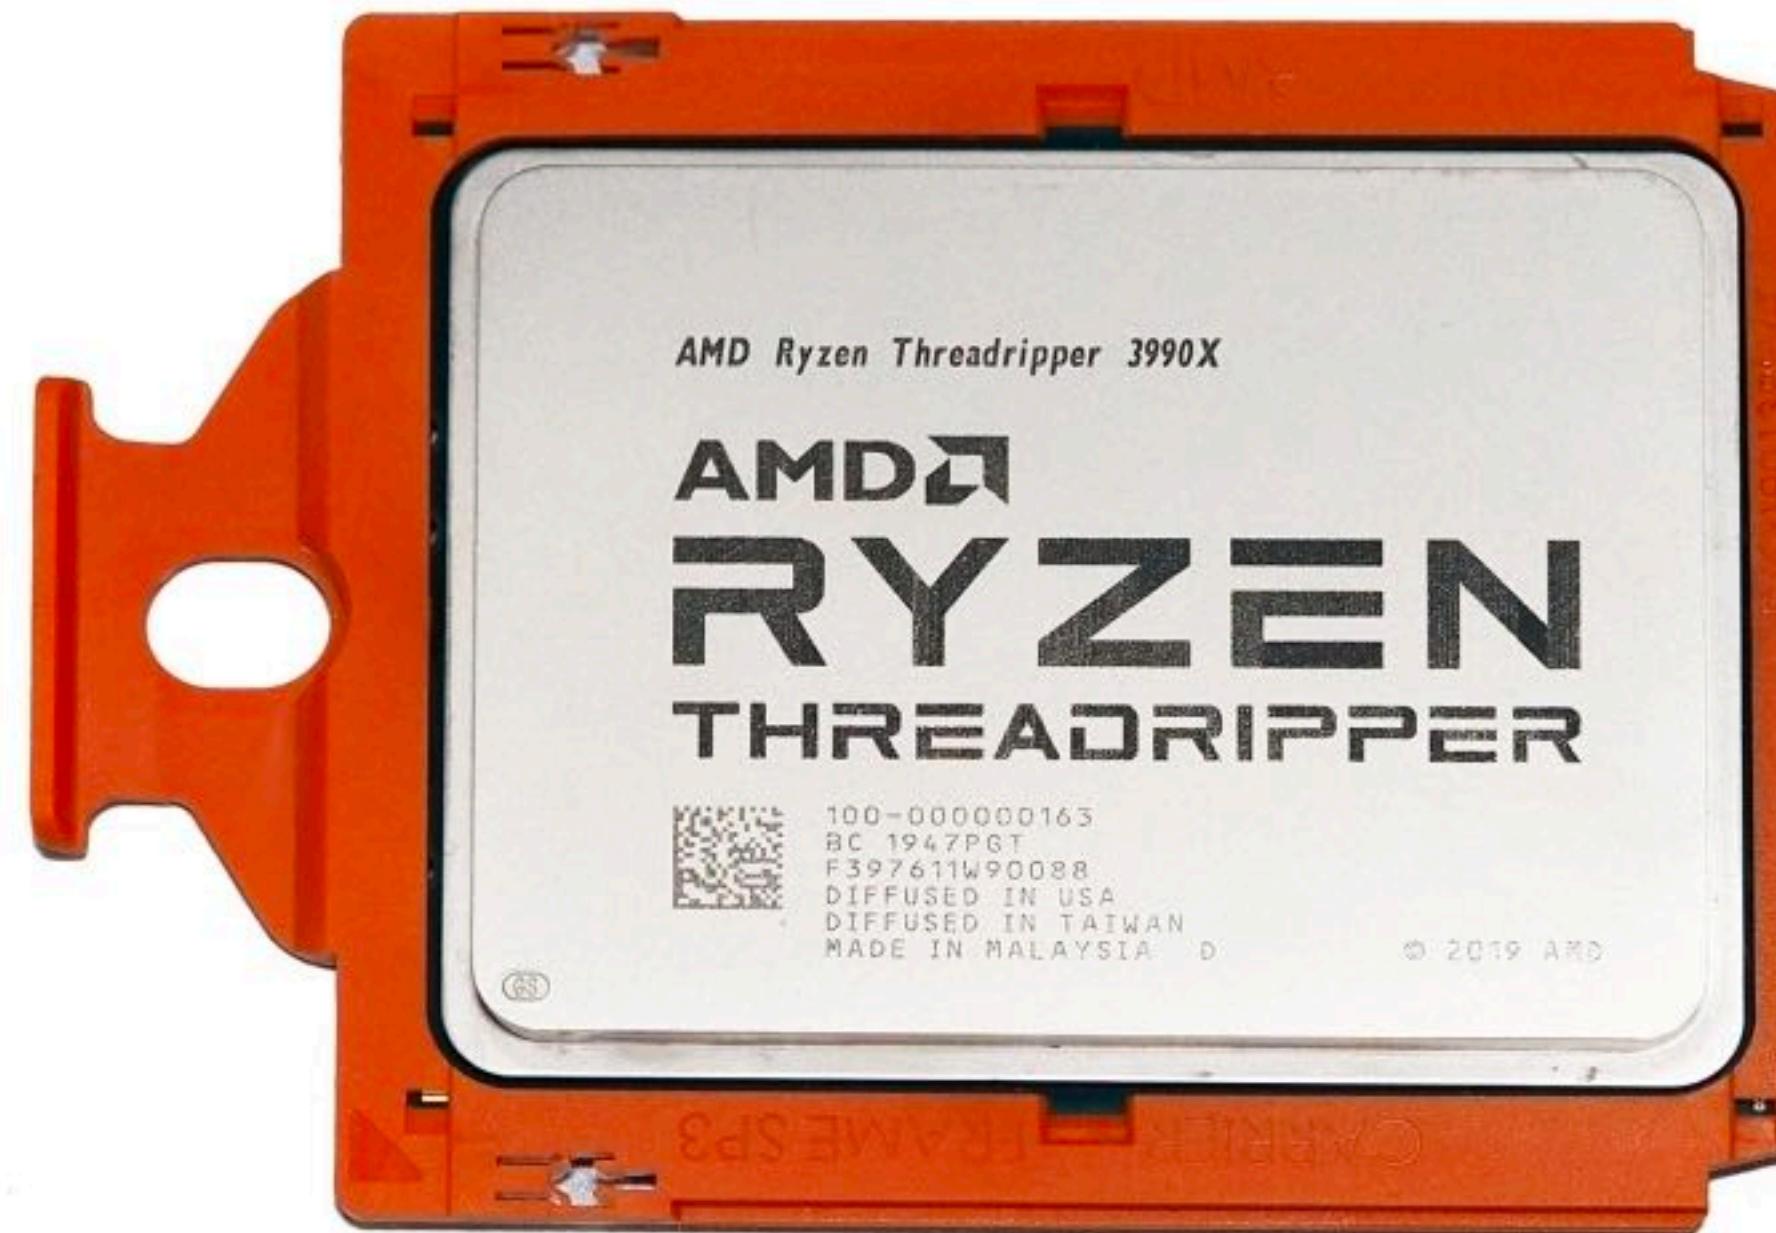
~32-40x faster!

We'll talk about these terms next time!



AMD Ryzen Threadripper 3990X

64 cores, 4.3 GHz



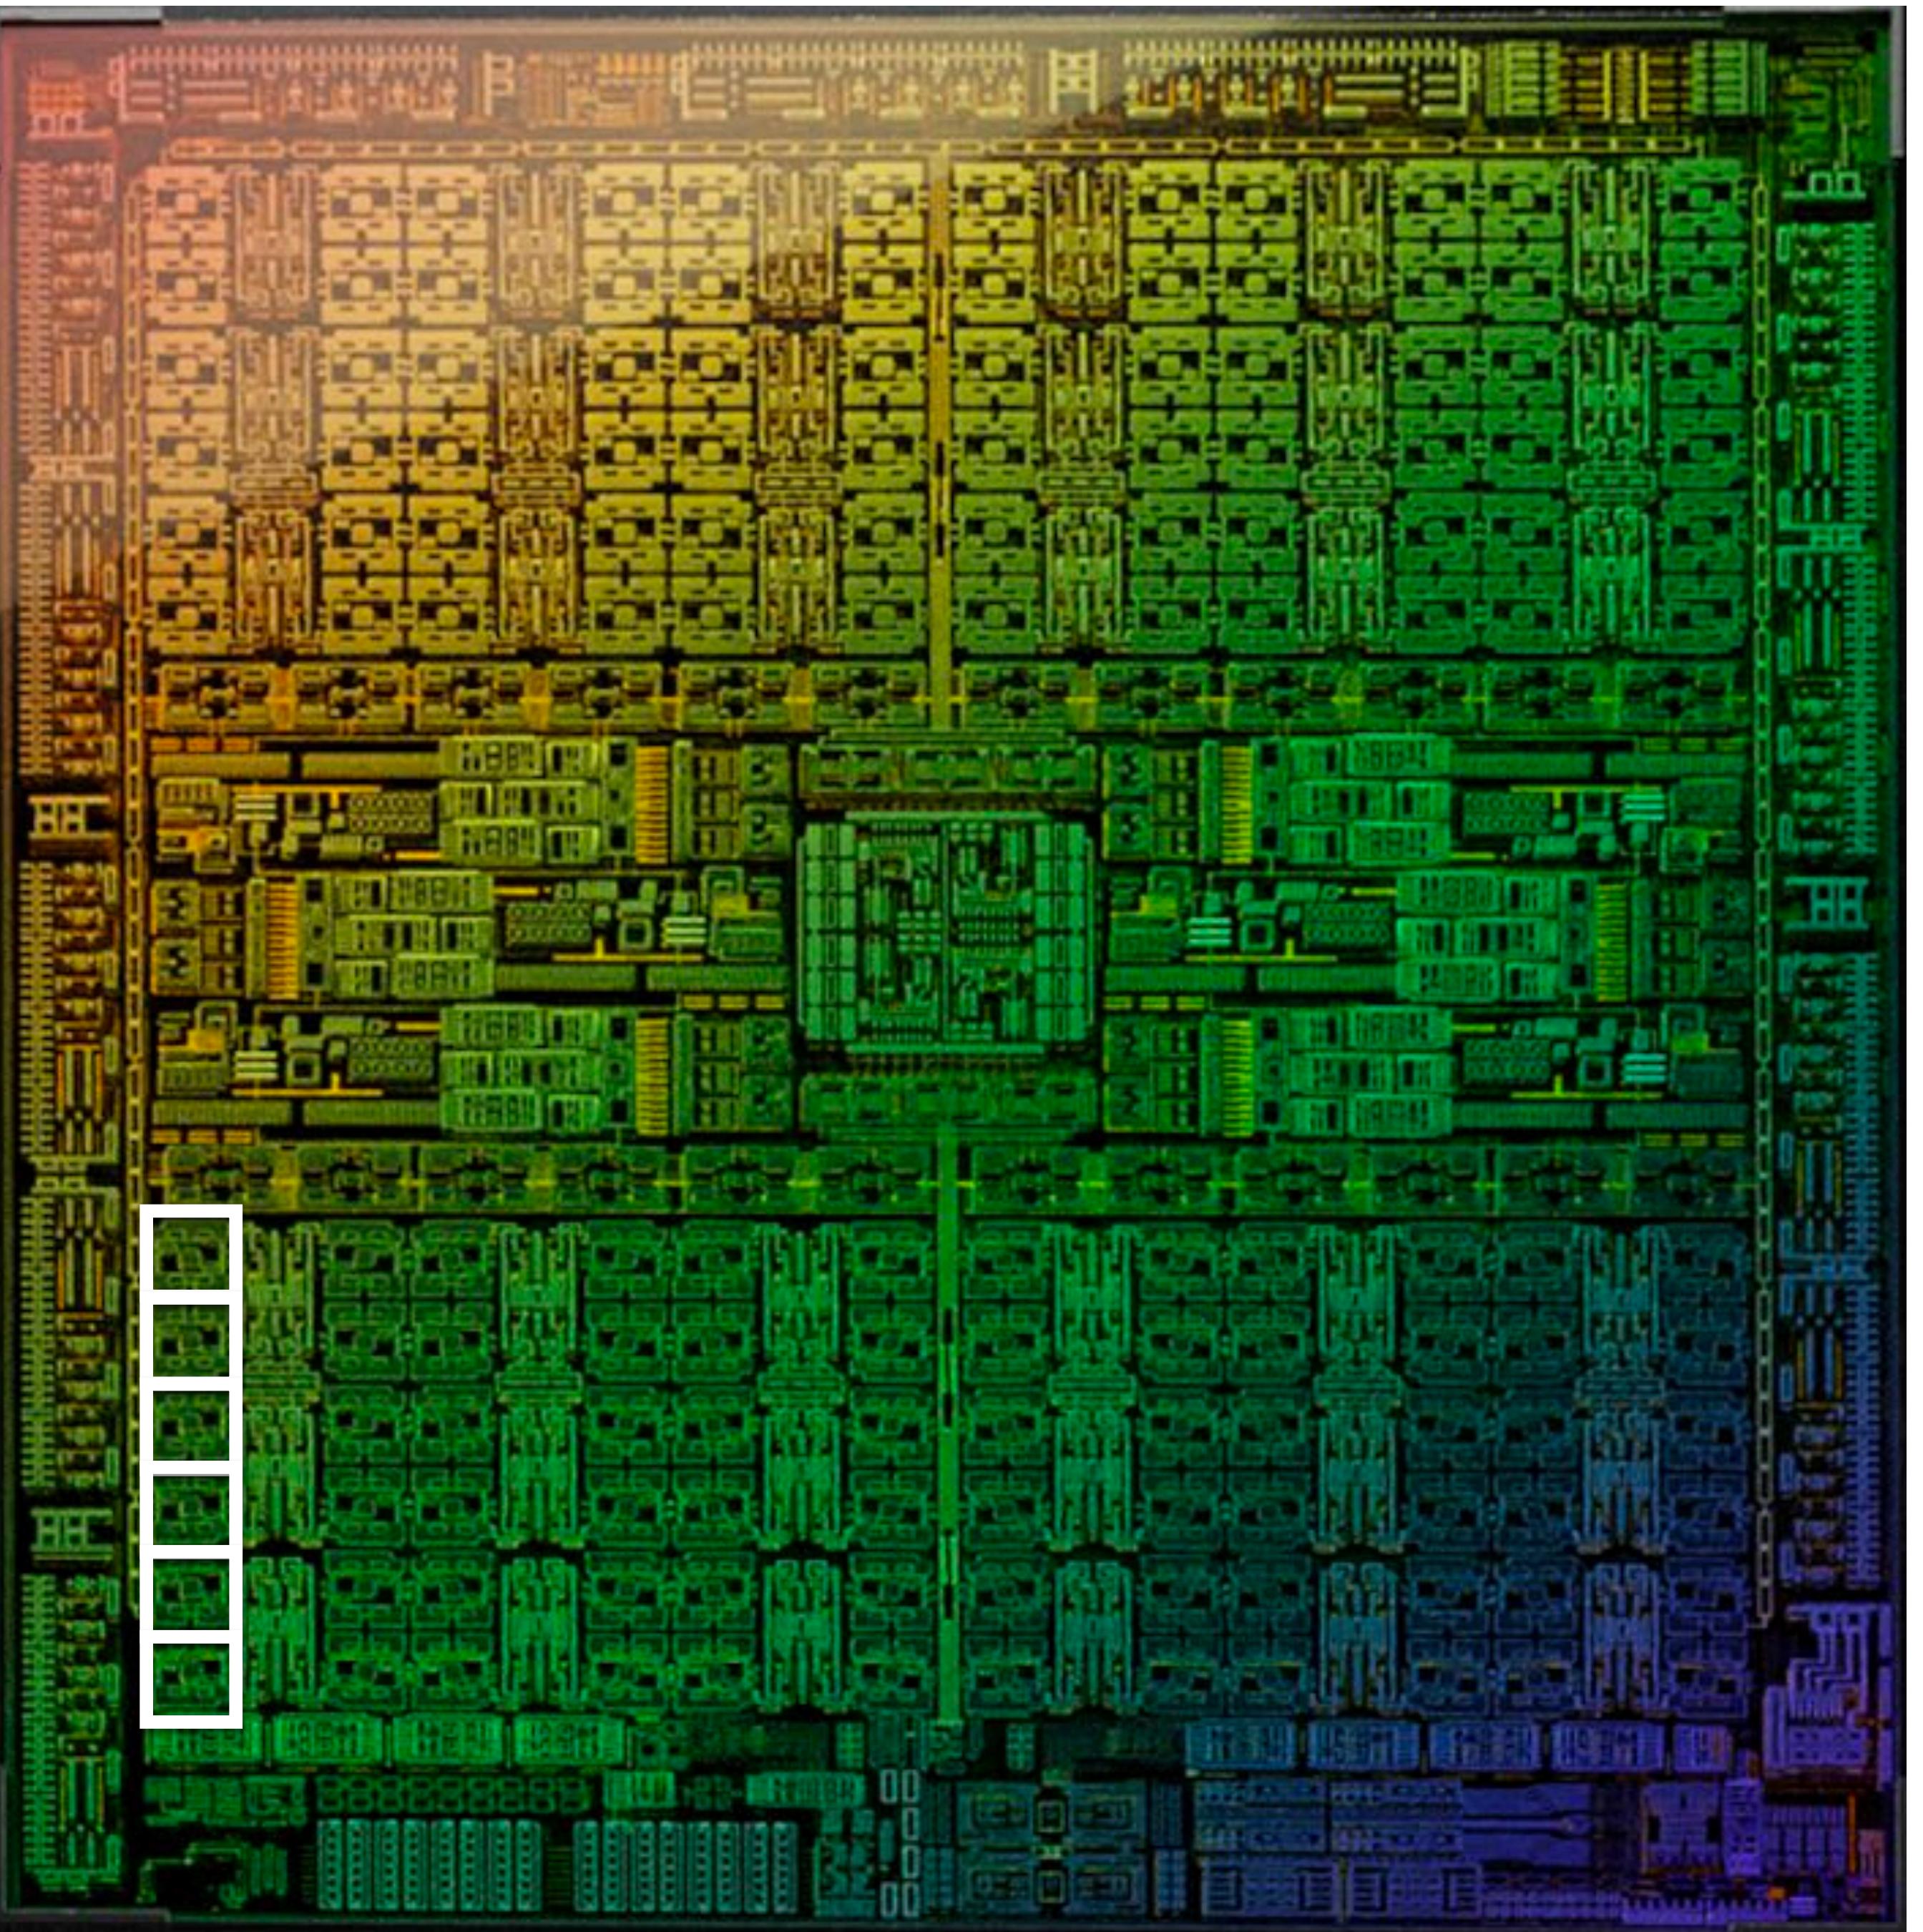
Four 8-core chiplets

NVIDIA AD102 GPU

GeForce RTX 4090 (2022)

76 billion transistors

18,432 fp32 multipliers organized in
144 processing blocks (called SMs)



GPU-accelerated supercomputing



Frontier (at Oak Ridge National Lab)
(world's #1 in Fall 2022)

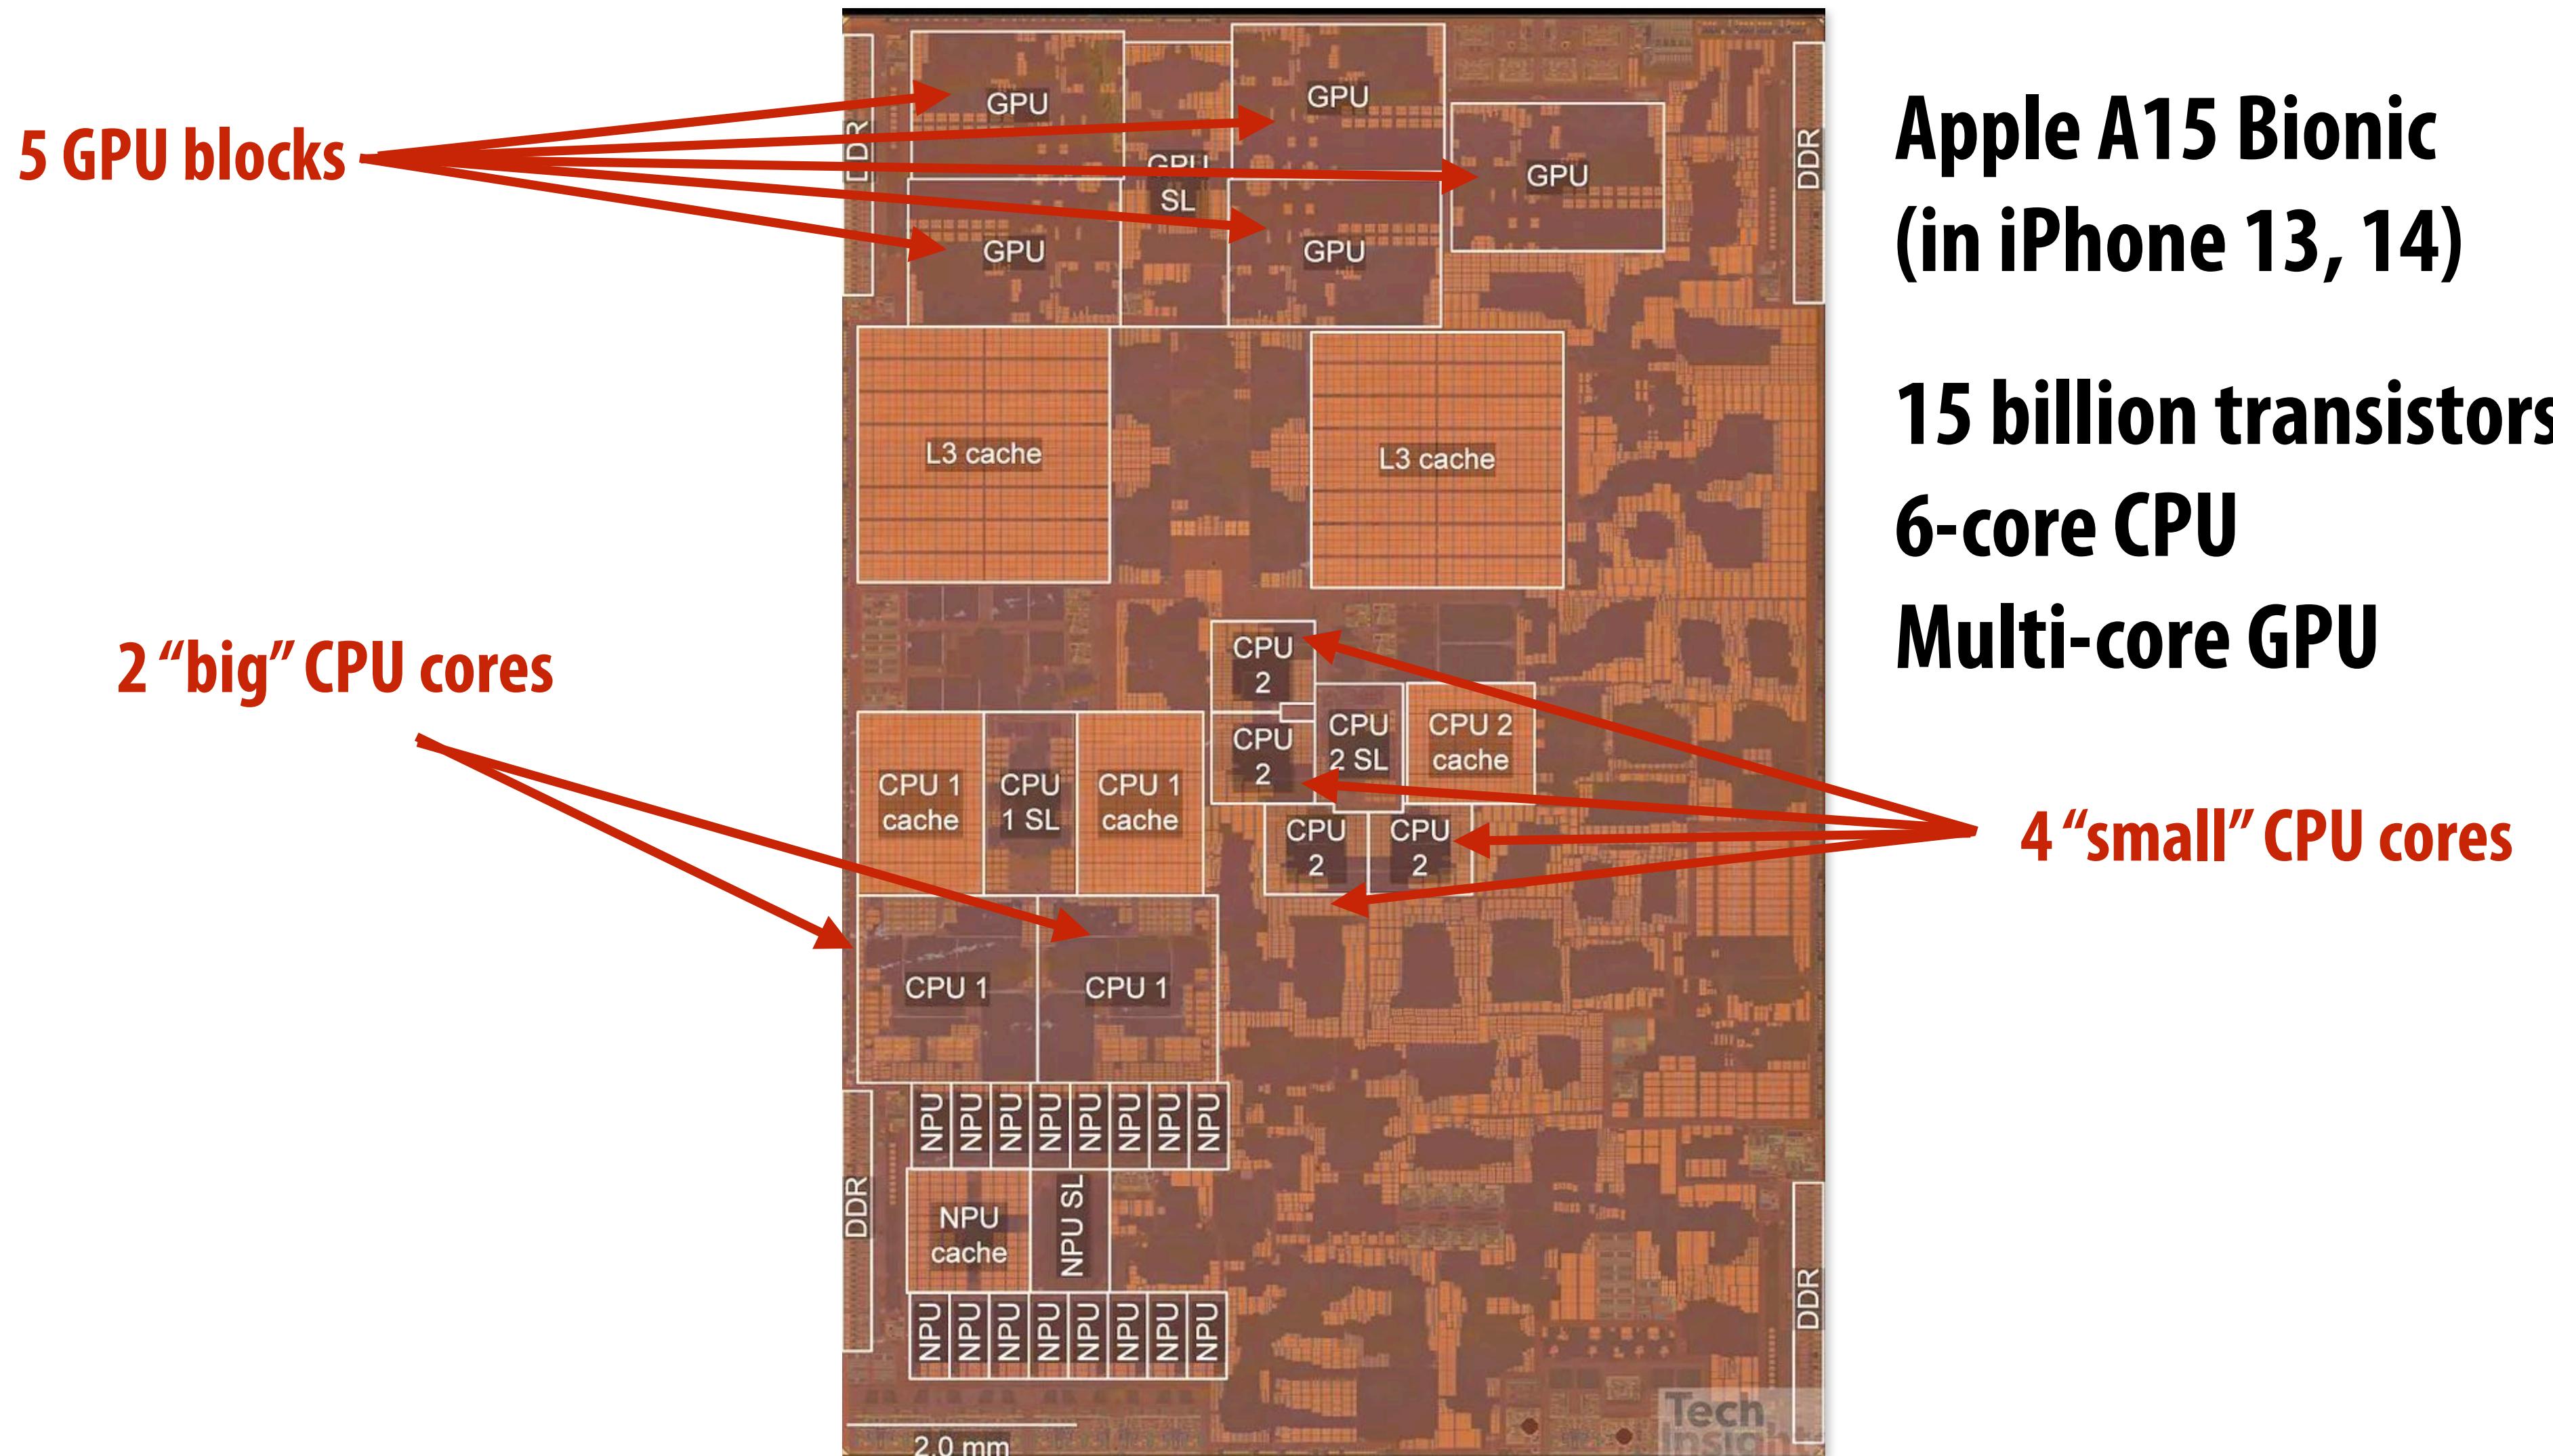
9472 x 64 core AMD CPUs (606,208 CPU cores)

37,888 Radeon GPUs

21 Megawatts

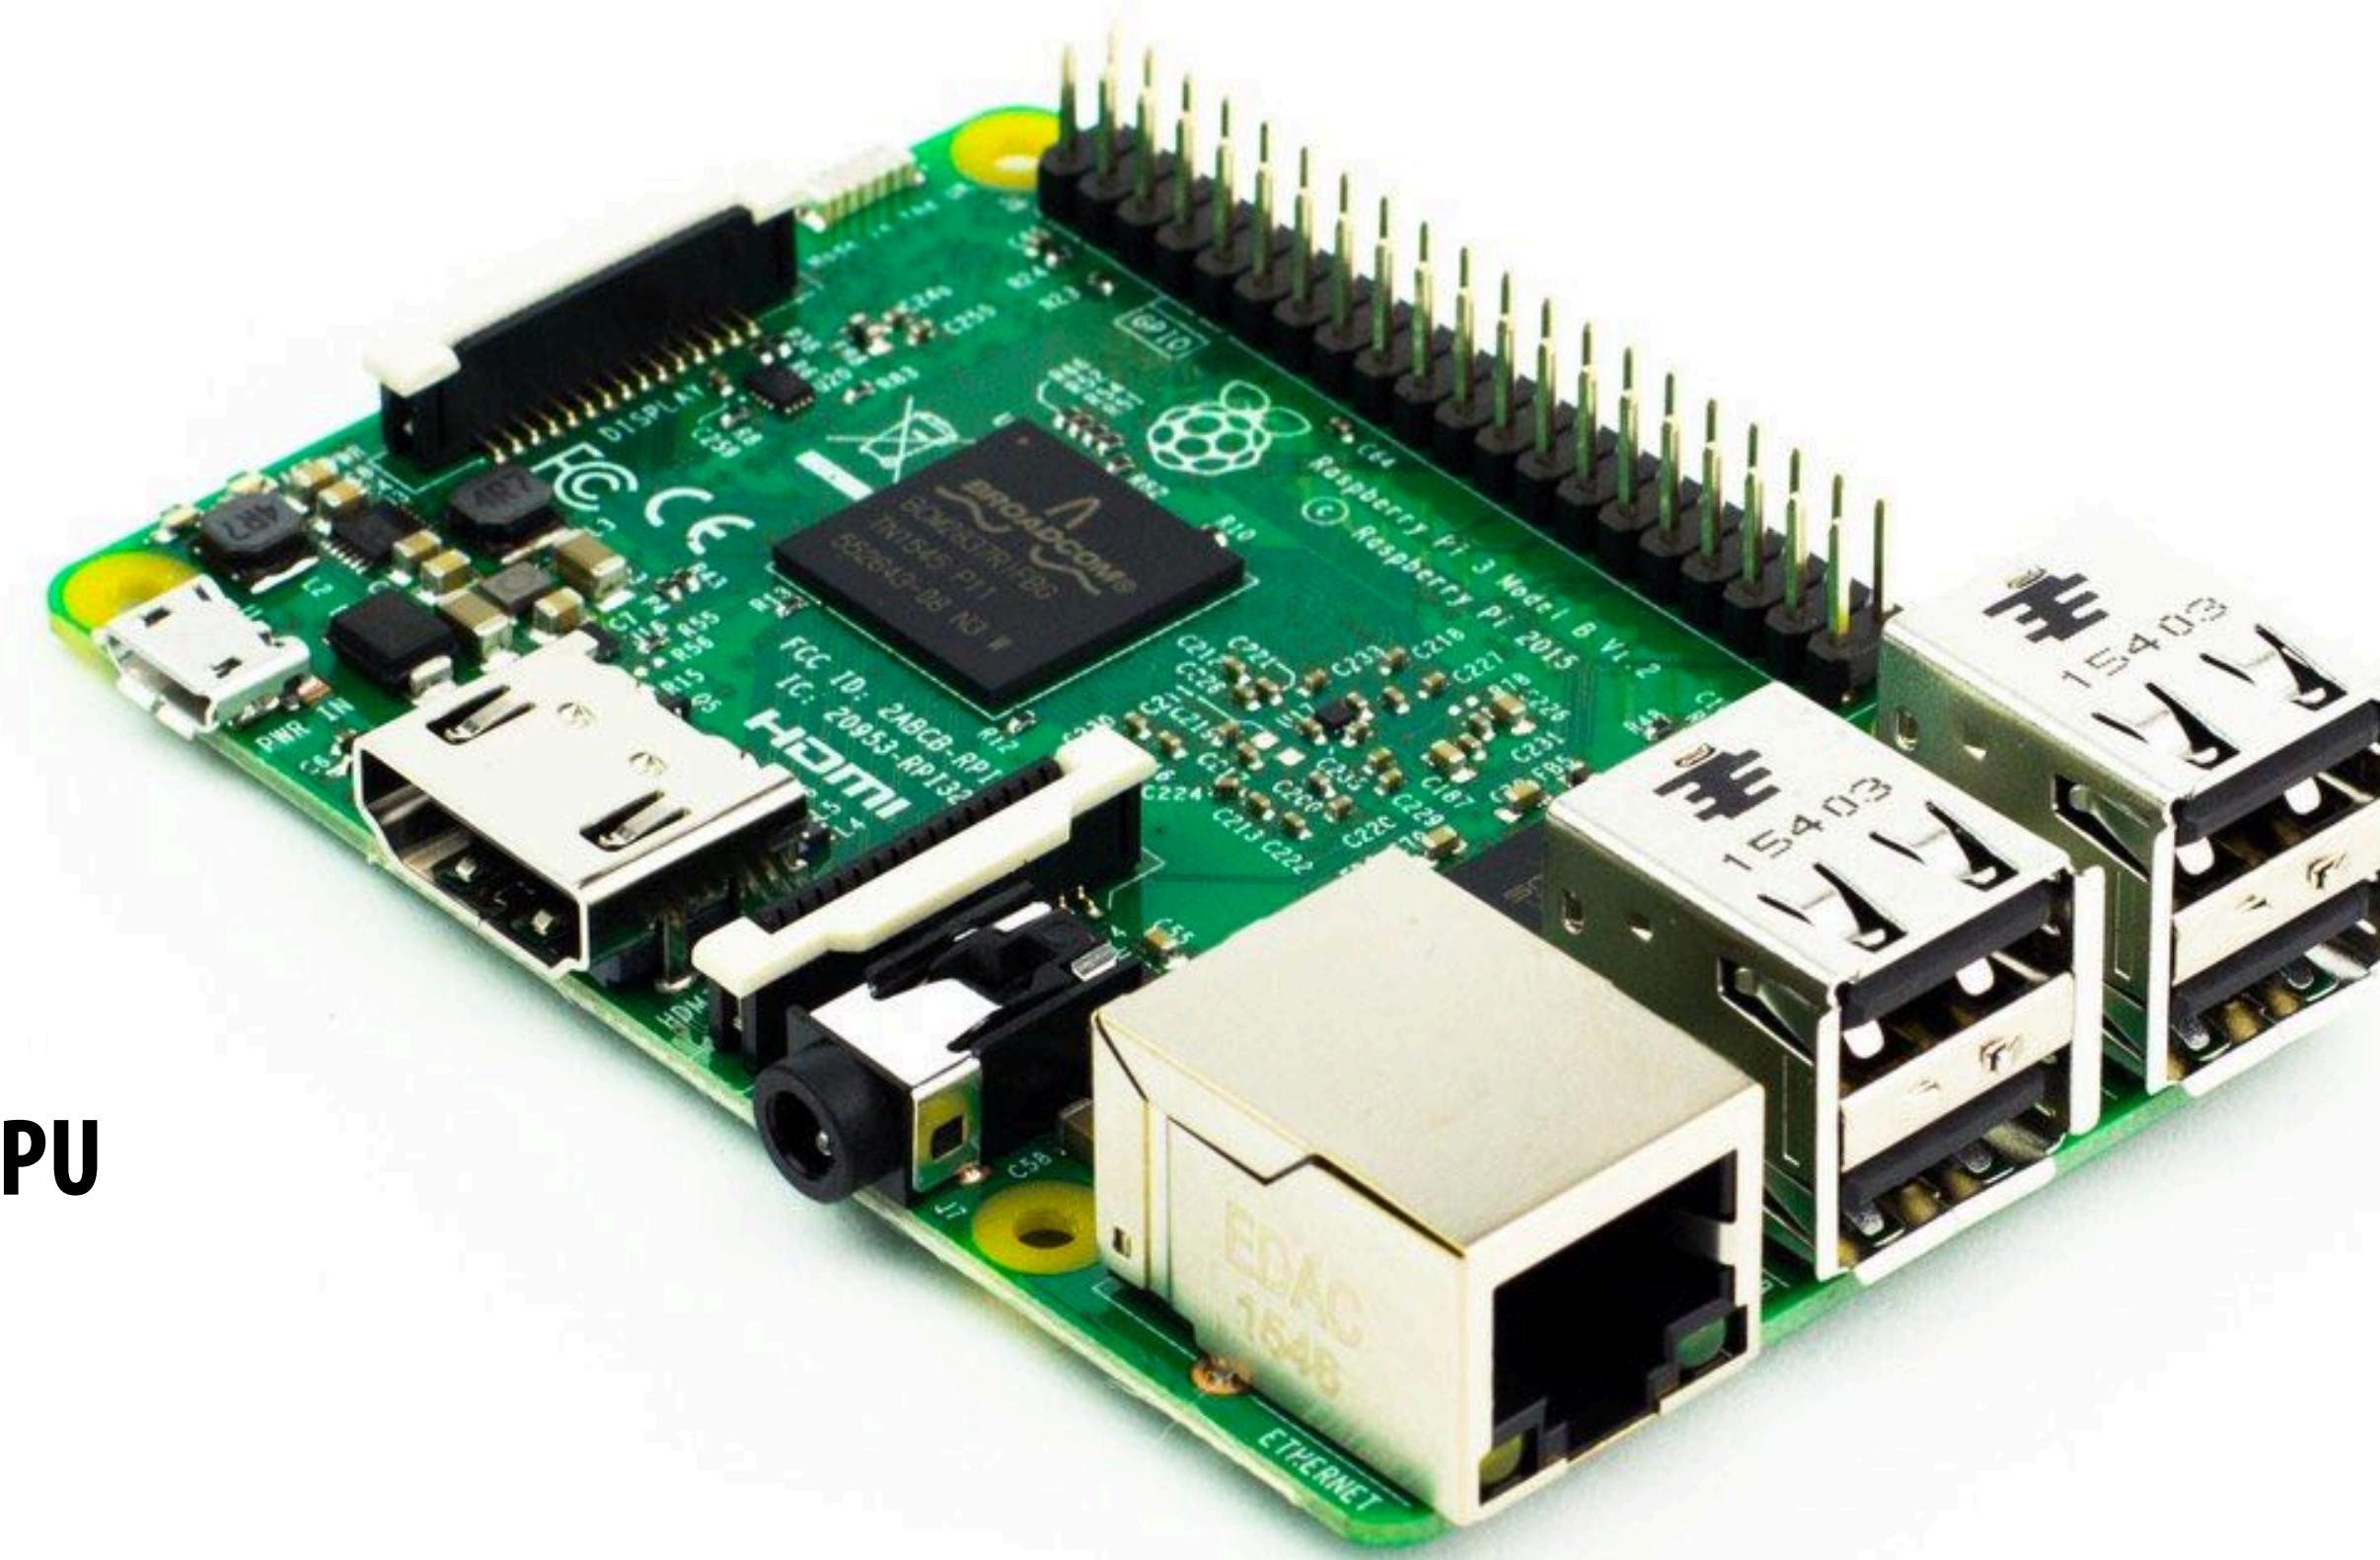
Mobile parallel processing

Power constraints also heavily influence the design of mobile systems



Mobile parallel processing

Raspberry Pi 3
Quad-core ARM A53 CPU



**But in modern computing
software must be more than just parallel...**

IT MUST ALSO BE EFFICIENT

all

Q. What is a big concern in mobile computing?

A. Power

Two reasons to save power

Run at *higher performance*
for a *fixed* amount of time.



Power = heat
If a chip gets too hot, it must be
clocked down to cool off *

Run at *sufficient performance*
for a *longer* amount of time.



Power = battery
Long battery life is a desirable
feature in mobile devices

* Another reason: hotter systems cost more to cool.

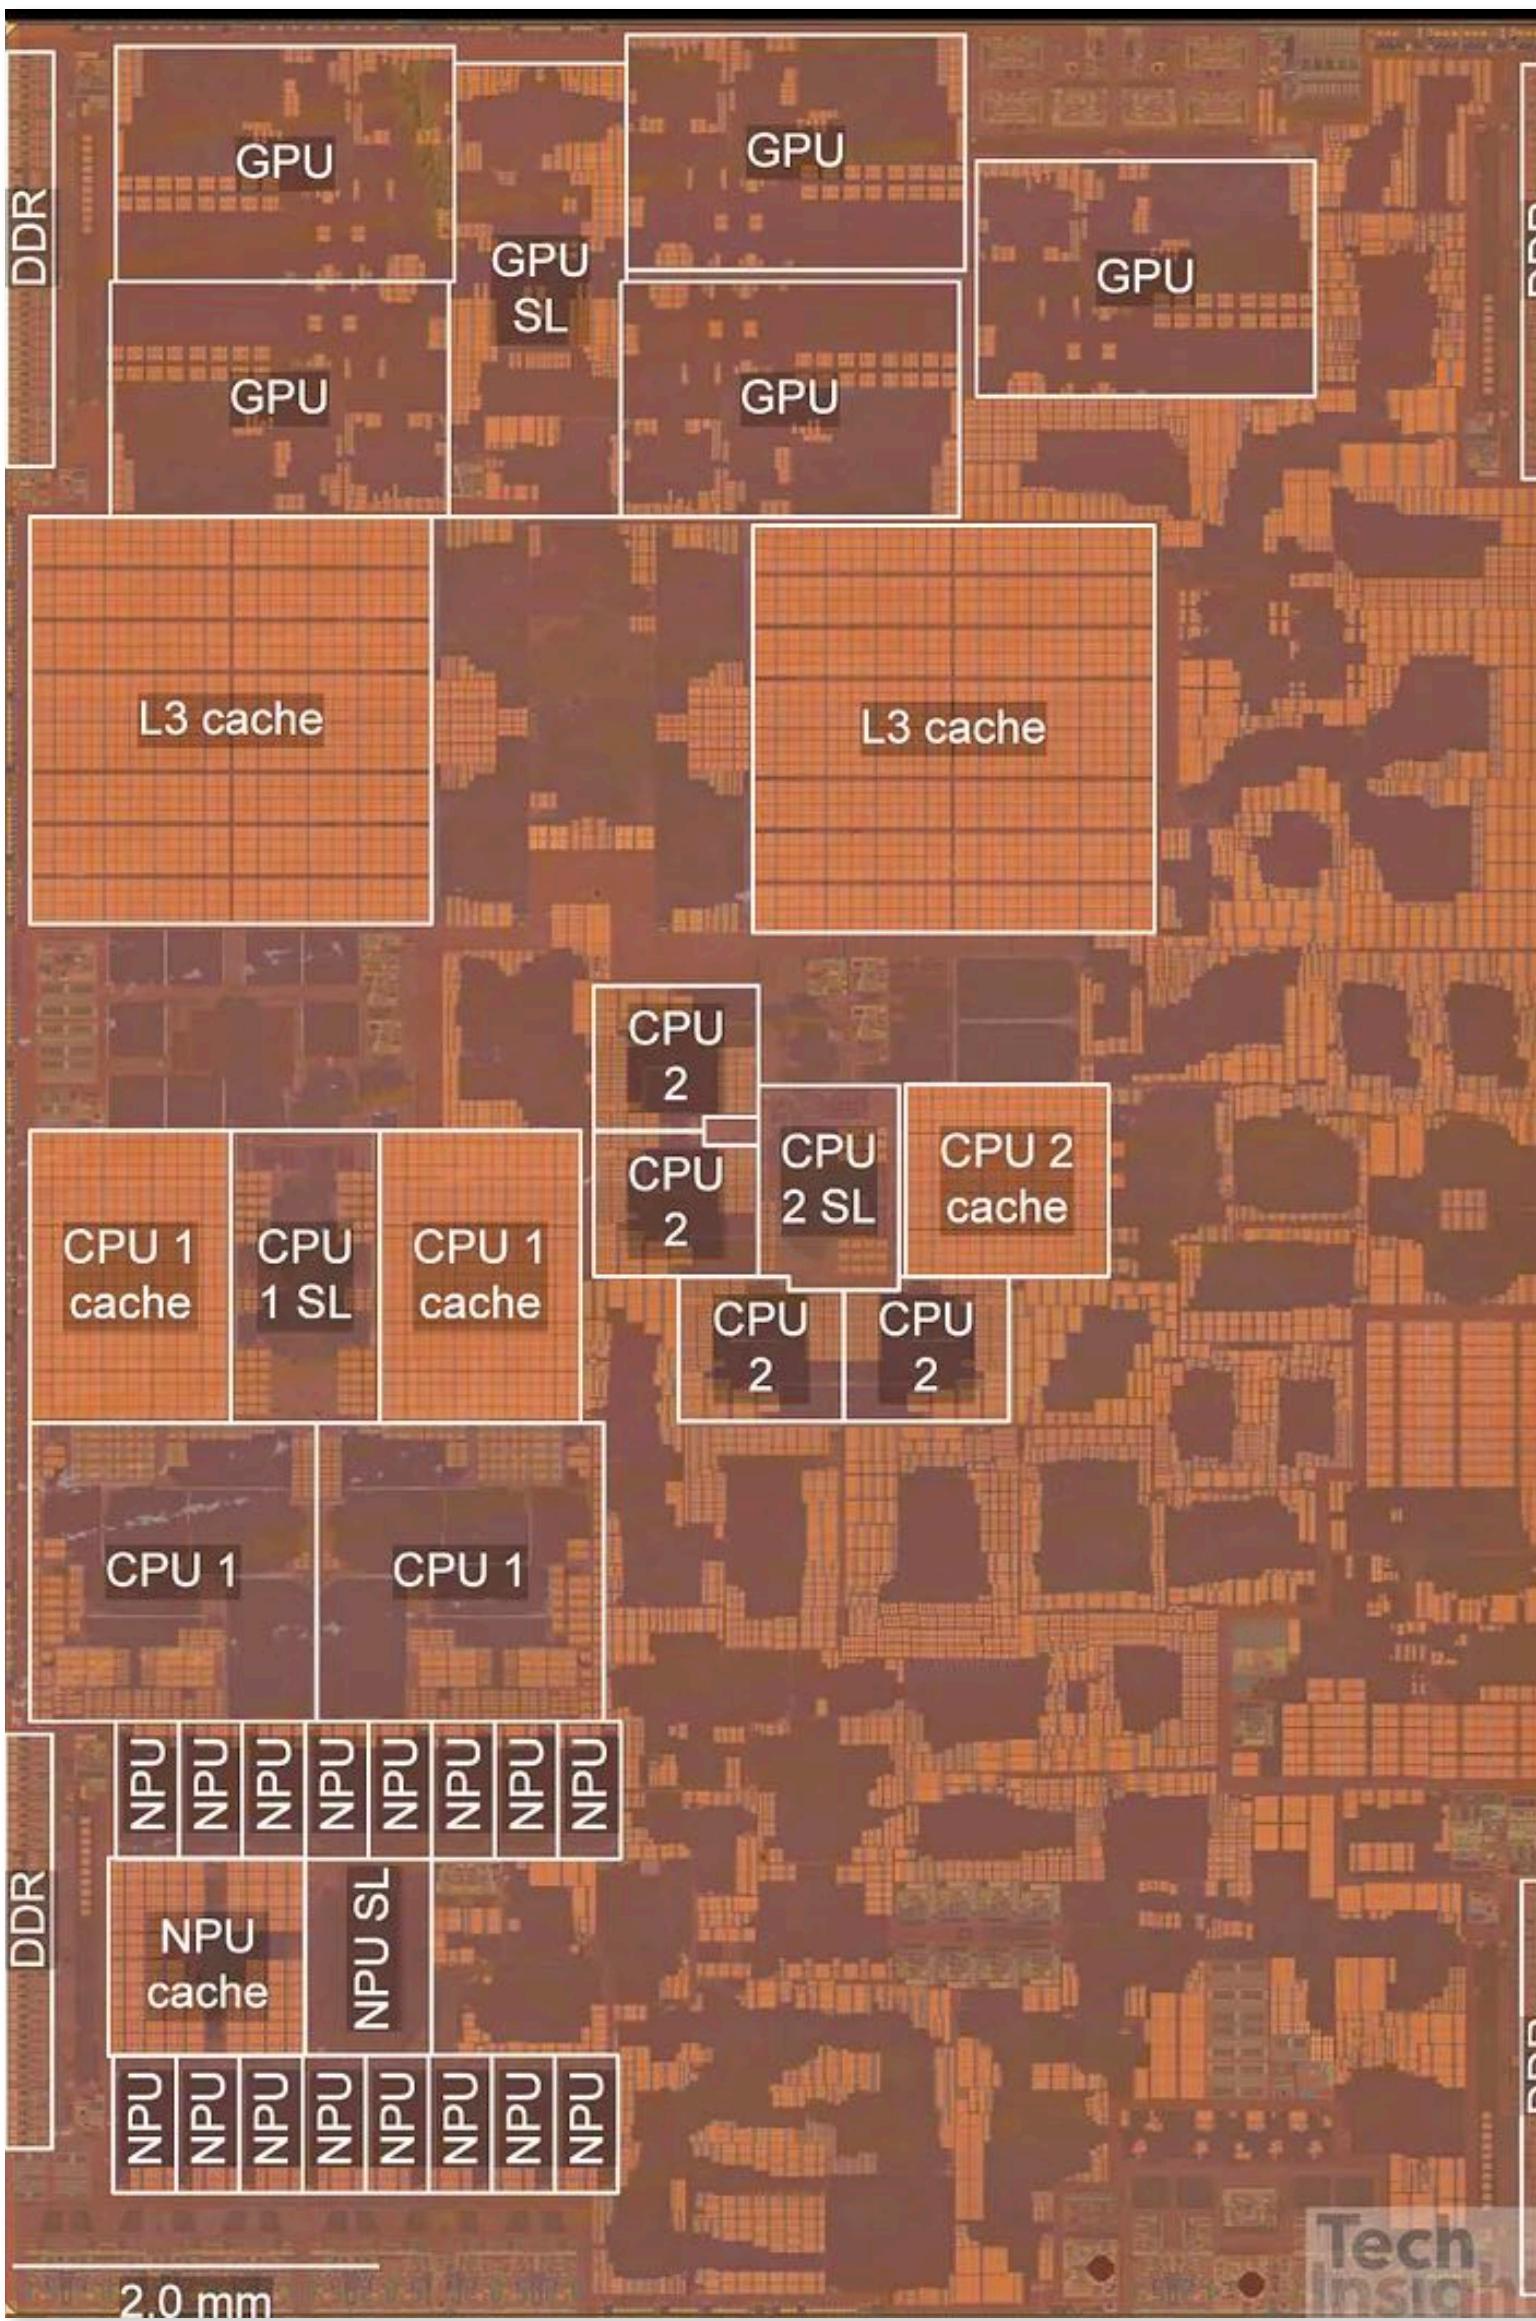
Mobile phone example

Apple iPhone 13



**3227 mAmp hours
(12.4 Watt hours)**

Specialized processing is ubiquitous in mobile systems



**Apple A15 Bionic
(in iPhone 13, 14)**

15 billion transistors

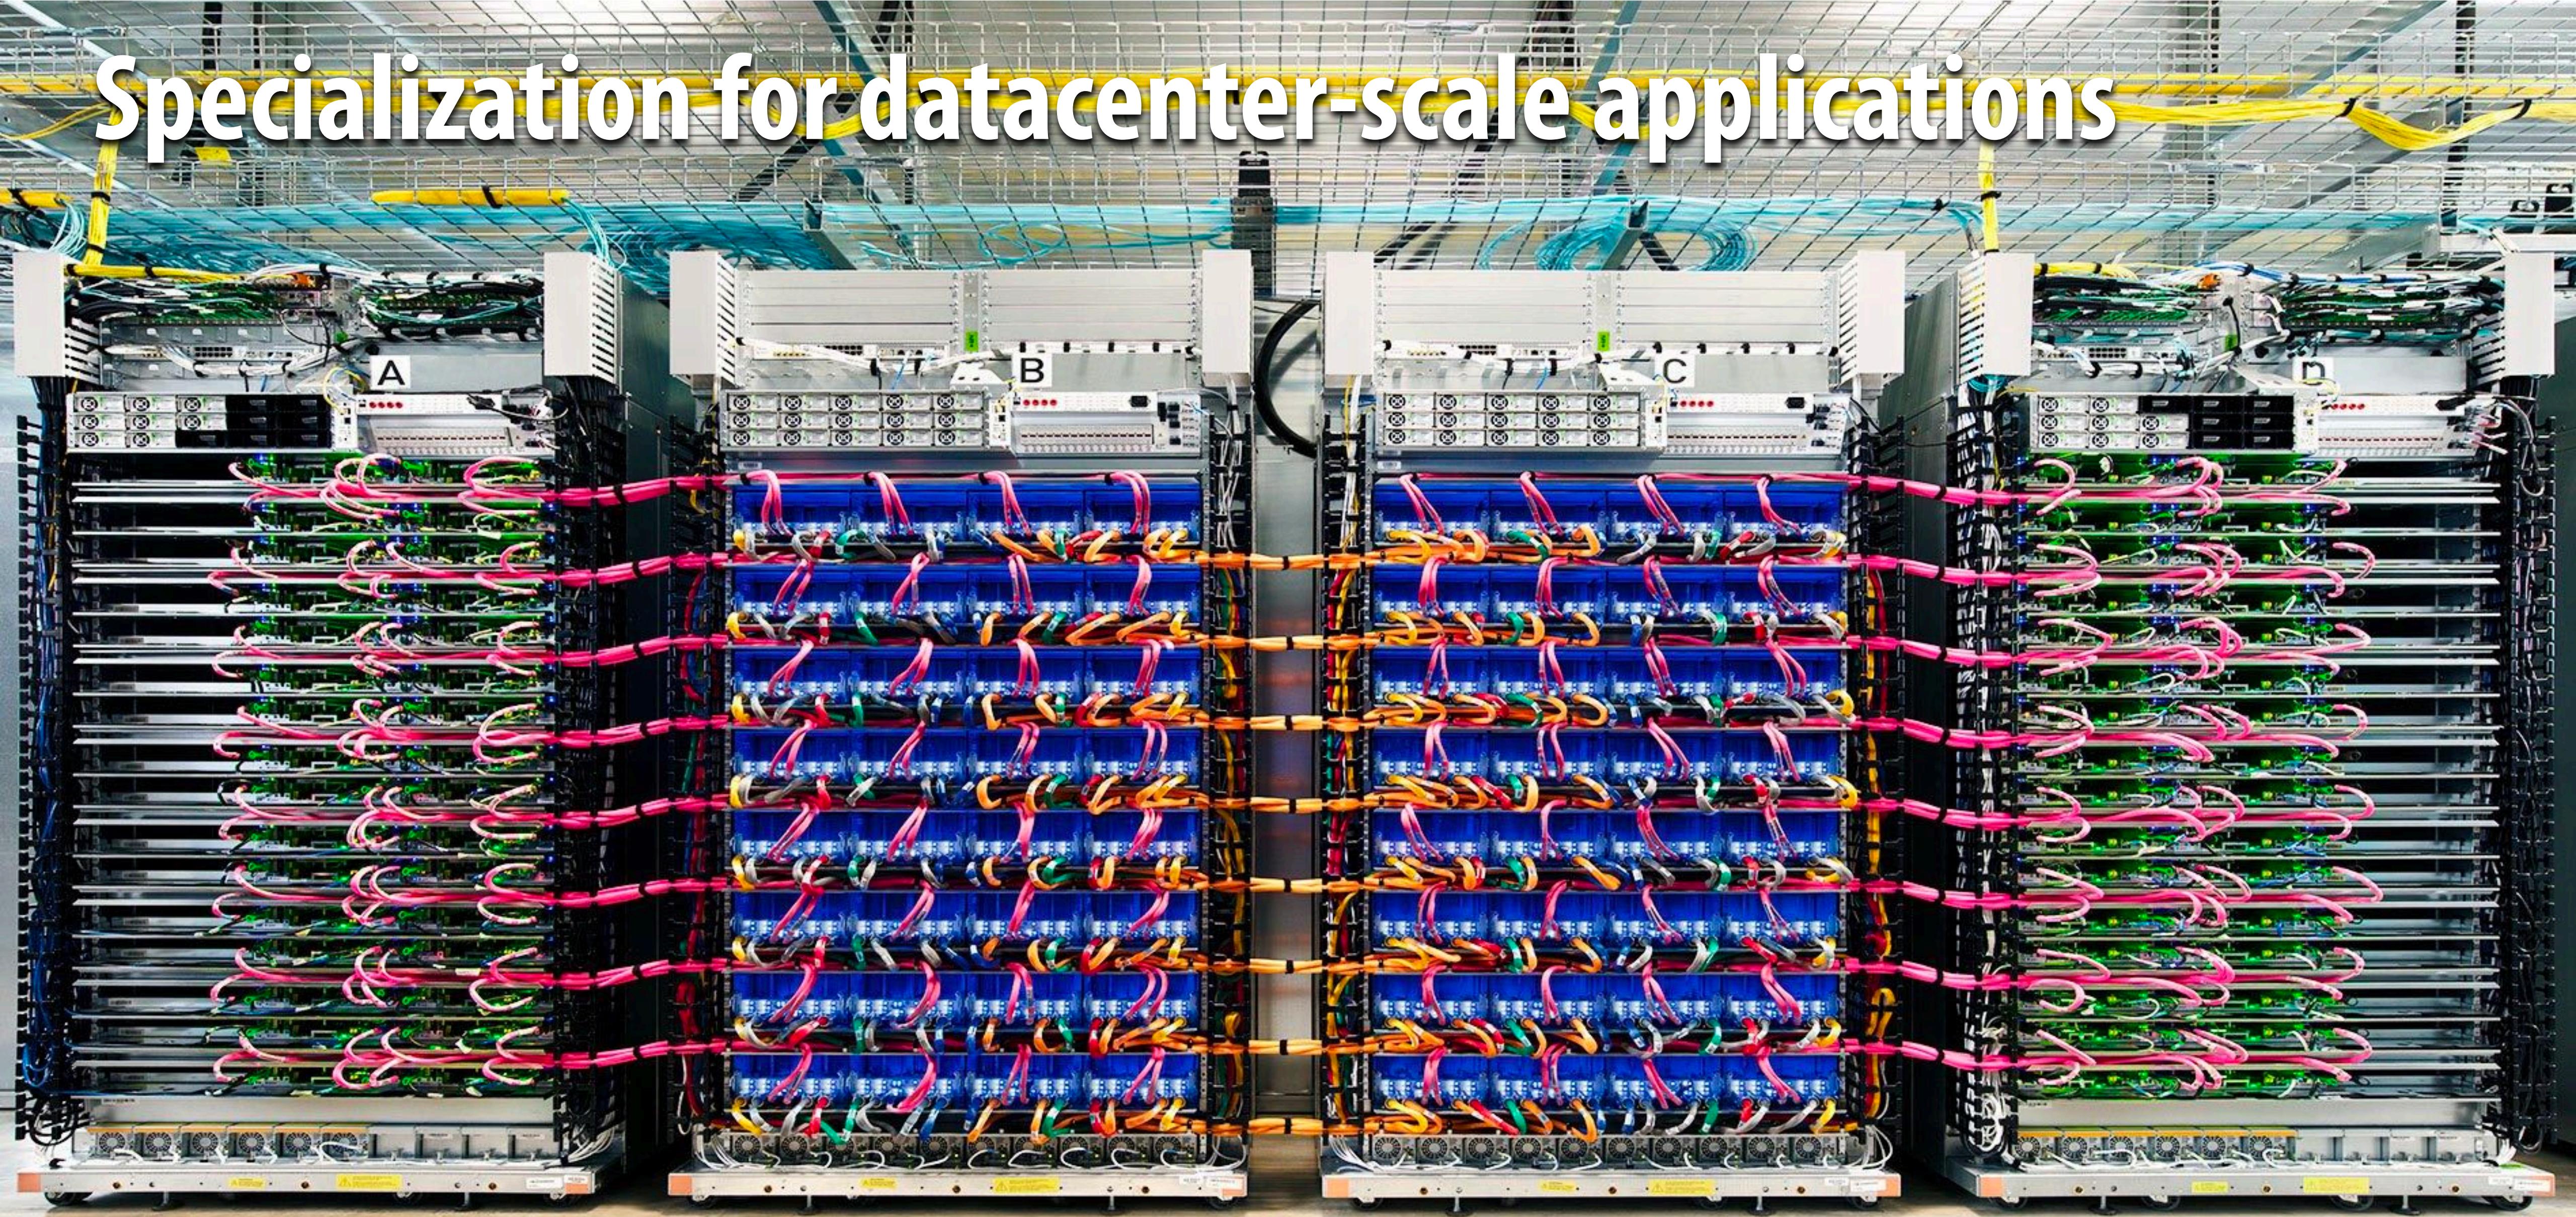
6-core GPU
2 “big” CPU cores
4 “small” CPU cores

Apple-designed multi-core GPU
Neural Engine (NPU) for DNN acceleration +
Image/video encode/decode processor +
Motion (sensor) processor

Parallel + specialized HW

- Achieving high efficiency will be a key theme in this class
- We will discuss how modern systems not only use many processing units, but also utilize specialized processing units to achieve high levels of power efficiency

Specialization for datacenter-scale applications

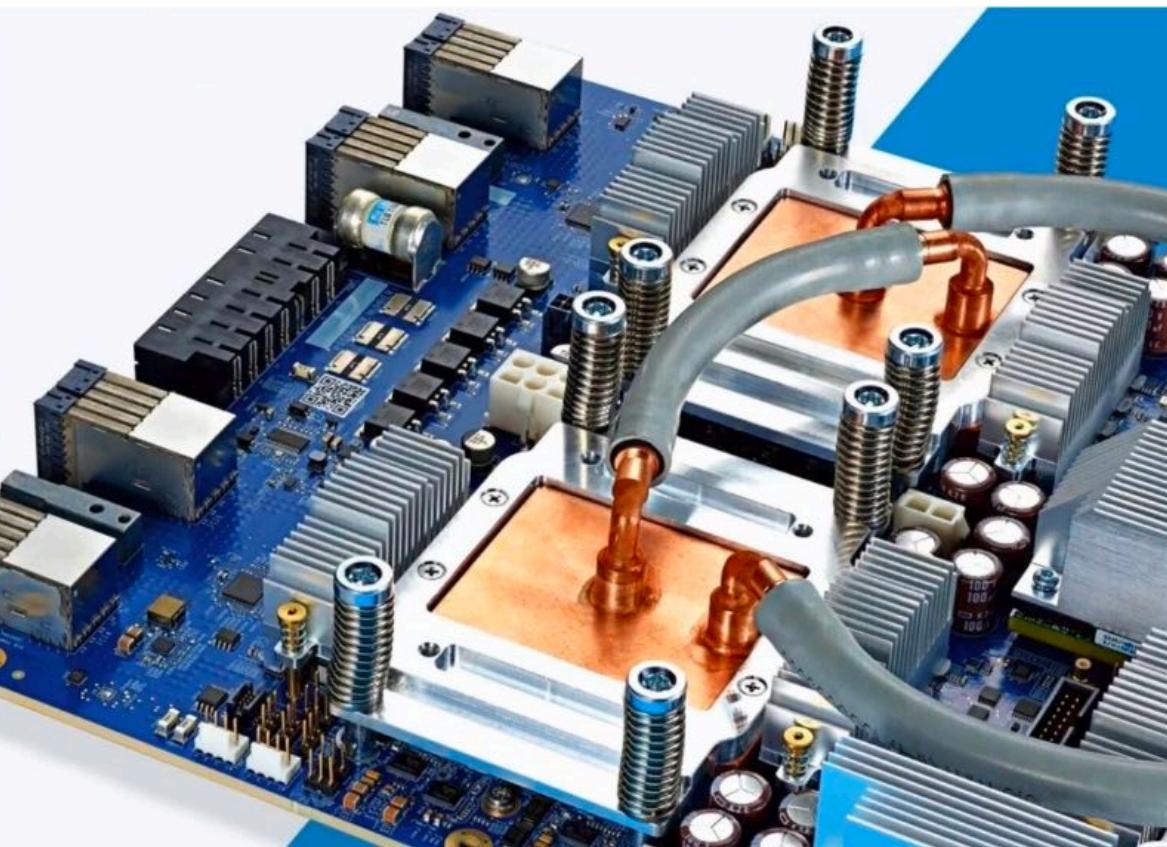


Google TPU pods

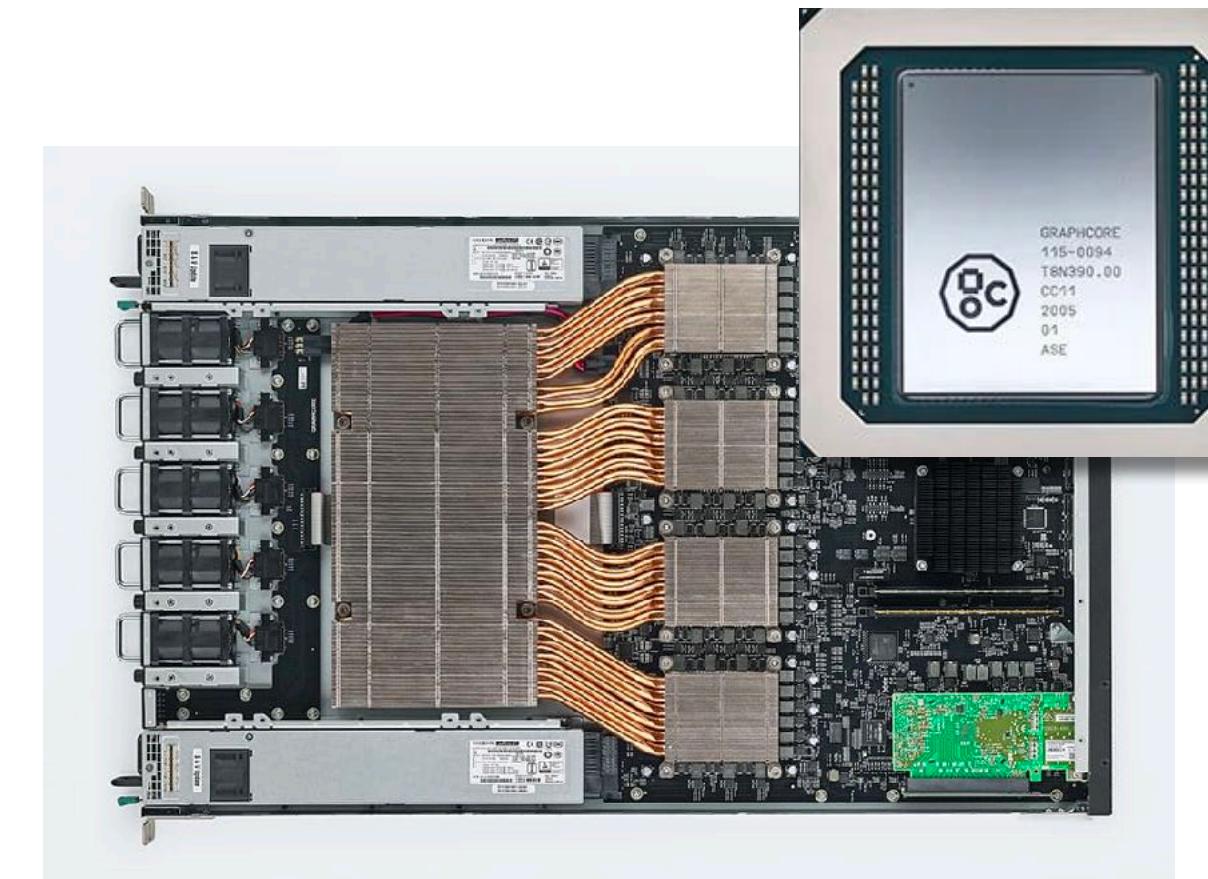
TPU = Tensor Processing Unit: specialized processor for ML computations

Image Credit: TechInsights Inc.

Specialized hardware to accelerate DNN inference/training



Google TPU3



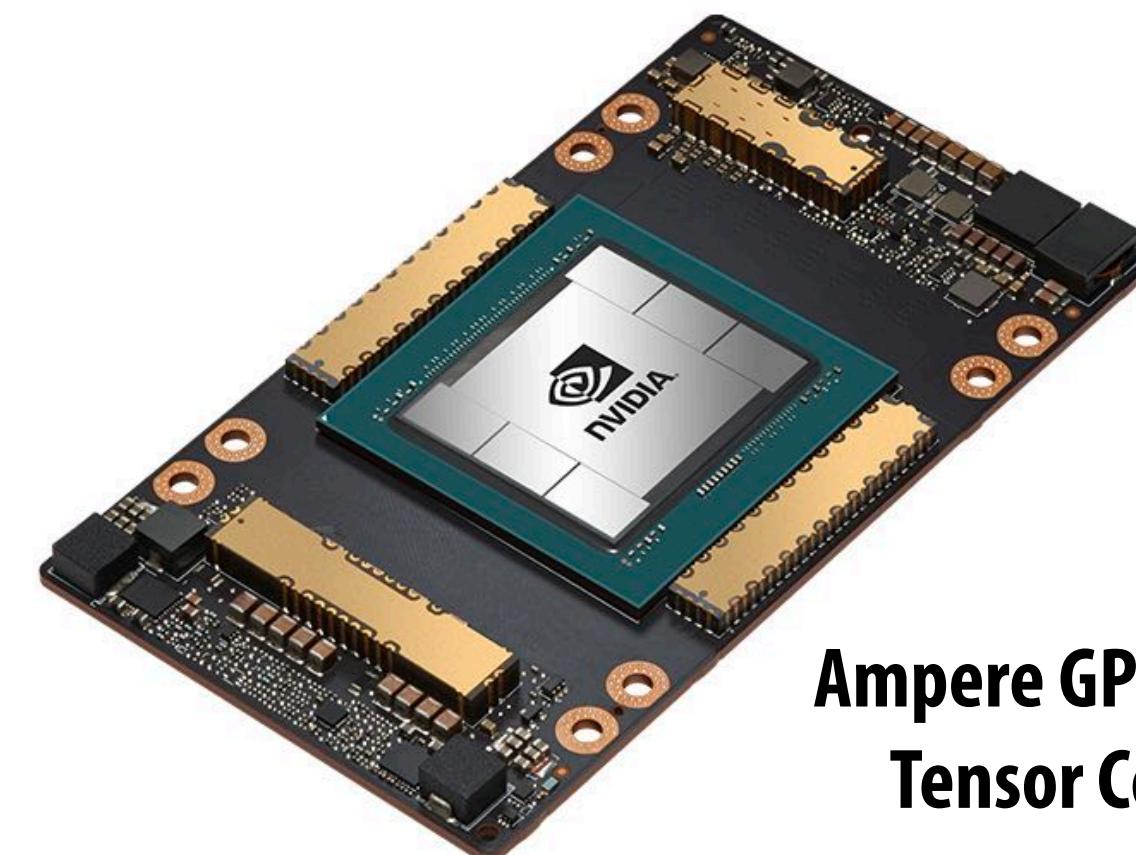
GraphCore IPU



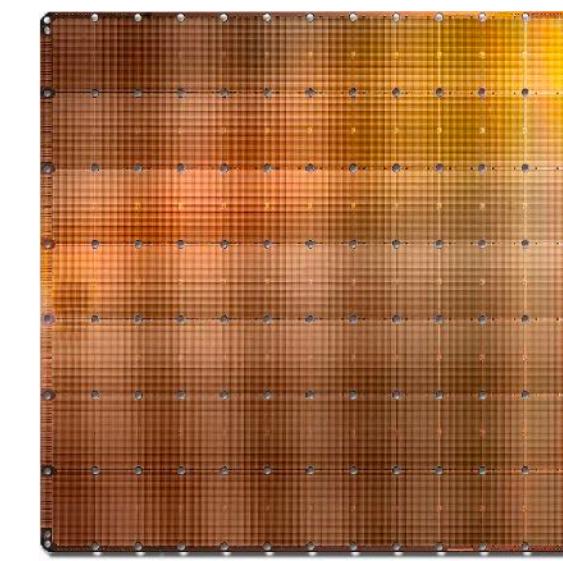
Apple Neural Engine



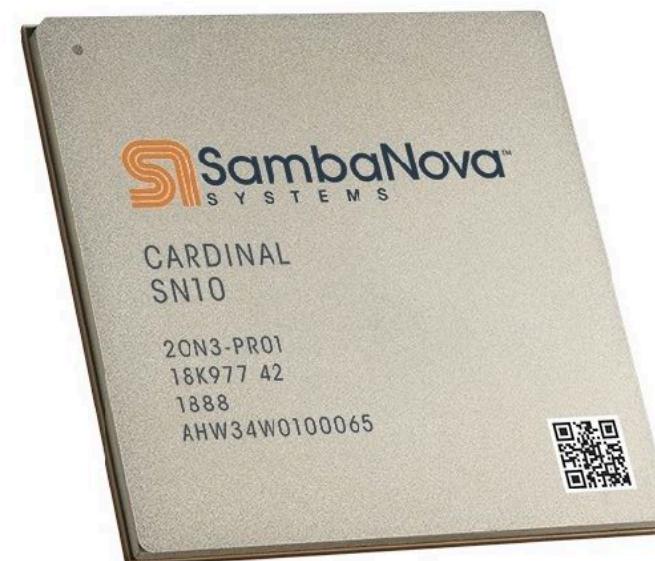
AWS Trainium



Ampere GPU with
Tensor Cores



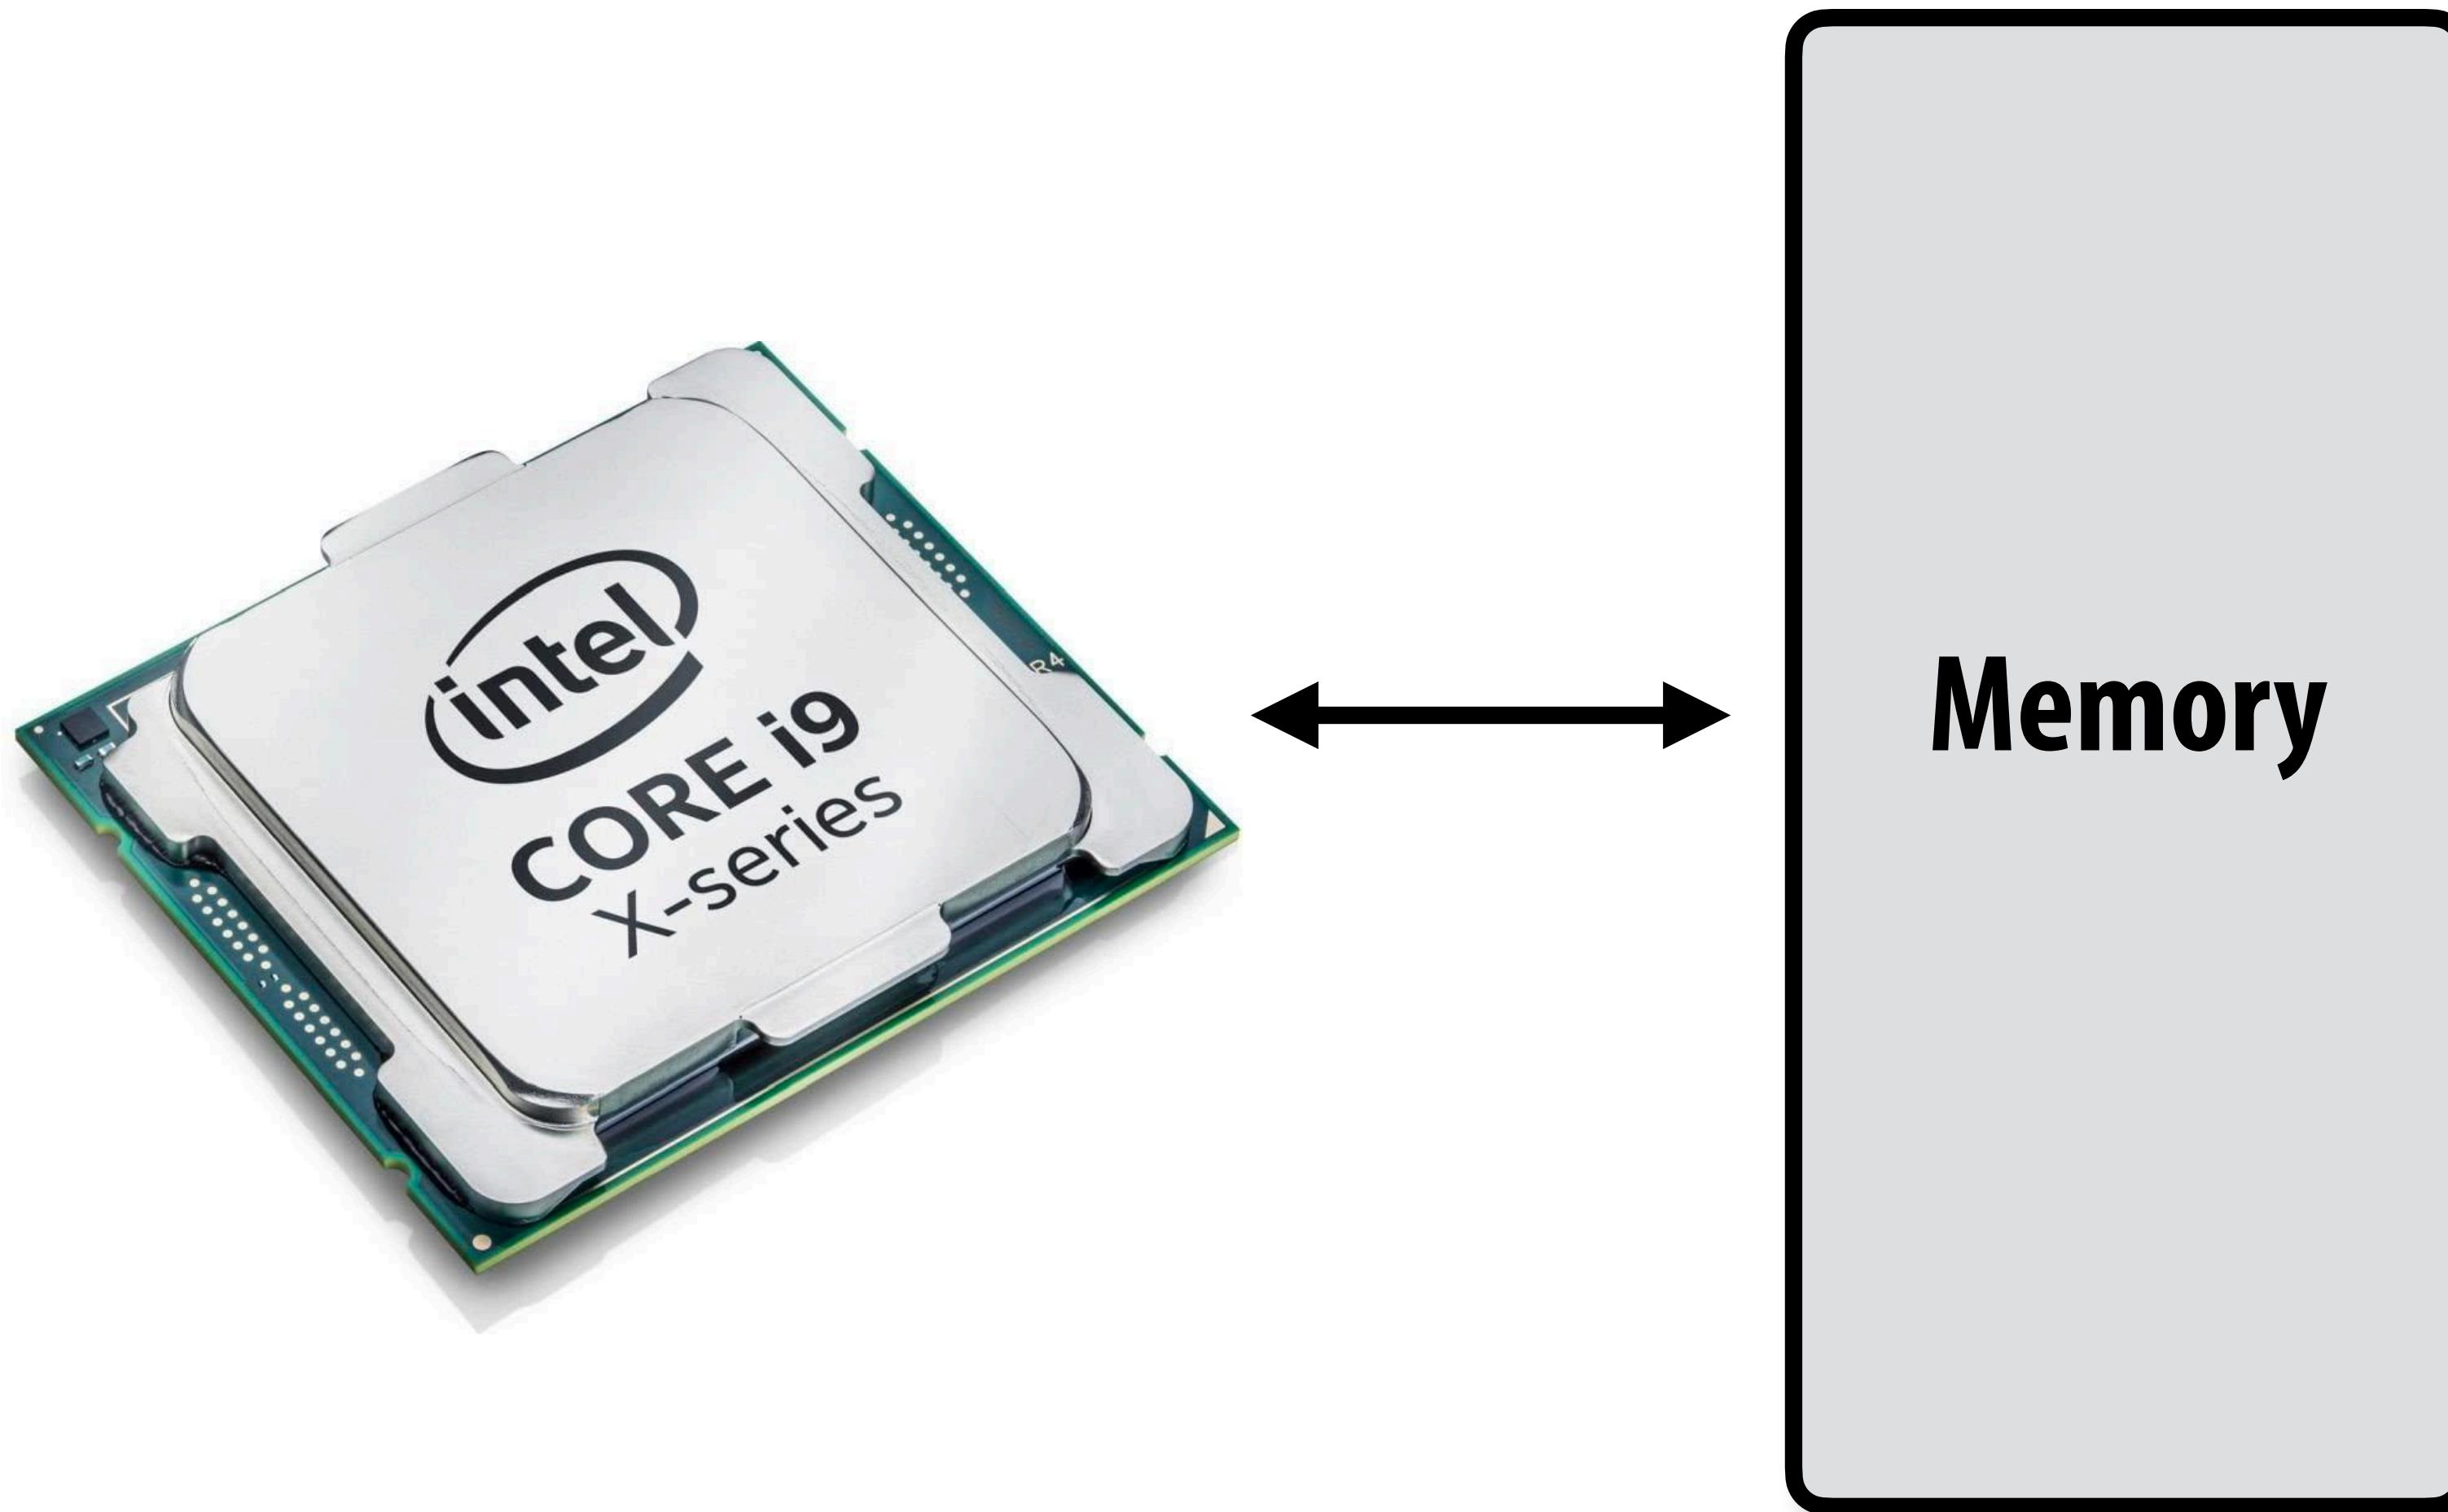
Cerebras Wafer Scale Engine



SambaNova
Cardinal SN10

**Achieving efficient processing
almost always comes down to
accessing data efficiently.**

What is memory?



A program's memory address space

- A computer's memory is organized as an array of bytes
- Each byte is identified by its "address" in memory
(its position in this array)
(We'll assume memory is byte-addressable)

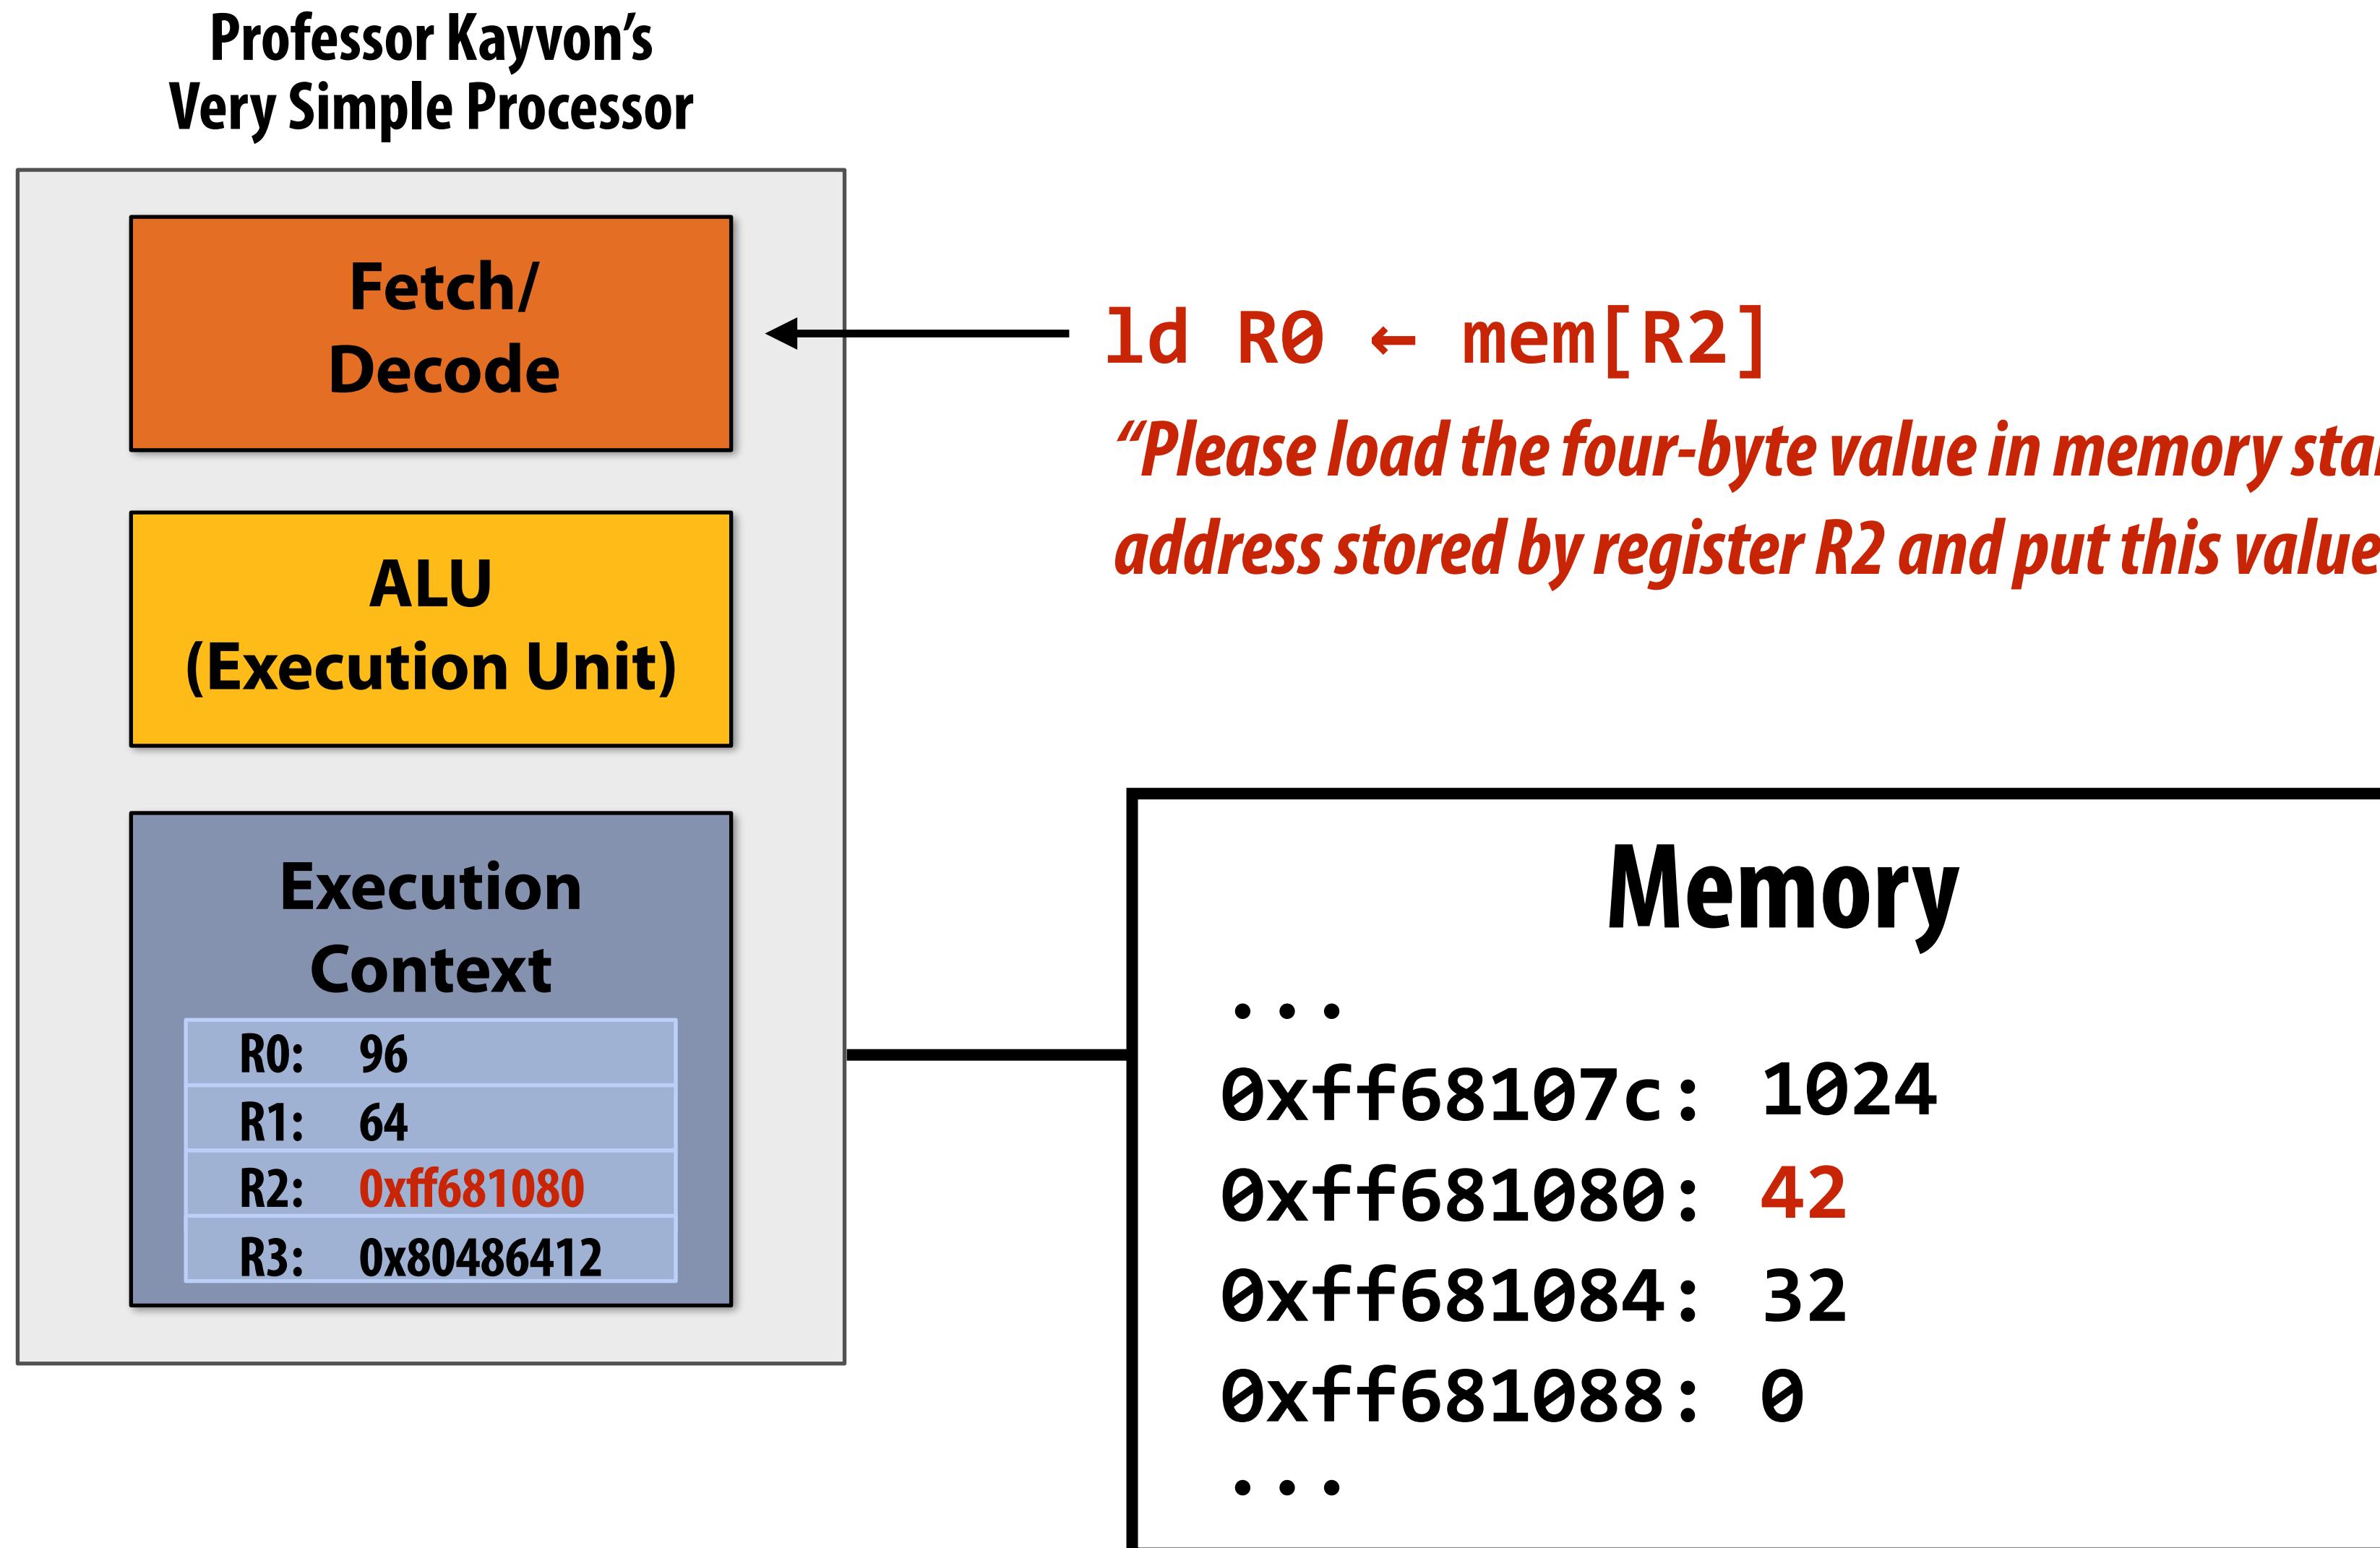
"The byte stored at address 0x8 has the value 32."

"The byte stored at address 0x10 (16) has the value 128."

In the illustration on the right, the program's
memory address space is 32 bytes in size
(so valid addresses range from 0x0 to 0x1F)

Address	Value
0x0	16
0x1	255
0x2	14
0x3	0
0x4	0
0x5	0
0x6	6
0x7	0
0x8	32
0x9	48
0xA	255
0xB	255
0xC	255
0xD	0
0xE	0
0xF	0
0x10	128
:	:
0x1F	0

Load: an instruction for accessing the contents of memory



Terminology

■ Memory access latency

- The amount of time it takes the memory system to provide data to the processor
- Example: 100 clock cycles, 100 nsec



Stalls 停顿/停火

- A processor “stalls” (can’t make progress) when it cannot run the next instruction in an instruction stream because future instructions depend on a previous instruction that is not yet complete.
- Accessing memory is a major source of stalls

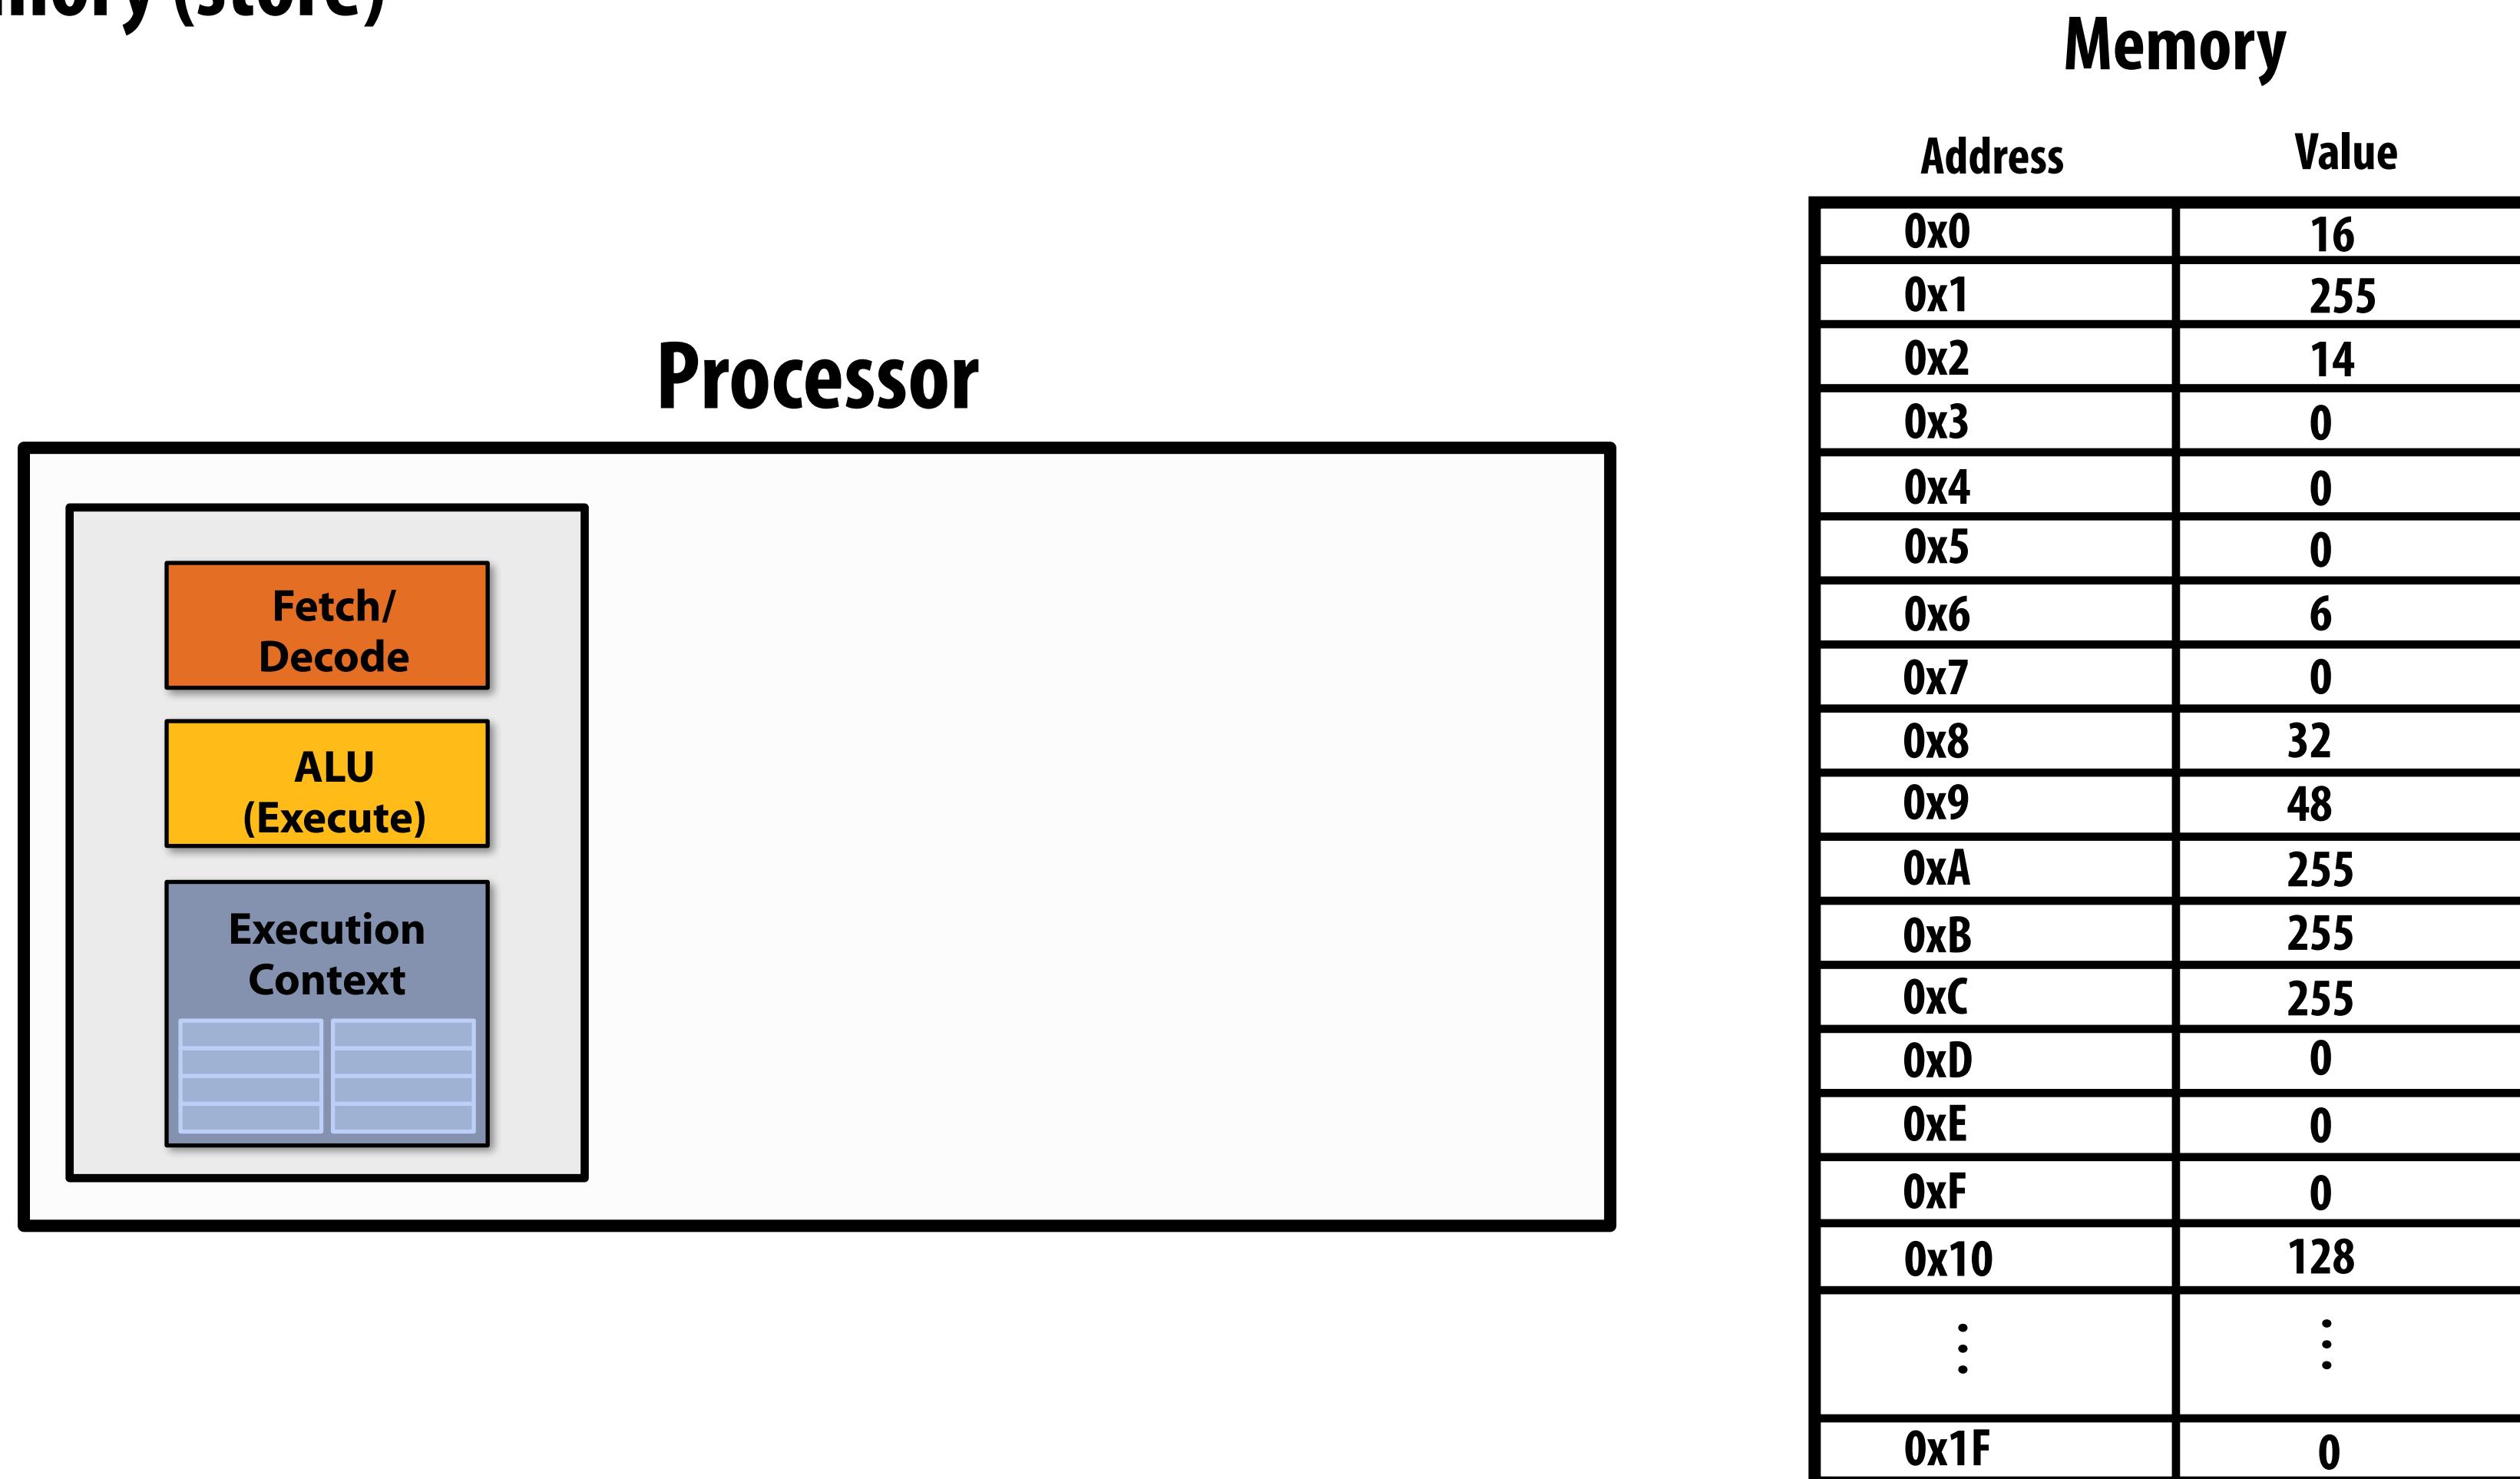
```
1d r0 mem[r2]    ←  
1d r1 mem[r3]    ←  
add r0, r0, r1    |
```

Dependency: cannot execute ‘add’ instruction until data from mem[r2] and mem[r3] have been loaded from memory

- Memory access times ~ 100’s of cycles
 - Memory “access time” is a measure of latency

What are caches?

- Recall memory is just an array of values
- And a processor has instructions for moving data from memory into registers (load) and storing data from registers into memory (store)

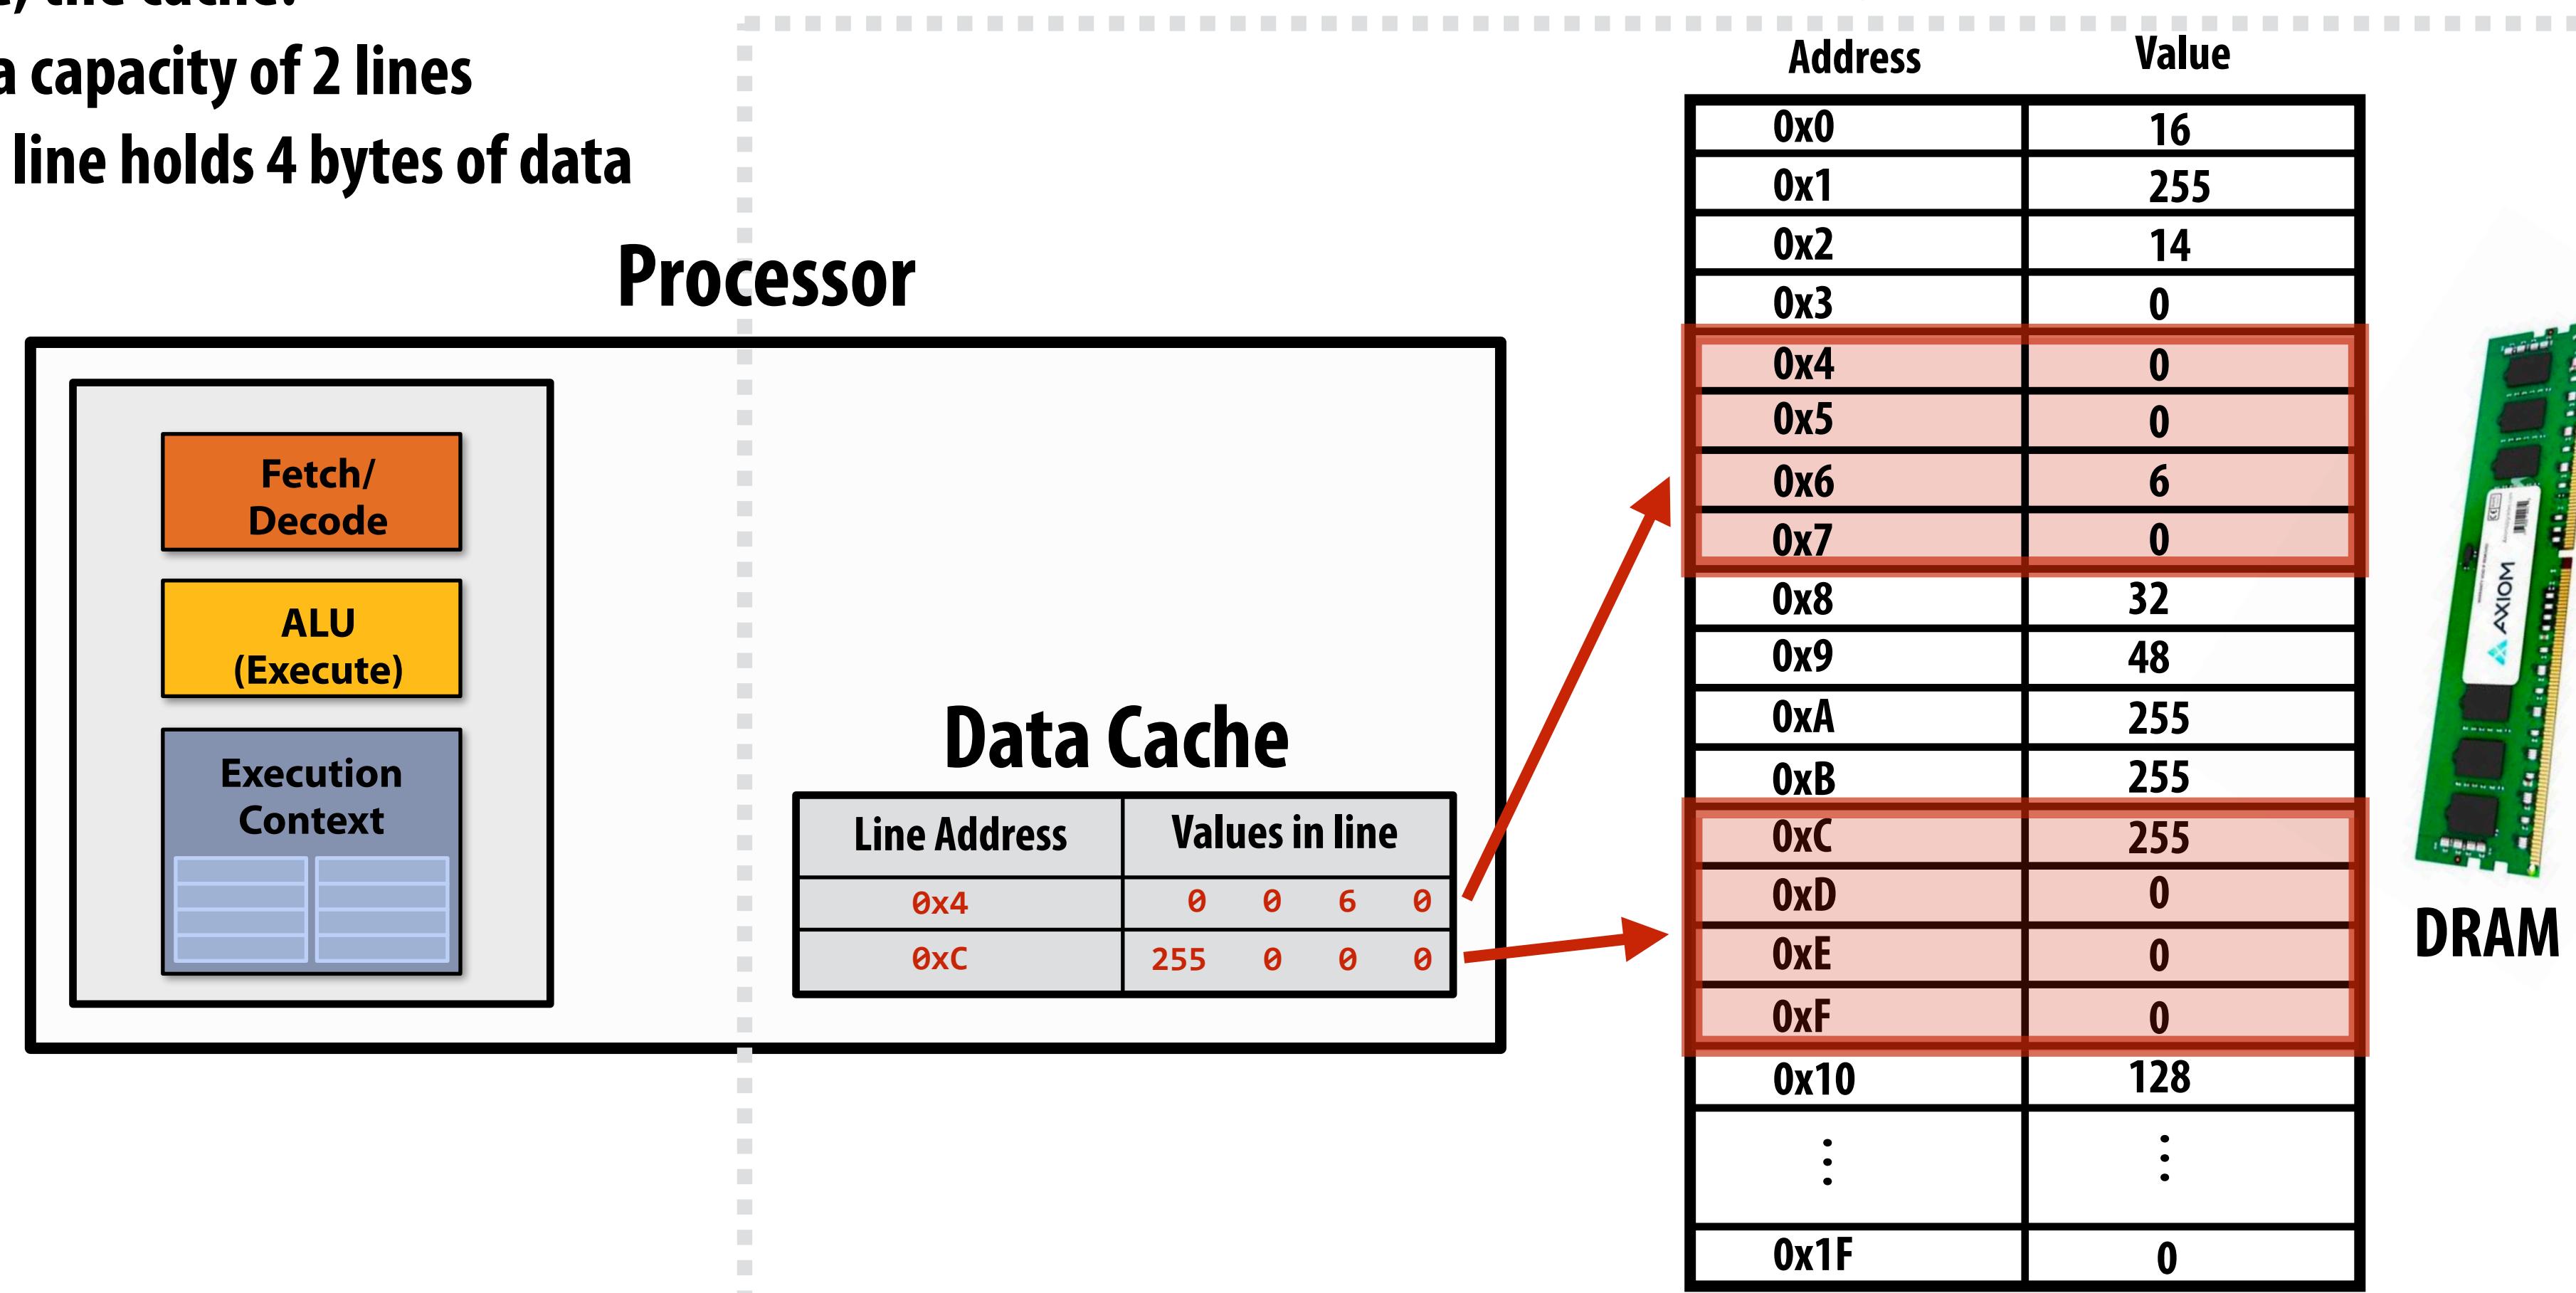


What are caches?

- A cache is a hardware implementation detail that does not impact the output of a program, only its performance
- Cache is on-chip storage that maintains a copy of a subset of the values in memory
- If an address is stored “in the cache” the processor can load/store to this address more quickly than if the data resides only in DRAM
- Caches operate at the granularity of “cache lines”.

In the figure, the cache:

- Has a capacity of 2 lines
- Each line holds 4 bytes of data



How does a processor decide what data to keep in cache?

- Outside the scope of this course, but I suggest googling these terms...
 - Direct mapped cache
 - Set-associative cache
 - Cache line
- For now, just assume that the cache of size N bytes stores values for the last N addresses accessed
 - LRU replacement policy ("least recently used") - to make room for new data, throw out the data in the cache that was accessed the longest time ago

Cache example 1

Array of 16 bytes in memory

Address	Value
0x0	16
0x1	255
0x2	14
0x3	0
Line 0x4	0
0x4	0
0x5	0
0x6	6
Line 0x8	0
0x8	32
0x9	48
0xA	255
Line 0xC	255
0xB	255
0xC	255
0xD	0
Line 0xE	0
0xE	0
Line 0xF	0
0xF	0

Assume:

Total cache capacity of 8 bytes

Cache with 4-byte cache lines
(So 2 lines fit in cache)

Least recently used (LRU)
replacement policy

time

Address accessed	Cache action	Cache state (after load is complete)
0x0	"cold miss", load 0x0	0x0 ●●●●
0x1	hit	0x0 ●●●●
0x2	hit	0x0 ●●●●
0x3	hit	0x0 ●●●●
0x2	hit	0x0 ●●●●
0x1	hit	0x0 ●●●●
0x4	"cold miss", load 0x4	0x4 ●●●●
0x1	hit	0x4 ●●●●

There are two forms of "data locality" in this sequence:

Spatial locality: loading data in a cache line "preloads" the data needed for subsequent accesses to different addresses in the same line, leading to cache hits

Temporal locality: repeated accesses to the same address result in hits.

Cache example 2

Array of 16 bytes in memory

Address	Value
0x0	16
0x1	255
0x2	14
0x3	0
Line 0x0	0x4
	0
	0
	0x6
Line 0x4	6
	0
	0
	0x7
Line 0x8	0
	32
	48
	0xA
Line 0xC	255
	255
	0xD
	0
Line 0xE	0
	0
	0xF
	0

Assume:

Total cache capacity of 8 bytes

Cache with 4-byte cache lines
(So 2 lines fit in cache)

Least recently used (LRU)
replacement policy

Address accessed	Cache action	Cache state (after load is complete)
0x0	"cold miss", load 0x0	0x0 ●●●●
0x1	hit	0x0 ●●●●
0x2	hit	0x0 ●●●●
0x3	hit	0x0 ●●●●
0x4	"cold miss", load 0x4	0x0 ●●●● 0x4 ●●●●
0x5	hit	0x0 ●●●● 0x4 ●●●●
0x6	hit	0x0 ●●●● 0x4 ●●●●
0x7	hit	0x0 ●●●● 0x4 ●●●●
0x8	"cold miss", load 0x8 (evict 0x0)	0x8 ●●●● 0x4 ●●●●
0x9	hit	0x8 ●●●● 0x4 ●●●●
0xA	hit	0x8 ●●●● 0x4 ●●●●
0xB	hit	0x8 ●●●● 0x4 ●●●●
0xC	"cold miss", load 0xC (evict 0x4)	0x8 ●●●● 0xC ●●●●
0xD	hit	0x8 ●●●● 0xC ●●●●
0xE	hit	0x8 ●●●● 0xC ●●●●
0xF	hit	0x8 ●●●● 0xC ●●●●
0x0	"capacity miss", load 0x0 (evict 0x8)	0x0 ●●●● 0xC ●●●●

time ↓

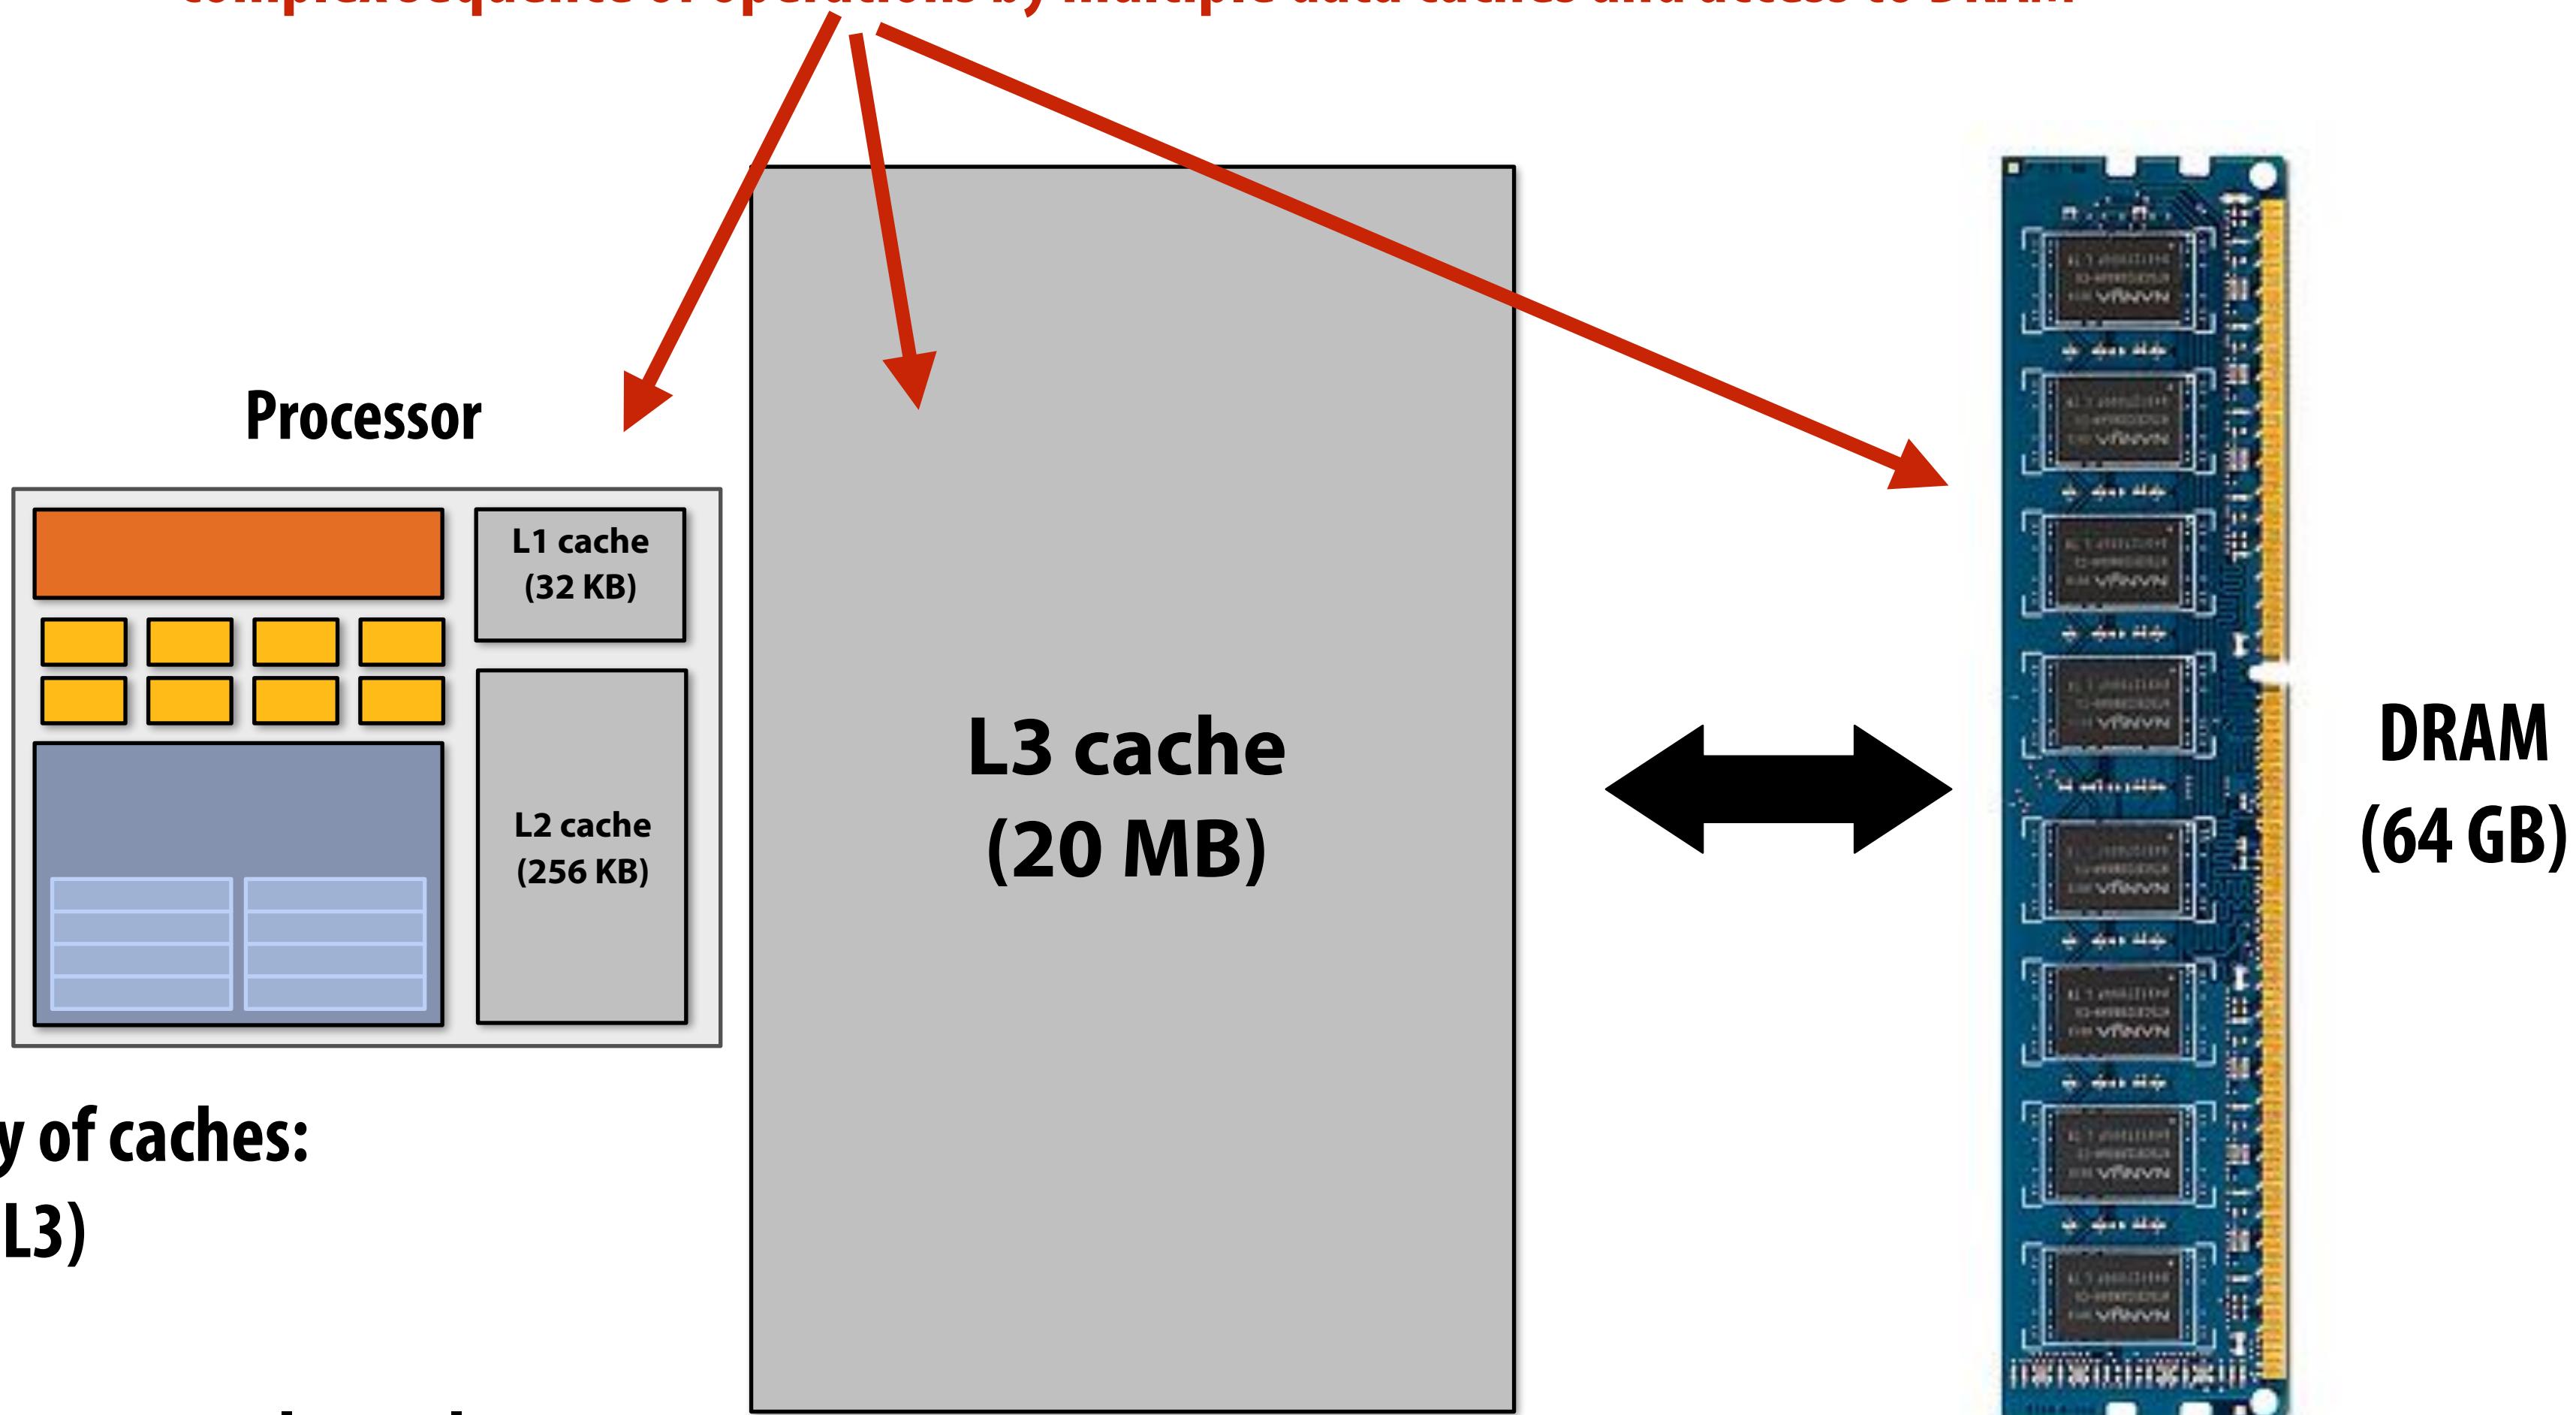
Caches reduce length of stalls (reduce memory access latency)

- Processors run efficiently when they access data that is resident in caches
- Caches reduce memory access latency when processors accesses data that they have recently accessed! *

* Caches also provide high bandwidth data transfer

The implementation of the linear memory address space abstraction on a modern computer is complex

The instruction “load the value stored at address X into register R0” might involve a complex sequence of operations by multiple data caches and access to DRAM



Common organization: hierarchy of caches:

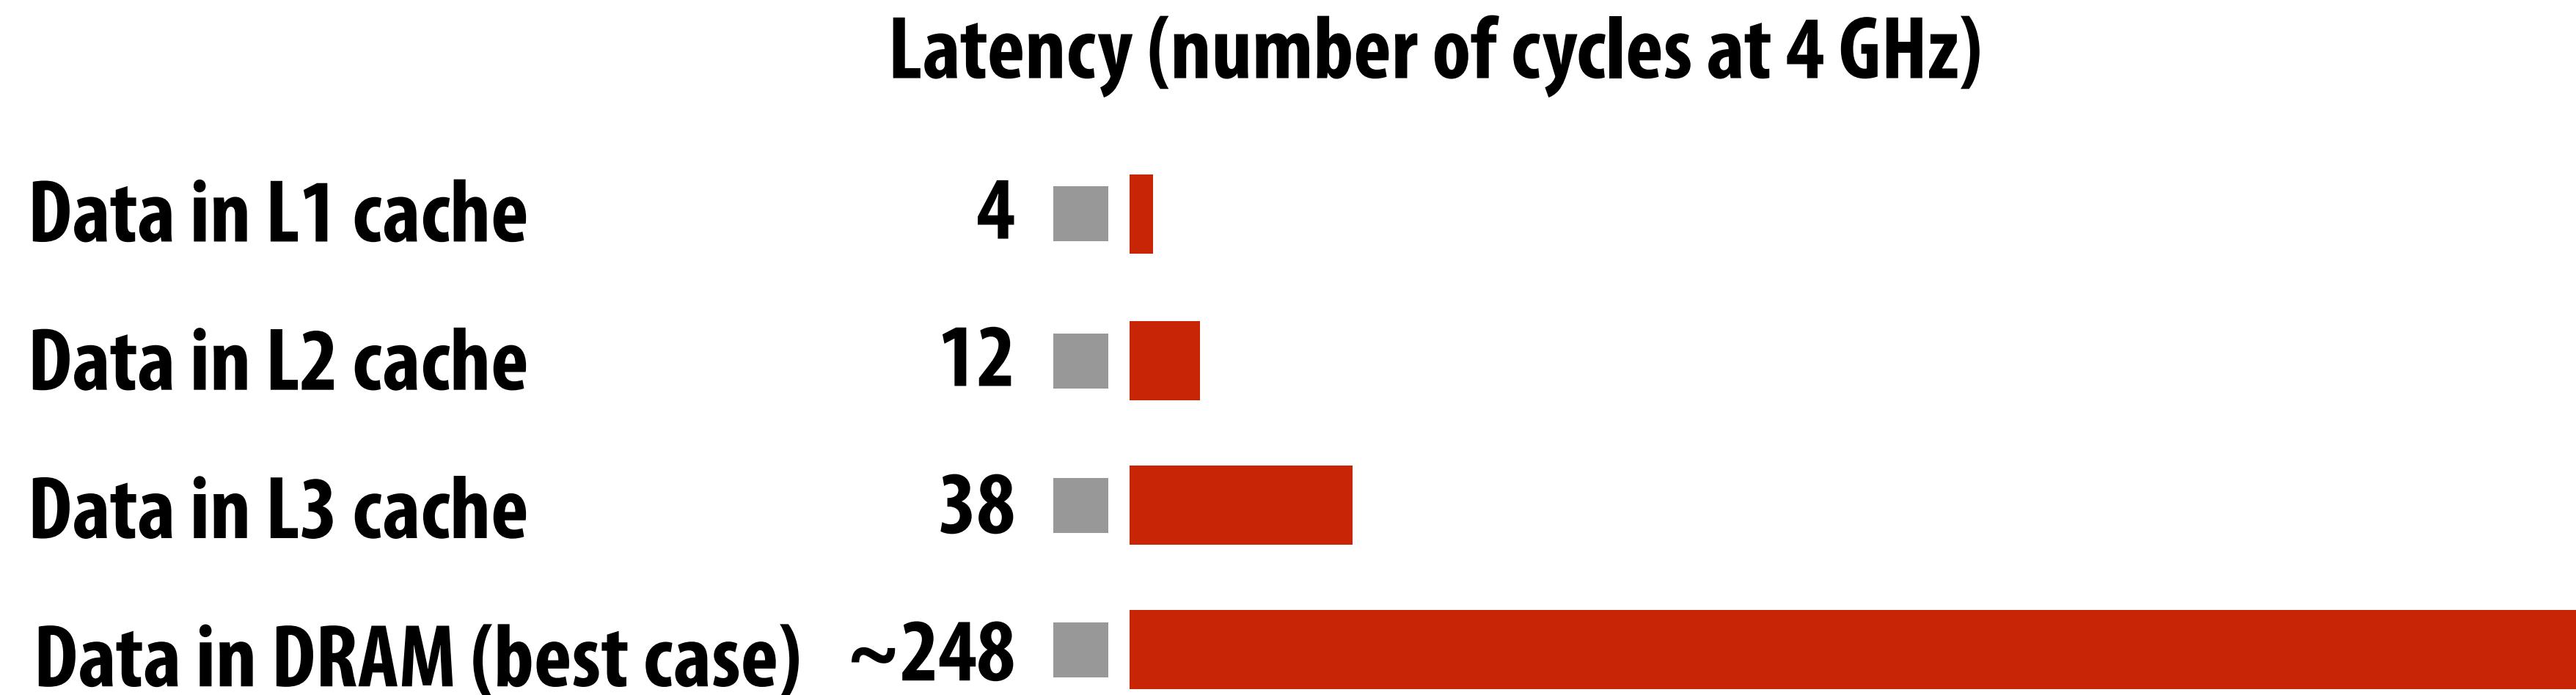
Level 1 (L1), level 2 (L2), level 3 (L3)

Smaller capacity caches near processor → lower latency

Larger capacity caches farther away → larger latency

Data access times

(Kaby Lake CPU)



Data movement has high energy cost

- Rule of thumb in modern system design: always seek to reduce amount of data movement in a computer
- “Ballpark” numbers
 - Integer op: ~ 1 pJ *
 - Floating point op: ~20 pJ *
 - Reading 64 bits from small local SRAM (1mm away on chip): ~ 26 pJ
 - Reading 64 bits from low power mobile DRAM (LPDDR): ~1200 pJ
- Implications
 - Reading 10 GB/sec from memory: ~1.6 watts
 - Entire power budget for mobile GPU: ~1 watt
(remember phone is also running CPU, display, radios, etc.)
 - iPhone 6 battery: ~7 watt-hours (note: my Macbook Pro laptop: 99 watt-hour battery)
 - Exploiting locality matters!!!

[Sources: Bill Dally (NVIDIA), Tom Olson (ARM)]

* Cost to just perform the logical operation, not counting overhead of instruction decode, load data from registers, etc.

Summary

- Today, single-thread-of-control performance is improving very slowly
 - To run programs significantly faster, programs must utilize multiple processing elements or specialized processing hardware
 - Which means you need to know how to reason about and write parallel and efficient code
- Writing parallel programs can be challenging
 - Requires problem partitioning, communication, synchronization
 - Knowledge of machine characteristics is important
 - In particular, understanding data movement!
- I suspect you will find that modern computers have tremendously more processing power than you might realize, if you just use it efficiently!

Welcome to CS149!

- Your goal between now and Thursday: Find yourself a partner!
(remember, we can do it for you!)



Prof. Kayvon



Prof. Olukotun